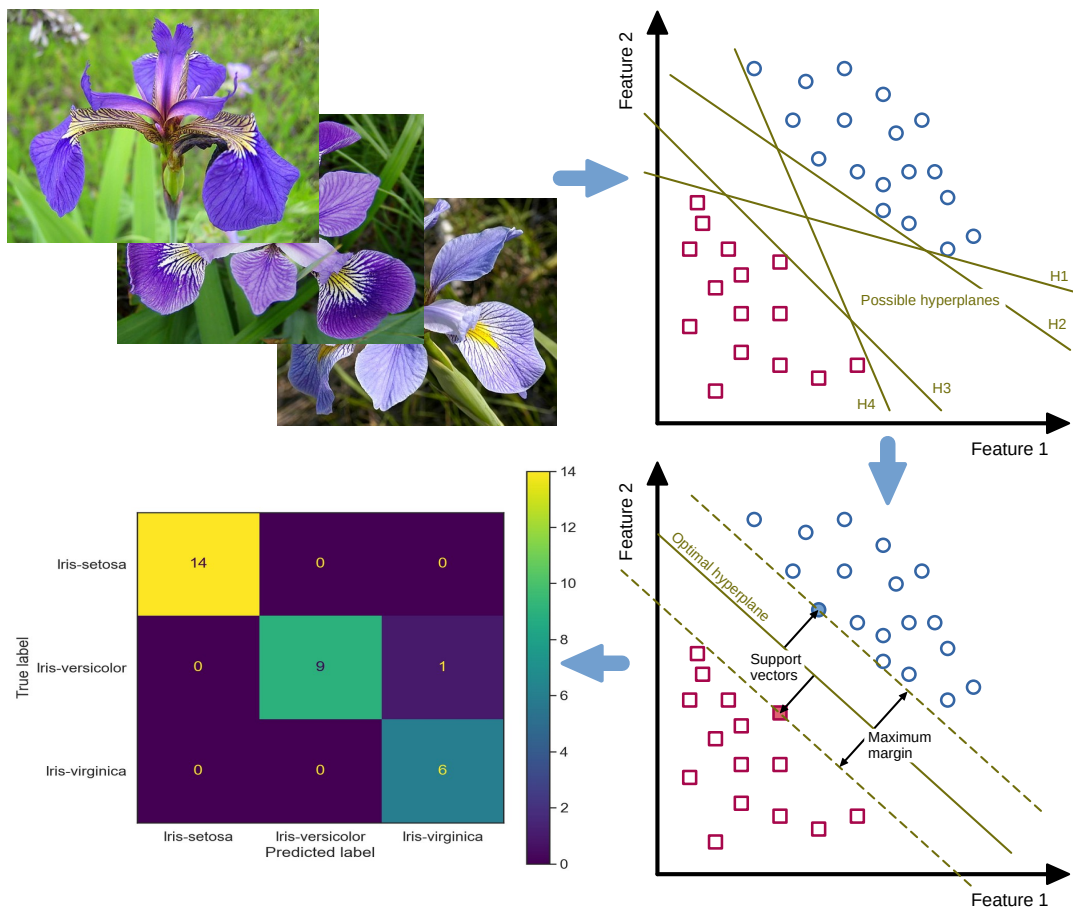# Getting started with Machine Learning (ML) and Support Vector Classifiers (SVC) - A systematic step-by-step approach

Dipl.-Ing. Björn Kasper (kasper.bjoern@bgetem.de)

Test and Certification Body for Electrical Engineering at BG ETEM

August 14, 2022

Anyone who wants to seriously deal with the hypothetical topic of our time "Artificial Intelligence (AI)" or "Machine Learning (ML)" cannot avoid dealing with the basic ML algorithms, corresponding software tools, libraries and programming systems. However, someone who opens the door for the first time to this equally very exciting as well as arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. Here, it is a good idea to consult introductory and systematic tutorials. Therefore, this Getting Started tutorial systematically demonstrates the typical ML work process step-by-step using the very powerful and performant "Support Vector Classifier (SVC)" and the widely known and exceptionally beginner-friendly "Iris Dataset". Furthermore, the selection of the "correct" SVC kernel and its parameters are described and their effects on the classification result are shown.

Contents

1 Introduction

1.1 English introduction

In the **digitized work environment**, there is an increasing demand for **Work equipment** to be able to adapt independently and in a task-related manner to changing work situations. This **situational adaptivity** can often only be realized through the use of **Artificial Intelligence (AI)** or **Machine Learning (ML)**, depending on the degree of flexibility.

Examples of such AI applications in work environments can range from comparatively simple **voice assistance systems** (similar, for example, to Siri or Alexa from the private sphere) to partially or even **fully autonomous systems**. Such fully autonomous systems are, for example, so-called **driverless transport systems**, which are autonomously driving logistics vehicles in larger industrial plants.

In addition to the numerous very interesting advantages in terms of economic efficiency, workload reduction, etc., such fully autonomous systems are characterized by a very high level of technical complexity. This concerns both their **operating functions** (e.g. autonomous navigation through complex industrial environments with shared use of the roadways by other human-controlled vehicles) and their **safety functions** (e.g. evaluation of interlinked imaging and non-imaging safety sensors for monitoring the driving space to avoid collisions).

Very high requirements are placed on such autonomous systems and the AI algorithms used for this purpose with regard to **functional safety**. However, the requirements for safety evaluability in terms of **transparency** and **explainability** of decisions made by AI are very difficult or impossible to meet, depending on the AI algorithms applied. For example, current research projects are investigating the transparency and explainability of **deep neural networks**. Furthermore, today's AI algorithms, in terms of their **recognition rates** and thus their **reliabilities**, very often do not meet the functional safety requirements to achieve higher safety levels (e.g. Performance Level d (PLd) according to ISO 13849), even under the most convenient conditions.

An appropriate assessment or even **testing** with regard to the required functional safety according to uniform and ideally standardized criteria has numerous consequences for the future orientation and organization of technical **occupational safety and health (OSH)** in Germany and in Europe. In addition to the currently still very difficult safety-related assessability, an important point is that the previous clear separation between **placing on the market law** (see e.g. Machinery Directive) and **occupational safety and health law** (see European Framework Directive for Occupational Safety and Health and German Ordinance on Occupational Safety and Health) can no longer be continued in this way. The reason for this is that the **safety-relevant properties** of the autonomous systems will change due to new or **adapted behaviors** learned during operation.

For these reasons, especially the actors of technical occupational safety and health who will deal with the evaluation of such adaptive, autonomous systems or system components with AI algorithms in the future should familiarize themselves with the AI or ML algorithms in depth as early as possible. This is the only way to ensure that the rapid development of adaptive systems capable of learning can be accompanied by OSH and their testing authorities in a constructive, critical and technically appropriate manner. If this is omitted, it must be assumed on the basis of the experiences of recent years that the OSH system will be ruthlessly circumvented or undermined by the economic interests of globally operating software giants. This would have the consequence that serious or fatal occupational accidents are likely to occur due to inadequately designed AI-based work systems.

The safety-related evaluation of such learning-capable systems requires a deeper technical entry into the world of **Artificial Intelligence (AI)** or **Machine Learning (ML)**. For this purpose, it is necessary to deal with the basic operation of typical ML algorithms, corresponding software tools, libraries and programming systems.

However, someone who opens the door for the first time to this equally very exciting as well as arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. In addition to reading general technical literature, it is advisable to consult introductory and systematic tutorials.

This Getting Started tutorial has exactly this goal, demonstrating systematically and step-by-step the typical ML workflow using the very powerful **Support Vector Classifier (SVC)** as an example.

This tutorial will be presented in the context of a workshop at the **Conference "Artificial Intelligence"**, hosted by the German Social Accident Insurance (DGUV), probably in November 2022 in Dresden. The workshop addresses interested ML novices in the technical occupational safety and health of the social accident insurance institutions.

Besides the **deep neural networks**, which are very present in the media, there is a very rich diversity of other very powerful ML algorithms - suitable for the particular use case. For a more generally comprehensible introduction, the SVC algorithm was deliberately chosen for the target audience of the workshop. Its operating principles are easy to convey to ML novices as well as in the time frame given for the workshop - quite in contrast to the entry into the world of deep neural networks.

The following main sections will demonstrate the typical ML workflow step-by-step. In **Step 0**, specific guidance is provided for selecting hardware and software suitable for machine learning. To allow an ML novice to first familiarize themselves with the ML algorithms, tools, libraries, and programming systems, the ready-made and very beginner-friendly **Iris dataset** is involved in **Step 1**. Only after a comprehensive acquaintance with the application of ML tools would it make sense to examine one's own environment for ML-suitable applications and to obtain suitable datasets from them. However, this is beyond the scope of this introductory tutorial.

One of the most important steps in the entire ML process is **Step 2**, in which the dataset included in Step 1 is examined using typical data analysis tools. In addition to exploring the **data structure** and **internal correlations** in the dataset, errors such as gaps, duplications, or obvious misentries must also be found and corrected where possible. This is enormously important so that the classification can later provide plausible results.

After exploring the dataset, in **step 3** one has to decide on a specific ML algorithm based on certain selection criteria. Among other ML algorithms suitable for the Iris dataset (such as the decision-tree-based **random-forests classifier**), the reasoned choice here in the tutorial falls on the **support vector classifier (SVC)**. A dedicated SVC model is now being implemented.

In **step 4** the dataset is prepared for the actual classification by SVC. Depending on the selected ML algorithm as well as the data structure, it may be necessary to prepare the data before training (e.g., by standardization, normalization, or binarization based on thresholds). After splitting the dataset into a training and test dataset, the SVC model is trained with the training dataset in **step 5**. Subsequently, classification predictions are made with the trained SVC model based on the test data. In **step 6**, the quality of the classification result is evaluated using known **metrics** such as the **confusion matrix**.

Since the classification in step 5 was initially performed with standard parameters (so-called **hyper-parameters**), their meaning is explained in **step 7** and then their effect on the classification result is demonstrated by manually varying the individual hyper-parameters.

In the final **Step 8**, two approaches to systematic hyper-parameter search are presented: **Grid Search** and **Randomized Search**. While the former exhaustively considers all parameter combinations for given values, the latter selects a number of candidates from a parameter space with a particular random distribution.

1.2 German introduction

Von den **Arbeitsmitteln** in der **digitalisierten Arbeitswelt** wird immer stärker gefordert, dass sie sich selbstständig und aufgabenbezogen an sich ändernde Arbeitssituationen anpassen können. Diese **situative Adaptivität** kann je nach Stärke des Flexibilisierungsgrades oft nur durch Anwendung von **Künstlicher Intelligenz (KI)** bzw. **Maschinellem Lernen (ML)** realisiert werden.

Beispiele für solche KI-Anwendungen in der Arbeitswelt reichen von vergleichsweise einfachen **Sprachassistenzsystemen** (ähnlich z. B. Siri oder Alexa aus dem privaten Umfeld) bis hin zu teil- oder gar **vollautonomen Systemen**. Solche vollautonomen Systeme sind beispielsweise sog. **fahrerlose Transportsysteme**, bei denen es sich um autonom fahrende Logistikfahrzeuge in größeren Industrieanlagen handelt.

Neben den vielen sehr interessanten Vorteilen bzgl. Wirtschaftlichkeit, Arbeitserleichterung usw. kennzeichnet solche vollautonomen Systeme eine sehr hohe technische Komplexität. Diese betrifft sowohl ihre **Betriebsfunktionen** (z. B. autonome Navigation durch komplexe industrielle Umgebungen bei gemeinsamer Nutzung der Fahrwege durch andere menschlich gesteuerte Fahrzeuge) als auch ihre **Sicherheitsfunktionen** (z. B. Auswertung miteinander verknüpfter bildgebender und nicht-bildgebender Sicherheitssensorik zur Überwachung des Fahrraums zur Kollisionsvermeidung).

An solche autonomen Systeme und die hierfür eingesetzten KI-Algorithmen werden sehr hohe Anforderungen hinsichtlich der **funktionalen Sicherheit** gestellt. Jedoch sind die Anforderungen für eine sicherheitstechnische Bewertbarkeit bezüglich der **Transparenz** und **Erklärbarkeit** der durch KI getroffenen Entscheidungen je nach verwendeten KI-Algorithmen sehr schwer bis unmöglich erreichbar. Beispielsweise werden durch aktuell laufende Forschungsprojekte die Transparenz und Erklärbarkeit von **tiefen neuronalen Netzen** untersucht. Weiterhin erfüllen heutige KI-Algorithmen hinsichtlich ihrer **Erkennungsraten** und damit ihrer **Zuverlässigkeiten** selbst unter günstigsten Bedingungen sehr oft nicht die Anforderungen an die funktionale Sicherheit, um höhere Safety-Level (z. B. Performance Level d (PLd) nach ISO 13849) zu erreichen.

Eine hinsichtlich der geforderten funktionalen Sicherheit angemessene Bewertung oder gar **Prüfung** nach einheitlichen und idealerweise genormten Maßstäben hat viele Konsequenzen für die zukünftige Ausrichtung und Gestaltung des **technischen Arbeitsschutzes** in Deutschland und in Europa. Neben der derzeit noch sehr schwierigen sicherheitstechnischen Bewertbarkeit von KI-Algorithmen ist ein wichtiger Punkt, dass die bisherige klare Trennung zwischen **Inverkehrbringensrecht** (siehe z. B. Maschinenrichtlinie) und **betrieblichem Arbeitsschutzrecht** (siehe Arbeitsschutz-Rahmenrichtlinie und Betriebssicherheitsverordnung) so nicht mehr aufrechterhalten werden kann. Grund hierfür ist, dass sich auch die **sicherheitsrelevanten Eigenschaften** der autonomen Systeme durch während des Betriebs erlernte, neue oder **angepasste Verhaltensweisen** verändern werden.

Aus diesen Gründen sollten sich insbesondere die Akteure des technischen Arbeitsschutzes, die sich zukünftig mit der Prüfung solcher lernfähigen, autonomen Systeme oder Systemkomponenten mit KI-Algorithmen befassen werden, möglichst frühzeitig mit den KI- bzw. ML-Algorithmen vertieft auseinandersetzen. Nur dadurch lässt sich erreichen, dass die stürmische Entwicklung lernfähiger, adaptiver Systeme durch den Arbeitsschutz und dessen Prüfinstitute konstruktiv, kritisch und fachlich angemessen begleitet werden kann. Wird dies versäumt, muss aufgrund der Erfahrungen der vergangenen Jahre davon ausgegangen werden, dass das Arbeitsschutzsystem durch die wirtschaftlichen Interessen global agierender Softwaregiganten skrupellos umgangen oder ausgehebelt werden wird. Dies hätte die Folge, dass schwere oder tödliche Arbeitsunfälle wegen unzulänglich gestalteter KI-basierter Arbeitssysteme wahrscheinlich werden.

Allerdings erfordert die sicherheitstechnische Bewertung solcher lernfähigen Systeme einen tiefer gehenden fachlichen Einstieg in die Welt von **Künstlicher Intelligenz (KI)** bzw. **Maschinellem Lernen (ML)**. Hierzu muss sich mit den grundlegenden Funktionsweisen typischer ML-Algorithmen, entsprechenden Software-Werkzeugen, Bibliotheken und Programmiersystemen auseinander gesetzt werden.

Wer jedoch zum ersten Mal die Tür zu dieser ebenso spannenden wie beliebig komplexen und auf den ersten Blick verwirrenden Welt öffnet, wird sehr schnell überfordert sein. Hier empfiehlt es sich neben dem Lesen allgemeiner Fachliteratur, einführende und systematische Anleitungen zu Rate zu ziehen.

Genau dieses Ziel verfolgt das vorliegende Getting-Started-Tutorial, indem systematisch und Schritt-für-Schritt der typische ML-Arbeitsablauf am Beispiel des sehr leistungsfähigen **Support Vector Classifier (SVC)** demonstriert wird.

Dieses Tutorial wird im Rahmen eines Workshops auf der **Fachtagung "Künstliche Intelligenz"**, ausgerichtet durch die Deutsche Gesetzliche Unfallversicherung (DGUV), voraussichtlich im November 2022 in Dresden vorgestellt. Der Workshop richtet sich an interessierte ML-Neulinge im technischen Arbeitsschutz der gesetzlichen Unfallversicherungsträger.

Neben den medial sehr präsenten **tiefen neuronalen Netzen** gibt es eine sehr reichhaltige Auswahl anderer sehr leistungsfähiger ML-Algorithmen - passend für den jeweiligen Anwendungsfall. Für einen allgemein verständlicheren Einstieg wurde für die Zielgruppe des Workshops der SVC-Algorithmus bewusst gewählt. Dessen Arbeitsweise ist sowohl für ML-Neulinge als auch in dem für den Workshop vorgegebenen Zeitrahmen leicht vermittelbar - ganz im Gegensatz zum Einstieg in die Welt der tiefen neuronalen Netze.

Die folgenden Hauptabschnitte demonstrieren den typischen ML-Arbeitsablauf Schritt-für-Schritt. Im **Schritt 0** werden konkrete Hinweise für die Auswahl der für das maschinelle Lernen geeigneten Hardware und Software gegeben. Damit sich ein ML-Neuling zunächst mit den ML-Algorithmen, Werkzeugen, Bibliotheken und Programmiersystemen vertraut machen kann, wird im **Schritt 1** der fertige und sehr einsteigerfreundliche **Iris-Datensatz** hinzugezogen. Erst nach einer umfassenden Einarbeitung in die Anwendung der ML-Werkzeuge wäre es sinnvoll, die eigene Umgebung auf ML-taugliche Anwendungen hin zu untersuchen und daraus geeignete Datensätze zu gewinnen. Dies geht jedoch über den Rahmen dieses einführenden Tutorials hinaus.

Mit der wichtigste Schritt im gesamten ML-Prozess ist **Schritt 2**, in dem der in Schritt 1 einbezogene Datensatz mit Hilfe typischer Datenanalyse-Werkzeuge untersucht wird. Neben der Erkundung der **Datenstruktur** sowie **innerer Zusammenhänge** im Datensatz müssen auch Fehler wie z. B. Lücken, Dopplungen oder offensichtliche Fehleingaben gefunden und nach Möglichkeit behoben werden. Dies ist enorm wichtig, damit die Klassifikation später plausible Ergebnisse liefern kann.

Nach der Erkundung des Datensatzes muss man sich im **Schritt 3** anhand bestimmter Auswahlkriterien für einen konkreten ML-Algorithmus entscheiden. Neben anderen für den Iris-Datensatz passenden ML-Algorithmen (wie z. B. der entscheidungsbaum-basierte **Random-forests-Classifier**) fällt die begründete Auswahl hier im Tutorial auf den **Support-Vector-Classifier (SVC)**. Ein entsprechendes SVC-Modell wird nun implementiert.

Im **Schritt 4** wird der Datensatz für die eigentliche Klassifikation per SVC vorbereitet. Je nach gewähltem ML-Algorithmus sowie der Datenstruktur kann es erforderlich sein, dass die Daten vor dem Training aufbereitet werden müssen (z. B. durch Standardisierung, Normalisierung oder Binärisierung anhand von Schwellwerten). Nach der Aufteilung des Datensatzes in einen Trainings- und Testdatensatz, wird das SVC-Modell im **Schritt 5** mit dem Trainingsdatensatz trainiert. Anschließend werden mit dem trainierten SVC-Modell anhand der Testdaten Klassifikationsvorhersagen getroffen. Im **Schritt 6** wird die Güte des Klassifikationsergebnisses anhand bekannter **Metriken** wie z. B. der **Konfusionsmatrix** evaluiert.

Da die Klassifikation im Schritt 5 zunächst mit Standard-Parametern (den sog. **Hyper-Parametern**) durchgeführt wurde, wird ihre Bedeutung im **Schritt 7** erklärt und danach ihr Einfluss auf das Klassifikationsergebnis durch manuelle Variation der einzelnen Hyper-Parameter demonstriert.

Im abschließenden **Schritt 8** werden zwei Ansätze zur systematischen Hyper-Parameter-Suche vorgestellt: **Grid Search** und **Randomized Search**. Während bei ersterer für gegebene Werte erschöpfend alle Parameterkombinationen betrachtet werden, wird beim zweiten Ansatz eine Anzahl von Kandidaten aus einem Parameterraum mit einer bestimmten zufälligen Verteilung ausgewählt.

1.3 Steps of the systematic ML process

The following **steps of the systematic ML process** are covered in the next main sections:

2 Load globally used libraries and set plot parameters

```python
[24]: import time

      from IPython.display import HTML

      import pandas as pd
      import numpy as np
      import matplotlib.pyplot as plt
      from sklearn import svm, metrics
      import seaborn as sns
      %matplotlib inline

      from ipypublish import nb_setup
```

3 STEP 0: Select hardware and software suitable for ML

In this step, specific guidance is provided for selecting hardware and software suitable for machine learning.

3.1 Hardware

3.1.1 General hardware requirements

3.1.2 Desktop or server based

3.1.3 Embedded application

3.2 Software

3.2.1 General requirements to the operating system

3.2.2 Programming IDEs

R and RStudio

Python and JupyterLab

3.2.3 Packages for data analytics and libraries for ML (Python only)

NumPy

Pandas

matplotlib and seaborn

Scikit-Learn

TensorFlow The package `TensorFlow` offers, among other things, the possibility to create and train **artificial neural networks (ANN)** based on Google AI. However, the installation and application is very much beyond the scope of this beginner tutorial. Further information can be found here: `https://www.tensorflow.org`.

4 STEP 1: Acquire the ML dataset

To allow an ML novice to first familiarize themselves with the ML algorithms, tools, libraries, and programming systems, the ready-made and very beginner-friendly **Iris dataset** is involved in this step. Only after a comprehensive acquaintance with the application of ML tools would it make sense to examine one's own environment for ML-suitable applications and to obtain suitable datasets from them. However, this is beyond the scope of this introductory tutorial.

Several details, for example, on the history of the creation of the Iris dataset can be found here: Iris flower datasets.

It can be downloaded on Kaggle: Iris Flower Dataset. Furthermore, the dataset is available via Python in the machine learning package Scikit-learn, so that users can access it without having to find a special source for it.

```
[2]: # import some data to play with
     irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

5 STEP 2: Explore the ML dataset

One of the most important steps in the entire ML process is this step, in which the dataset included in Step 1 is examined using typical data analysis tools. In addition to exploring the **data structure** and **internal correlations** in the dataset, errors such as **gaps**, **duplications**, or obvious **misentries** must also be found and corrected where possible. This is enormously important so that the classification can later provide plausible results.

5.1 Goals of exploration

The objectives of the exploration of the dataset are as follows:

1. Clarify the **origins history**:
   - Where did the data come from? => Contact persons and licensing permissions?
   - Who obtained the data and with which (measurement) methods? => Did systematic errors occur during the acquisition?
   - What were they originally intended for? => Can they be used for my application?
2. Overview of the internal **structure and organisation** of the data:
   - Which columns are there? => With which methods can they be read in (e.g. import of CSV files)?
   - What do they contain for (physical) measured variables? => Which technical or physical correlations exist?
   - Which data formats or types are there? => Do they have to be converted?
   - In which value ranges do the measurement data vary? => Are normalizations necessary?
3. Identify **anomalies** in the datasets:
   - Do the data have **gaps** or **duplicates**? => Does the dataset needs to be cleaned?
   - Are there obvious erroneous entries or measurement outliers? => Does (statistical) filtering have to be carried out?
4. Avoidance of **tendencies due to bias**:
   - Are all possible classes included in the dataset and equally distributed? => Does the dataset need to be enriched with additional data for balance?

5. Find a first rough **idea of which correlations** could be in the dataset

5.2 Clarify the origins history

The **Iris flower datasets** is a multivariate dataset introduced by the British statistician and biologist *Ronald Fisher* in his paper "The use of multiple measurements in taxonomic problems as an example of linear discriminant analysis" (1936). It is sometimes called *Anderson's Iris dataset* because Edgar Anderson collected the data to quantify the morphologic variation of Iris flowers of three related species (source: Iris flower dataset).

The dataset is published in Public Domain with a CC0-License.

This dataset became a typical test case for many statistical classification techniques in machine learning such as **support vector machines**.

[..] measurements of the flowers of fifty plants each of the two species *Iris setosa* and *I. versicolor*, found **growing together in the same colony** and measured by Dr E. Anderson [..] (source: R. A. Fisher (1936). "The use of multiple measurements in taxonomic problems". Annals of Eugenics)

[..] *Iris virginica*, differs from the two other samples in **not being taken from the same natural colony** [..] (source: ibidem)

5.3 Overview of the internal structure and organisation of the data

The dataset consists of 50 samples from each of three species of Iris (*Iris setosa*, *Iris virginica* and *Iris versicolor*), so there are 150 total samples. Four features were measured from each sample: the length and the width of the sepals and petals, in centimetres.
Here is a principle illustration of a flower with sepal and petal:

Figure 1: Principle illustration of a flower with sepal and petal (source: Mature_flower_diagram.svg, license: public domain)

Here are pictures of the three different Iris species (*Iris setosa*, *Iris virginica* and *Iris versicolor*). Given the dimensions of the flower, it will be possible to predict the class of the flower.

Figure 2: Left: *Iris setosa* (source: Irissetosa1.jpg, license: public domain); middle: *Iris versicolor* (source: Iris_versicolor_3.jpg, license: CC-SA 3.0); right: *Iris virginica* (source: Iris_virginica.jpg, license: CC-SA 2.0)

5.3.1 Inspect structure of dataframe

Print first or last 5 rows of dataframe:

```
[3]: irisdata_df.head()
```

```
[3]:    sepal_length  sepal_width  petal_length  petal_width      species
     0           5.1          3.5           1.4          0.2  Iris-setosa
     1           4.9          3.0           1.4          0.2  Iris-setosa
     2           4.7          3.2           1.3          0.2  Iris-setosa
     3           4.6          3.1           1.5          0.2  Iris-setosa
     4           5.0          3.6           1.4          0.2  Iris-setosa
```

```
[4]: irisdata_df.tail()
```

```
[4]:      sepal_length  sepal_width  petal_length  petal_width        species
     145           6.7          3.0           5.2          2.3  Iris-virginica
     146           6.3          2.5           5.0          1.9  Iris-virginica
     147           6.5          3.0           5.2          2.0  Iris-virginica
     148           6.2          3.4           5.4          2.3  Iris-virginica
     149           5.9          3.0           5.1          1.8  Iris-virginica
```

While printing a dataframe - only an abbreviated view of the dataframe is shown :(
Default setting in the pandas library makes it to display only 5 lines from head and from tail.

```
[5]: irisdata_df
```

```
[5]:      sepal_length  sepal_width  petal_length  petal_width         species
     0             5.1          3.5           1.4          0.2     Iris-setosa
     1             4.9          3.0           1.4          0.2     Iris-setosa
     2             4.7          3.2           1.3          0.2     Iris-setosa
     3             4.6          3.1           1.5          0.2     Iris-setosa
     4             5.0          3.6           1.4          0.2     Iris-setosa
     ..            ...          ...           ...          ...             ...
     145           6.7          3.0           5.2          2.3  Iris-virginica
     146           6.3          2.5           5.0          1.9  Iris-virginica
     147           6.5          3.0           5.2          2.0  Iris-virginica
     148           6.2          3.4           5.4          2.3  Iris-virginica
     149           5.9          3.0           5.1          1.8  Iris-virginica

     [150 rows x 5 columns]
```

To print all rows of a dataframe, the option `display.max_rows` has to set to `None` in pandas:

```
[6]: pd.set_option('display.max_rows', None)
     irisdata_df
```

[6]:

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa
10	5.4	3.7	1.5	0.2	Iris-setosa
11	4.8	3.4	1.6	0.2	Iris-setosa
12	4.8	3.0	1.4	0.1	Iris-setosa
13	4.3	3.0	1.1	0.1	Iris-setosa
14	5.8	4.0	1.2	0.2	Iris-setosa
15	5.7	4.4	1.5	0.4	Iris-setosa
16	5.4	3.9	1.3	0.4	Iris-setosa
17	5.1	3.5	1.4	0.3	Iris-setosa
18	5.7	3.8	1.7	0.3	Iris-setosa
19	5.1	3.8	1.5	0.3	Iris-setosa
20	5.4	3.4	1.7	0.2	Iris-setosa
21	5.1	3.7	1.5	0.4	Iris-setosa
22	4.6	3.6	1.0	0.2	Iris-setosa
23	5.1	3.3	1.7	0.5	Iris-setosa
24	4.8	3.4	1.9	0.2	Iris-setosa
25	5.0	3.0	1.6	0.2	Iris-setosa
26	5.0	3.4	1.6	0.4	Iris-setosa
27	5.2	3.5	1.5	0.2	Iris-setosa
28	5.2	3.4	1.4	0.2	Iris-setosa
29	4.7	3.2	1.6	0.2	Iris-setosa
30	4.8	3.1	1.6	0.2	Iris-setosa
31	5.4	3.4	1.5	0.4	Iris-setosa
32	5.2	4.1	1.5	0.1	Iris-setosa
33	5.5	4.2	1.4	0.2	Iris-setosa
34	4.9	3.1	1.5	0.1	Iris-setosa
35	5.0	3.2	1.2	0.2	Iris-setosa
36	5.5	3.5	1.3	0.2	Iris-setosa
37	4.9	3.1	1.5	0.1	Iris-setosa
38	4.4	3.0	1.3	0.2	Iris-setosa
39	5.1	3.4	1.5	0.2	Iris-setosa
40	5.0	3.5	1.3	0.3	Iris-setosa
41	4.5	2.3	1.3	0.3	Iris-setosa
42	4.4	3.2	1.3	0.2	Iris-setosa
43	5.0	3.5	1.6	0.6	Iris-setosa
44	5.1	3.8	1.9	0.4	Iris-setosa
45	4.8	3.0	1.4	0.3	Iris-setosa
46	5.1	3.8	1.6	0.2	Iris-setosa
47	4.6	3.2	1.4	0.2	Iris-setosa
48	5.3	3.7	1.5	0.2	Iris-setosa
49	5.0	3.3	1.4	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
51	6.4	3.2	4.5	1.5	Iris-versicolor
52	6.9	3.1	4.9	1.5	Iris-versicolor

53	5.5	2.3	4.0	1.3	Iris-versicolor
54	6.5	2.8	4.6	1.5	Iris-versicolor
55	5.7	2.8	4.5	1.3	Iris-versicolor
56	6.3	3.3	4.7	1.6	Iris-versicolor
57	4.9	2.4	3.3	1.0	Iris-versicolor
58	6.6	2.9	4.6	1.3	Iris-versicolor
59	5.2	2.7	3.9	1.4	Iris-versicolor
60	5.0	2.0	3.5	1.0	Iris-versicolor
61	5.9	3.0	4.2	1.5	Iris-versicolor
62	6.0	2.2	4.0	1.0	Iris-versicolor
63	6.1	2.9	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
65	6.7	3.1	4.4	1.4	Iris-versicolor
66	5.6	3.0	4.5	1.5	Iris-versicolor
67	5.8	2.7	4.1	1.0	Iris-versicolor
68	6.2	2.2	4.5	1.5	Iris-versicolor
69	5.6	2.5	3.9	1.1	Iris-versicolor
70	5.9	3.2	4.8	1.8	Iris-versicolor
71	6.1	2.8	4.0	1.3	Iris-versicolor
72	6.3	2.5	4.9	1.5	Iris-versicolor
73	6.1	2.8	4.7	1.2	Iris-versicolor
74	6.4	2.9	4.3	1.3	Iris-versicolor
75	6.6	3.0	4.4	1.4	Iris-versicolor
76	6.8	2.8	4.8	1.4	Iris-versicolor
77	6.7	3.0	5.0	1.7	Iris-versicolor
78	6.0	2.9	4.5	1.5	Iris-versicolor
79	5.7	2.6	3.5	1.0	Iris-versicolor
80	5.5	2.4	3.8	1.1	Iris-versicolor
81	5.5	2.4	3.7	1.0	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
83	6.0	2.7	5.1	1.6	Iris-versicolor
84	5.4	3.0	4.5	1.5	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
90	5.5	2.6	4.4	1.2	Iris-versicolor
91	6.1	3.0	4.6	1.4	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor
93	5.0	2.3	3.3	1.0	Iris-versicolor
94	5.6	2.7	4.2	1.3	Iris-versicolor
95	5.7	3.0	4.2	1.2	Iris-versicolor
96	5.7	2.9	4.2	1.3	Iris-versicolor
97	6.2	2.9	4.3	1.3	Iris-versicolor
98	5.1	2.5	3.0	1.1	Iris-versicolor
99	5.7	2.8	4.1	1.3	Iris-versicolor
100	6.3	3.3	6.0	2.5	Iris-virginica
101	5.8	2.7	5.1	1.9	Iris-virginica
102	7.1	3.0	5.9	2.1	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
104	6.5	3.0	5.8	2.2	Iris-virginica
105	7.6	3.0	6.6	2.1	Iris-virginica
106	4.9	2.5	4.5	1.7	Iris-virginica
107	7.3	2.9	6.3	1.8	Iris-virginica
108	6.7	2.5	5.8	1.8	Iris-vrginica
109	7.2	3.6	6.1	2.5	Iris-virginica

110	6.5	3.2	5.1	2.0	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica
113	5.7	2.5	5.0	2.0	Iris-virginica
114	5.8	2.8	5.1	2.4	Iris-virginica
115	6.4	3.2	5.3	2.3	Iris-virginica
116	6.5	3.0	5.5	1.8	Iris-virginica
117	7.7	3.8	6.7	2.2	Iris-virginica
118	7.7	2.6	6.9	2.3	Iris-virginica
119	6.0	2.2	5.0	1.5	Iris-virginica
120	6.9	3.2	5.7	2.3	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
122	7.7	2.8	6.7	2.0	Iris-virginica
123	6.3	2.7	4.9	1.8	Iris-virginica
124	6.7	3.3	5.7	2.1	Iris-virginica
125	7.2	3.2	6.0	1.8	Iris-virginica
126	6.2	2.8	4.8	1.8	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica
128	6.4	2.8	5.6	2.1	Iris-virginica
129	7.2	3.0	5.8	1.6	Iris-virginica
130	7.4	2.8	6.1	1.9	Iris-virginica
131	7.9	3.8	6.4	2.0	Iris-virginica
132	6.4	2.8	5.6	2.2	Iris-virginica
133	6.3	2.8	5.1	1.5	Iris-virginica
134	6.1	2.6	5.6	1.4	Iris-virginica
135	7.7	3.0	6.1	2.3	Iris-virginica
136	6.3	3.4	5.6	2.4	Iris-virginica
137	6.4	3.1	5.5	1.8	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
139	6.9	3.1	5.4	2.1	Iris-virginica
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

5.3.2 Get data types

```
[7]: irisdata_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column        Non-Null Count  Dtype
---  ------        --------------  -----
 0   sepal_length  150 non-null    float64
 1   sepal_width   150 non-null    float64
 2   petal_length  150 non-null    float64
 3   petal_width   150 non-null    float64
 4   species       150 non-null    object
dtypes: float64(4), object(1)
memory usage: 5.3+ KB
```

```
[8]: irisdata_df.describe()
```

```
[8]:        sepal_length  sepal_width  petal_length  petal_width
count    150.000000   150.000000    150.000000   150.000000
mean       5.843333     3.054000      3.758667     1.198667
std        0.828066     0.433594      1.764420     0.763161
min        4.300000     2.000000      1.000000     0.100000
25%        5.100000     2.800000      1.600000     0.300000
50%        5.800000     3.000000      4.350000     1.300000
75%        6.400000     3.300000      5.100000     1.800000
max        7.900000     4.400000      6.900000     2.500000
```

5.3.3 Get data ranges with Boxplots

Boxplots can be used to explore the data ranges in the dataset. These also provide information about **outliers**.

```
[9]: sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.0})
sns.set_style("whitegrid")
#sns.set_style("white")


fig, axs = plt.subplots(2, 2, figsize=(12, 10))


fn = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
box1 = sns.boxplot(x = 'species', y = 'sepal_length',
                   data = irisdata_df, order = cn, ax = axs[0,0])
box2 = sns.boxplot(x = 'species', y = 'sepal_width',
                   data = irisdata_df, order = cn, ax = axs[0,1])
box3 = sns.boxplot(x = 'species', y = 'petal_length',
                   data = irisdata_df, order = cn, ax = axs[1,0])
box4 = sns.boxplot(x = 'species', y = 'petal_width',
                   data = irisdata_df,  order = cn, ax = axs[1,1])

# add some spacing between subplots
fig.tight_layout(pad=2.0)

plt.show()
```

5.4 Identify anomalies in the datasets

5.4.1 Find gaps in dataset

This section was inspired by Working with Missing Data in Pandas.

Checking for missing values using `isnull()`   In order to check for missing values in Pandas DataFrame, we use the function `isnull()`. This function returns a dataframe of Boolean values which are True for **NaN values**.

```
[10]: pd.set_option('display.max_rows', 40)
pd.set_option('display.min_rows', 30)
```

```
[11]: irisdata_df.isnull()
```

```
[11]:    sepal_length  sepal_width  petal_length  petal_width  species
0         False        False         False        False    False
```
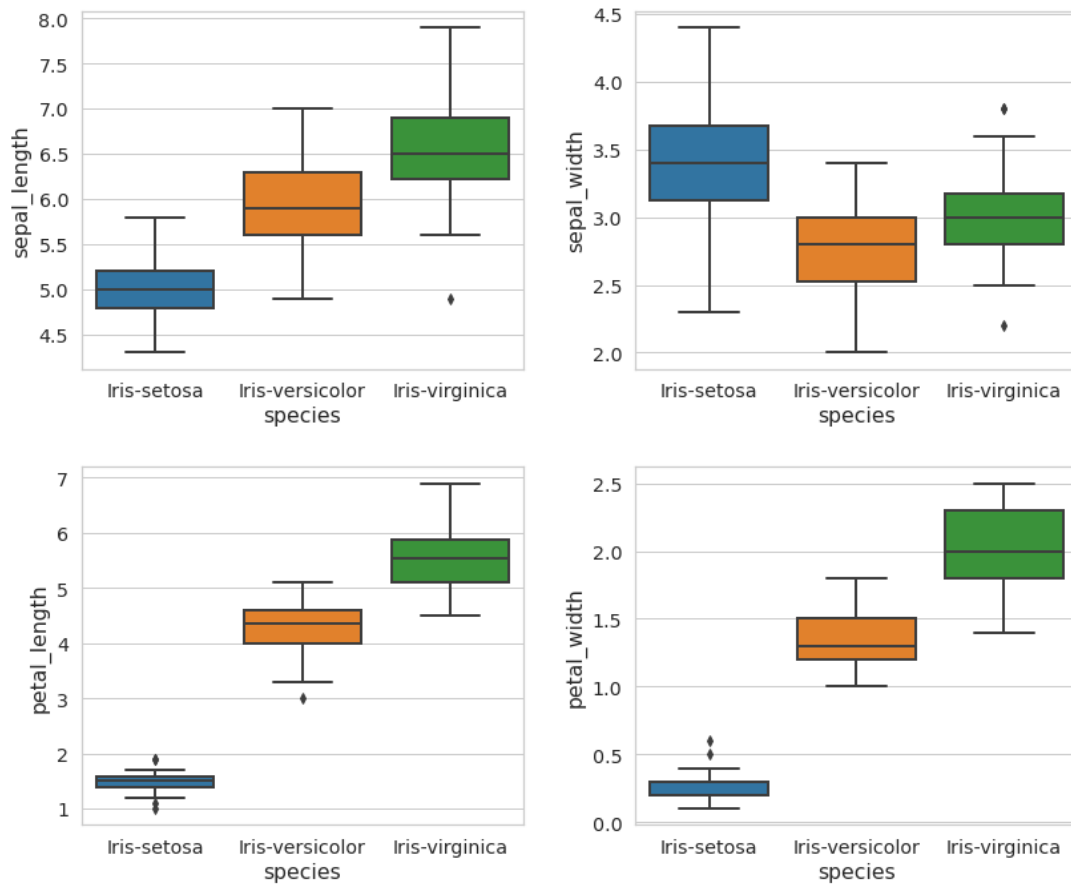
Figure 3: Boxplots used to explore the data ranges in the Iris dataset

1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
5	False	False	False	False	False
6	False	False	False	False	False
7	False	False	False	False	False
8	False	False	False	False	False
9	False	False	False	False	False
10	False	False	False	False	False
11	False	False	False	False	False
12	False	False	False	False	False
13	False	False	False	False	False
14	False	False	False	False	False
..	...	...	...	...	...
135	False	False	False	False	False
136	False	False	False	False	False
137	False	False	False	False	False
138	False	False	False	False	False
139	False	False	False	False	False
140	False	False	False	False	False
141	False	False	False	False	False
142	False	False	False	False	False
143	False	False	False	False	False
144	False	False	False	False	False
145	False	False	False	False	False

146	False	False	False	False	False
147	False	False	False	False	False
148	False	False	False	False	False
149	False	False	False	False	False

```
[150 rows x 5 columns]
```

Show only the gaps:

```
[12]: irisdata_df_gaps = irisdata_df[irisdata_df.isnull().any(axis=1)]
      irisdata_df_gaps
```

```
[12]: Empty DataFrame
      Columns: [sepal_length, sepal_width, petal_length, petal_width, species]
      Index: []
```

Fine - this dataset seems to be complete :)

So let's look for something else for exercise: employes.csv

```
[13]: pd = nb_setup.setup_pandas(escape_latex = False)

      # import data to dataframe from csv file
      employees_df = pd.read_csv("./datasets/employees_edit.csv")

      # highlight cells with nan values
      employees_df_hl = employees_df.style.highlight_null('yellow')
      employees_df_hl

      #employees_df
```

```
[13]: <pandas.io.formats.style.Styler at 0xa7a14490>
```

Show only the gaps from this gappy dataset again:

```
[14]: employees_df_gaps = employees_df[employees_df.isnull().any(axis=1)]

      # highlight cells with nan values
      employees_df_gaps = employees_df_gaps.style.highlight_null('yellow')

      employees_df_gaps
```

```
[14]: <pandas.io.formats.style.Styler at 0xa7bc2e50>
```

Fill the missing values with `fillna()` Now we are going to fill all the null (NaN) values in Gender column with *"No Gender"*.

Warning: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

```
[15]: # filling a null values using fillna()
      employees_df["Gender"].fillna("No Gender", inplace = True)

      # highlight cells by condition
      employees_df_hl = employees_df.style.apply(lambda x: ["background: yellow" if v ==␣
       ↪'No Gender' else "" for v in x], axis = 1)
      employees_df_hl
```

```
#employees_df
```

[15]: <pandas.io.formats.style.Styler at 0xa77712b0>

Dropping missing values using `dropna()` In order to drop null values from a dataframe, we use `dropna()` function. This function drops rows or columns of datasets with NaN values in different ways.

Default is to drop rows with at least 1 null value (NaN). Giving the parameter `how = 'all'` the function drops rows with all data missing or contain null values (NaN).

[25]:
```
# making a new dataframe with dropped NaN values
employees_df_dropped = employees_df.dropna(axis = 0, how ='any')
employees_df_dropped
```

[25]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus % \
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
15	Lillian	Female	6/5/2016	6:09 AM	59414	1256.00
16	Jeremy	Male	9/21/2010	5:56 AM	90370	7369.00
17	Shawn	Male	12/7/1986	7:45 PM	111737	6414.00
...	...	...	...	...	...	...
989	Stephen	No Gender	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	No Gender	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
8	True	Engineering
9	True	Business Development
11	True	Legal

```
12              True       Human Resources
13              False                Sales
14              True               Finance
15              False              Product
16              False      Human Resources
17              False              Product
...               ...                  ...
989             False                Legal
990             False            Marketing
991             True               Finance
992             False      Human Resources
993             False                Legal
994             True        Client Services
995             True             Marketing
996             True               Finance
997             True            Engineering
998             True             Marketing
999             False           Distribution
1000            False              Finance
1001            False              Product
1002            False  Business Development
1003            True                 Sales

[903 rows x 8 columns]
```

Finally we compare the sizes of dataframes so that we learn how many rows had at least 1 Null value.

```
[26]: print("Old data frame length:", len(employees_df))
      print("New data frame length:", len(employees_df_dropped))
      print("Number of rows with at least 1 NaN value: ",
            (len(employees_df)-len(employees_df_dropped)))
```

```
Old data frame length: 1004
New data frame length: 903
Number of rows with at least 1 NaN value:  101
```

5.4.2 Find and remove duplicates in dataset

This section was inspired by: - How to Find Duplicates in Pandas DataFrame (With Examples) - How to Drop Duplicate Rows in a Pandas DataFrame

Checking for duplicate values using `duplicated()` In order to check for duplicate values in Pandas DataFrame, we use a function `duplicated()`. This function can be used in two ways: - find duplicate rows across **all columns** with duplicateRows = df[df.duplicated()] - find duplicate rows across **specific columns** duplicateRows = df[df.duplicated(subset=['col1', 'col2'])]

Find duplicate rows across **all columns**:

```
[18]: # import (again) data to dataframe from csv file
      employees_df = pd.read_csv("./datasets/employees_edit.csv")
```

```
[19]: # find duplicate rows across all columns
      duplicateRows = employees_df[employees_df.duplicated()]
      duplicateRows
```

```
[19]:      First Name  Gender  Start Date Last Login Time  Salary  Bonus %  \
      112       Karen  Female  11/30/1999         7:46 AM  102488  17653.0
```

```
127      Linda  Female     5/25/2000                5:45 PM  119009   12506.0
296    Brandon     NaN     11/3/1997                8:17 PM  121333   15295.0
580  Nicholas    Male      3/1/2013                9:26 PM  101036    2826.0

     Senior Management                   Team
112               True                Product
127               True  Business Development
296              False  Business Development
580               True       Human Resources
```

```
[20]:  # argument keep='last' displays the first duplicate rows instead of the last
       duplicateRows = employees_df[employees_df.duplicated(keep='last')]
       duplicateRows
```

```
[20]:      First Name  Gender  Start Date Last Login Time  Salary  Bonus %  \
       55       Karen  Female  11/30/1999         7:46 AM  102488  17653.0
       92       Linda  Female   5/25/2000         5:45 PM  119009  12506.0
       153    Brandon     NaN   11/3/1997         8:17 PM  121333  15295.0
       442   Nicholas    Male    3/1/2013         9:26 PM  101036   2826.0

           Senior Management                   Team
       55               True                Product
       92               True  Business Development
       153             False  Business Development
       442              True       Human Resources
```

Find duplicate rows across **specific columns**:

```
[21]:  # identify duplicate rows across 'First Name' and 'Last Login Time' columns
       duplicateRows = employees_df[employees_df.duplicated(
                       subset=['First Name', 'Last Login Time'])]
       duplicateRows
```

```
[21]:      First Name  Gender  Start Date Last Login Time   Salary  Bonus %  \
       112      Karen  Female  11/30/1999         7:46 AM   102488  17653.0
       127      Linda  Female   5/25/2000         5:45 PM   119009  12506.0
       296    Brandon     NaN   11/3/1997         8:17 PM   121333  15295.0
       577        NaN  Female   1/13/2009         1:01 PM   118736   7421.0
       580   Nicholas    Male    3/1/2013         9:26 PM   101036   2826.0
       632        NaN     NaN    9/2/1988        12:49 PM   147309   1702.0
       881        NaN    Male    9/5/1980         7:36 AM   114896  13823.0
       929        NaN  Female   8/23/2000         4:19 PM    95866  19388.0
       934      Nancy  Female   9/10/2001        11:57 PM    85213   2386.0
       973      Linda  Female    2/4/2010         8:49 PM    44486  17308.0

           Senior Management                   Team
       112              True                Product
       127              True  Business Development
       296             False  Business Development
       577               NaN       Client Services
       580              True       Human Resources
       632               NaN          Distribution
       881               NaN       Client Services
       929               NaN                 Sales
       934              True             Marketing
       973              True           Engineering
```

```
[22]: # argument keep='last' displays the first duplicate rows instead of the last
      duplicateRows = employees_df[employees_df.duplicated(
                        subset=['First Name', 'Last Login Time'], keep='last')]
      duplicateRows
```

```
[22]:      First Name  Gender  Start Date Last Login Time  Salary   Bonus %  \
      23          NaN    Male   6/14/2012         4:19 PM  125792   5042.00
      37        Linda  Female  10/19/1981         8:49 PM   57427   9557.00
      55        Karen  Female  11/30/1999         7:46 AM  102488  17653.00
      66        Nancy  Female  12/15/2012        11:57 PM  125250   2672.00
      92        Linda  Female   5/25/2000         5:45 PM  119009  12506.00
      153     Brandon     NaN   11/3/1997         8:17 PM  121333  15295.00
      222         NaN  Female   6/17/1991        12:49 PM   71945      5.56
      269         NaN    Male    2/4/2005         1:01 PM   40451  16044.00
      442    Nicholas    Male    3/1/2013         9:26 PM  101036   2826.00
      778         NaN  Female   6/18/2000         7:36 AM  106428  10867.00

          Senior Management                 Team
      23                NaN                  NaN
      37               True      Client Services
      55               True              Product
      66               True  Business Development
      92               True  Business Development
      153             False  Business Development
      222               NaN            Marketing
      269               NaN         Distribution
      442              True      Human Resources
      778               NaN                  NaN
```

Dropping duplicate values using `drop_duplicates()` In order to drop duplicate values from a dataframe, we use `drop_duplicates()` function.

This function can be used in two ways: - remove duplicate rows across **all columns** with `df.drop_duplicates()` - find duplicate rows across **specific columns** `df.drop_duplicates(subset=['col1', 'col2'])`

Warning: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

Remove duplicate rows across **all columns**:

```
[23]: # remove duplicate rows across all columns
      employees_df.drop_duplicates(inplace=True)
      employees_df
```

```
[23]:     First Name  Gender  Start Date Last Login Time  Salary   Bonus %  \
      0      Douglas    Male    8/6/1993        12:42 PM   97308   6945.00
      1       Thomas    Male   3/31/1996         6:53 AM   61933      4.17
      2        Maria  Female   4/23/1993        11:17 AM  130590  11858.00
      3        Jerry    Male    3/4/2005         1:00 PM  138705      9.34
      4        Larry    Male   1/24/1998         4:47 PM  101004   1389.00
      5       Dennis    Male   4/18/1987         1:35 AM  115163  10125.00
      6         Ruby  Female   8/17/1987         4:20 PM   65476  10012.00
      7          NaN  Female   7/20/2015        10:43 AM   45906  11598.00
      8       Angela  Female  11/22/2005         6:29 AM   95570  18523.00
      9      Frances  Female    8/8/2002         6:51 AM  139852   7524.00
      10      Louise  Female   8/12/1980         9:01 AM   63241  15132.00
      11       Julie  Female  10/26/1997         3:19 PM  102508  12637.00
```

```
12       Brandon   Male    12/1/1980       1:08 AM  112807  17492.00
13          Gary   Male     1/27/2008     11:40 PM  109831   5831.00
14       Kimberly Female    1/14/1999      7:13 AM   41426  14543.00
…           …       …           …            …        …        …
989       Stephen    NaN     7/10/1983      8:10 PM   85668   1909.00
990         Donna Female    11/26/1982      7:04 AM   82871  17999.00
991        Gloria Female     12/8/2014      5:08 AM  136709  10331.00
992         Alice Female     10/5/2004      9:34 AM   47638  11209.00
993        Justin    NaN     2/10/1991      4:58 PM   38344   3794.00
994         Robin Female     7/24/1987      1:35 PM  100765  10982.00
995          Rose Female     8/25/2002      5:12 AM  134505  11051.00
996       Anthony   Male    10/16/2011      8:35 AM  112769  11625.00
997          Tina Female     5/15/1997      3:53 PM   56450     19.04
998        George   Male     6/21/2013      5:47 PM   98874   4479.00
999         Henry    NaN    11/23/2014      6:09 AM  132483  16655.00
1000       Phillip  Male     1/31/1984      6:30 AM   42392  19675.00
1001       Russell  Male     5/20/2013     12:39 PM   96914   1421.00
1002         Larry  Male     4/20/2013      4:45 PM   60500  11985.00
1003        Albert  Male     5/15/2012      6:24 PM  129949  10169.00

      Senior Management                    Team
0                 True              Marketing
1                 True                    NaN
2                False                Finance
3                 True                Finance
4                 True        Client Services
5                False                  Legal
6                 True                Product
7                  NaN                Finance
8                 True            Engineering
9                 True   Business Development
10                True                    NaN
11                True                  Legal
12                True        Human Resources
13               False                  Sales
14                True                Finance
…                   …                      …
989              False                  Legal
990              False              Marketing
991               True                Finance
992              False        Human Resources
993              False                  Legal
994               True        Client Services
995               True              Marketing
996               True                Finance
997               True            Engineering
998               True              Marketing
999              False           Distribution
1000             False                Finance
1001             False                Product
1002             False   Business Development
1003              True                  Sales

[1000 rows x 8 columns]
```

Remove duplicate rows across **specific columns**:

[24]:
```
# remove duplicate rows across 'First Name' and 'Last Login Time' columns
employees_df.drop_duplicates(
    subset=['First Name', 'Last Login Time'], keep='last', inplace=True)
employees_df
```

[24]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus % \
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
...	...	...	...	...	...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...	...	...
989	False	Legal
990	False	Marketing

```
991              True              Finance
992             False     Human Resources
993             False               Legal
994              True     Client Services
995              True           Marketing
996              True             Finance
997              True         Engineering
998              True           Marketing
999             False        Distribution
1000            False             Finance
1001            False             Product
1002            False  Business Development
1003             True               Sales

[994 rows x 8 columns]
```

5.5 Avoidance of tendencies due to bias

The description of the Iris dataset says, that it consists of **50 samples** from **each of three species** of Iris (Iris setosa, Iris virginica and Iris versicolor), so there are **150 total samples**.

But how to prove it?

5.5.1 Count occurrences of unique values

To prove whether all possible classes included in the dataset and equally distributed, you can use the function `df.value_counts`.

Following parameters can be used for fine tuning: - `dropna=False` causes that NaN values are included - `normalize=True`: relative frequencies of the unique values are returned - `ascending=False`: sort resulting classes descending

```python
[25]:  # import (again) data to dataframe from csv file
       employees_df = pd.read_csv("./datasets/employees_edit.csv")
```

```python
[26]:  # count unique values without missing values in a column,
       # ordered descending and normalized
       irisdata_df['species'].value_counts(ascending=False, dropna=False, normalize=True)
```

```
[26]: Iris-setosa        0.333333
      Iris-versicolor    0.333333
      Iris-virginica     0.333333
      Name: species, dtype: float64
```

```python
[27]:  # count unique values and missing values in a column,
       # ordered descending and not absolute values
       employees_df['Team'].value_counts(ascending=False, dropna=False, normalize=False)
```

```
[27]: Client Services       106
      Business Development   103
      Finance               102
      Marketing              98
      Product                96
      Sales                  94
      Engineering            92
      Human Resources        92
      Distribution           90
```

```
Legal                    88
NaN                      43
Name: Team, dtype: int64
```

5.5.2 Display Histogram

This section was inspired by: Pandas Histogram – DataFrame.hist().

Histograms represent **frequency distributions** graphically. This requires the separation of the data into classes (so-called **bins**).

These classes are represented in the histogram as rectangles of equal or variable width. The height of each rectangle then represents the (relative or absolute) **frequency density**.

```
[28]: employees_df.hist(column=['Salary'])
      plt.show()
```


Figure 4: Histogram for frequency distribution of the salary

```
[29]: employees_df.hist(column='Salary', by='Gender')
      plt.show()
```


Figure 5: Histogram for the frequency distribution of the salary in comparison between men and women

5.6 First idea of correlations in dataset

To get a rough idea of the **dependencies** and **correlations** in the dataset, it can be helpful to visualize the whole dataset in a **correlation heatmap**. They show in a glance which variables are correlated, to what degree and in which direction.

Later, 2 particularly well correlated variables are selected from the dataset and plotted in a **scatterplot**.

5.6.1 Visualise data with correlation heatmap

This section was inspired by How to Create a Seaborn Correlation Heatmap in Python?.

> **Correlation matrices** are an **essential tool of exploratory data analysis**. Correlation heatmaps contain the same information in a visually appealing way. What more: they show in a glance which variables are correlated, to what degree, in which direction, and alerts us to potential multicollinearity problems (source: ibidem).

Simple correlation matrix Because **string values can never be correlated**, the class names (species) have to be converted first:

```
[30]:  # encoding the class column
       irisdata_df_enc = irisdata_df.replace({"species":  {"Iris-setosa":0,
                                                            "Iris-versicolor":1,
                                                            "Iris-virginica":2}})

       #irisdata_df_enc
```

```
[31]:  irisdata_df_enc.corr()
```

```
[31]:              sepal_length  sepal_width  petal_length  petal_width   species
       sepal_length     1.000000    -0.109369      0.871754     0.817954  0.782561
       sepal_width     -0.109369     1.000000     -0.420516    -0.356544 -0.419446
       petal_length     0.871754    -0.420516      1.000000     0.962757  0.949043
       petal_width      0.817954    -0.356544      0.962757     1.000000  0.956464
       species          0.782561    -0.419446      0.949043     0.956464  1.000000
```

Correlation heatmap Choose the color sets from color map.

```
[32]:  # increase the size of the heatmap
       plt.figure(figsize=(16, 6))

       # store heatmap object in a variable to easily access it
       # when you want to include more features (such as title)
       # set the range of values to be displayed on the colormap from -1 to 1,
       # and set 'annotation=True' to display the correlation values on the heatmap
       heatmap = sns.heatmap(irisdata_df_enc.corr(), vmin=-1, vmax=1,
                             annot=True, cmap='PRGn_r')

       # give a title to the heatmap
       # 'pad=12' defines the distance of the title from the top of the heatmap
       heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
       plt.show()
```

Triangle correlation heatmap When looking at the correlation heatmaps above, you would not lose any information by **cutting** away half of it **along the diagonal** line marked by 1-s.

The **numpy** function `np.triu()` can be used to isolate the upper triangle of a matrix while turning all the values in the lower triangle into 0.

Figure 6: Correlation heatmap to explore coherences between single variables in the iris dataset

```
[33]: np.triu(np.ones_like(irisdata_df_enc.corr()))
```

```
[33]: array([[1., 1., 1., 1., 1.],
             [0., 1., 1., 1., 1.],
             [0., 0., 1., 1., 1.],
             [0., 0., 0., 1., 1.],
             [0., 0., 0., 0., 1.]])
```

Use this mask to cut the heatmap along the diagonal:

```
[34]: plt.figure(figsize=(16, 6))

      # define the mask to set the values in the upper triangle to 'True'
      mask = np.triu(np.ones_like(irisdata_df_enc.corr(), dtype=bool))

      heatmap = sns.heatmap(irisdata_df_enc.corr(), mask=mask,
                            vmin=-1, vmax=1, annot=True, cmap='PRGn_r')

      heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
      plt.show()
```


Figure 7: Correlation heatmap, which was cut at its main diagonal without losing any information

As a result from the **heatmaps** we can see, that the shape of the **petals** are the **most correlationed**

columns (0.96) with the **type of flowers** (species classes).

Somewhat lower correlates **sepal length** with **petal length** (0.87).

5.6.2 Visualise data with scatter plot

In the following, Seaborn is applied which is a library for making statistical graphics in Python. It is built on top of matplotlib and closely integrated with pandas data structures.

To investigate whether there are dependencies (e.g. correlations) in `irisdata_df` between individual variables in the dataset, it is advisable to plot them in a **scatter plot**.

```python
[35]: # There are five preset seaborn themes: darkgrid, whitegrid, dark, white, and ticks.
sns.set_style("whitegrid")
# set scale of fonts
sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.5})

# 'sepal_length', 'petal_length' are iris feature data
# 'height' used to define height of graph
# 'hue' stores the class/label of iris dataset
sns.FacetGrid(irisdata_df, hue ="species",
              height = 7).map(plt.scatter,
                              'petal_width',
                              'petal_length').add_legend()


plt.title('Scatterplot of petal length and width')
plt.show()
```

5.6.3 Visualise data with pairs plot

For systematic investigation of dependencies, all variables (each against each) are plotted in separate scatter plots.

With this so called **pairs plot** it is possible to see both **relationships** between two variables and **distribution** of single variables.

This function will create a grid of Axes such that **each numeric variable** in `irisdata_df` will by shared in the y-axis across a single row and in the x-axis across a single column.

```python
[36]: sns.set(font_scale=1.0)
sns.set_style("white")

g = sns.pairplot(irisdata_df, diag_kind="kde", hue='species',
                 palette='Dark2', height=2.5)

g.map_lower(sns.kdeplot, levels=4, color=".2")
# y .. padding between title and plot
g.fig.suptitle('Pairs plot of the Iris dataset', y=1.05)
plt.show()
```

6 STEP 3: Choose and create the ML model

After exploring the dataset, in this step one has to decide on a specific ML algorithm based on certain selection criteria.

However, since the AI or ML world is so huge and impossible for a ML novice to overlook, a brief description of the **relationship between AI and ML** is given in the following sections. Furthermore, a **taxonomy** of the different **learning types** is presented by also providing some example algorithms.

Figure 8: Plotting two individual variables of the iris dataset in the scatterplot to explore the relationships between these two

6.1 Short overview of the AI world and its ML algorithms

6.1.1 Relationship between AI, ML and others

In the **science world**, the term **artificial intelligence (AI)** refers to machines and systems that are capable of performing tasks that are characteristic of human intelligence.

In the **business world**, on the other hand, AI typically refers to mechanisms that perceive environmental factors and take autonomous actions. This is seen as an opportunity to achieve **predefined goals** with maximum success - without human intervention. Ultimately, this view is a mapping of **input information** to controlled **output actions** of a system. This expectation of AI-driven systems is thus hardly higher than what can be expected from today's modern automation systems.

Machine Learning (ML), on the other hand, addresses the mathematical models and algorithms that enable a computer system to recognize (new) correlations in huge amounts of sample data from various sources by inferring them independently. For scientists, machine learning is a subset of AI.

The umbrella term AI covers a very large research area. It includes a number of techniques that enable computers to learn independently and solve complex problems:

- Computer-Vision (CV)
- Supervised and Unsupervised Learning
- Reinforcement Learning and Genetic Algorithms
- Computational Linguistics
- Robotics
- etc.

The following Venn diagram shows the relationship between Artificial Intelligence (AI), Machine Learning (ML) and other integrated technologies. The quantities that do not belong to the main category represent

Figure 9: Plot all individual variables of the Iris dataset in pairs plot to see both the relationships between two variables and the distribution of the individual variables

techniques that can function as stand-alone techniques and do not necessarily fall into the artificial intelligence group in all cases (for further details see Emerging technologies based on artificial intelligence to assess quality and consumer preference of beverages).

Figure 10: Venn diagram showing the relationship between Artificial Intelligence (AI), Machine Learning (ML) and other integrated technologies (source: Kasper, adapted from Emerging technologies based on artificial intelligence to assess quality and consumer preference of beverages, license: CC-BY-SA 4.0)

6.1.2 Taxonomy of machine learning

The field of machine learning can be divided into the following **types of learning**:

- Supervised learning
- Unsupervised learning
- Semi-supervised learning
- Reinforcement learning

Here are some further sources:

- Taxonomy of machine learning algorithms
- Comprehensive Survey of Machine Learning Approaches in Cognitive Radio-Based Vehicular Ad Hoc Networks
- A Taxonomy of Machine Learning Techniques
- ML Algorithms: One SD
- Machine Learning Map

Supervised learning The goal of **supervised learning (SL)** is to learn a **function** that maps a **input to an output**, based on example input-output pairs. This involves inferring a relationship describable by a mathematical function from **labeled training data** consisting of a set of training examples (see Supervised Learning).

A few well-known algorithms from the field of **supervised learning** are mentioned here:

- Naive Bayes
- Linear Regression
- Logistic Regression
- Artificial Neural Networks (ANN)
- Support Vector Classifier (SVC)
- Decision Trees
- Random Forests

Unsupervised learning The algorithms of this category look for internal structures in the data of a dataset, such as **grouping** or **clustering of data points**. These algorithms can thus learn relationships from test data that have not been labeled, classified, or categorized. Rather than responding to feedback (as in supervised learning), unsupervised learning algorithms detect **commonalities in the data** and respond based on the presence or absence of such commonalities in each new dataset (see Unsupervised learning).

Here are some algorithms from the field of **unsupervised learning**:

- K-means Clustering
- Spectral Clustering
- Hierarchical Clustering
- Principal Component Analysis (PCA)

Semi-supervised learning This type of learning falls between **unsupervised** learning (without any labeled training data) and **supervised** learning (with completely labeled training data). Some of the training examples are missing training labels, yet many machine-learning researchers have found that unlabeled data, when used in conjunction with a small amount of labeled data, can produce a considerable improvement in learning accuracy (source: Semi-supervised learning).

Reinforcement learning This is an area of machine learning concerned with how **intelligent agents** ought to **take actions in an environment** in order to maximize the notion of cumulative **reward**. Due to its generality, the field is studied in many other disciplines, such as **game theory** and **control theory**.

Reinforcement learning differs from supervised learning in **not needing labeled input/output pairs** be presented, and in not needing sub-optimal actions to be explicitly corrected. Instead the focus is

on **finding a balance** between **exploration** (of uncharted territory) and **exploitation** (of current knowledge) (source: Reinforcement learning).

Here are some algorithms from the field of **reinforcement learning**:

- Iterative Policy
- Q-Learning
- SARSA
- Learning Classifiers
- Stochastic Gradient
- Genetic Algorithm

6.2 Decision graph for selecting an suitable algorithm

Now that the iris dataset has been analyzed in terms of its data structure and internal correlations, the most difficult task on the way to solving a problem using machine learning arises: finding the "right" ML algorithm (also called **estimator**).

The diverse estimators available are more or less well qualified for the respective problems with their partly very different data types. The good news is that the ML software package **Scikit-Learn** provides the following **flowchart** as a rough **guide** in choosing the right estimator for the particular task (see: Choosing the right estimator).

However, it must also be emphasized that a considerable **level of experience** through systematic trial and error is crucial to be successful in finding an "optimal" estimator.
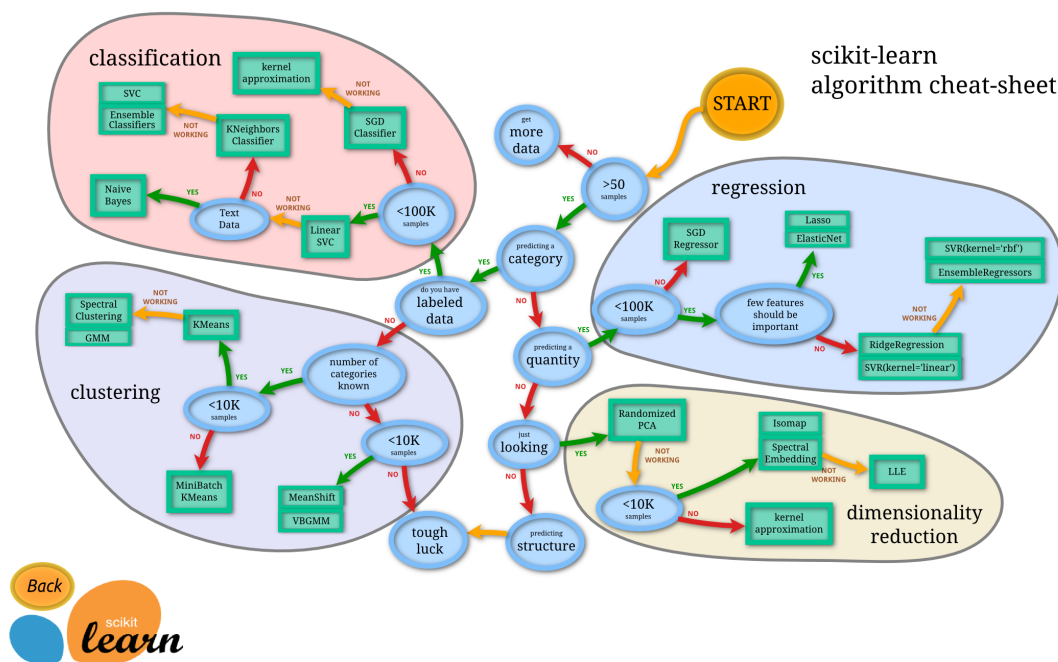
Figure 11: Decision graph for choosing an appropriate ML algorithm (source: Choosing the right estimator, license: unknown)

6.3 Reasons for choosing Support Vector Classifier (SVC)

Among other ML algorithms suitable for the Iris dataset (such as the decision-tree-based **random-forests classifier**), the reasoned choice here in this tutorial falls on the **support vector classifier (SVC)**.

The following **reasons** led to the decision for the **Support Vector Classifier (SVC)**:

- the aim is to predict the species using unlabeled test data, so the task is to **classify**

- the iris dataset is **fully labeled** (by designating the iris species)
- the dataset contains significantly **less than 100k samples**

But the most important reason is that it is **easy to understand** how it works - so it is exactly suitable for a beginner tutorial ;)

6.4 Operating principal of SVC

Support Vector Classifiers (SVC) try to **find the best hyperplane to separate** the different classes by maximizing the distance between sample points and the hyperplane (source: In Depth: Parameter tuning for SVC).

The figure **??** shows the operating principal of the SVC algorithm: the hyperplanes *H1* till *H4* (left graphic) do separate the classes. A good separation is achieved by the hyperplane that has the largest distance to the nearest training-data point of any class (so-called functional margin), since in general the larger the margin, the lower the generalization error of the classifier (source: Support-vector machine).

The right graphic shows the optimal hyperplane characterized by maximizing the margin between the classes. The perpendicular distance of the closest data points to the hyperplane determines their position and orientation. These perpendicular distances are the **support vectors** of the hyperplane - this is how the algorithm got its name.
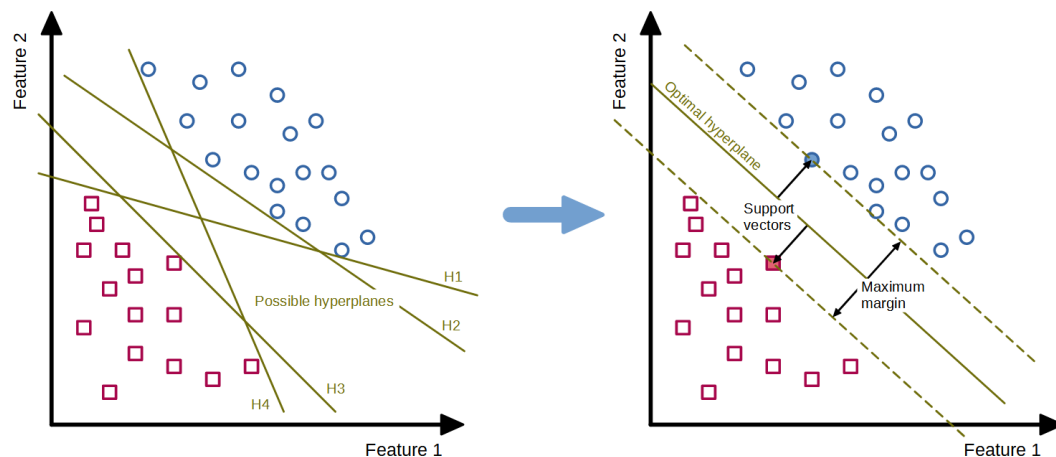
Figure 12: Support Vector Classifiers (SVC) separate the data points in classes by finding the best hyperplane by maximizing the margin to its support vectors (source: Kasper, license: CC-BY-SA 4.0)

6.5 Create the SVC model

In this step we create the SVC model choosing a **linear kernel** with default parameters.

```
[60]: from sklearn.svm import SVC
      classifier = SVC(kernel = 'linear', random_state = 0)
```

7 STEP 4: Prepare the dataset for training

In this step the dataset is prepared for the actual classification by SVC. Depending on the selected ML algorithm as well as the data structure, it may be necessary to prepare the data before training (e.g., by **standardization**, **normalization**, or **binarization** based on thresholds). Furthermore, errors in the dataset (e.g. **data gaps**, **duplicates** or obvious **misentries**) should be corrected now at the latest.

Through the intensive exploration of the data in (STEP 2: Explore the ML dataset), we know that special **preparation** of the data is **not necessary**. The values are complete and without gaps and there

are no duplicates. The values are in similar ranges, which **does not require normalization** of the data.

Furthermore, we know that the **classes** are very **evenly distributed** and thus bias tendencies should be avoided.

8 STEP 5: Carry out training, prediction and testing

8.1 Split the dataset

In the next very important step, the dataset is split into **2 subsets**: a **training dataset** and a **test dataset**. As the names suggest, the training dataset is used to train the ML algorithm. The test dataset is then used to check the quality of the trained ML algorithm (here the **recognition rate**). For this purpose, the **class labels** are **removed** from the training dataset - after all, these are to be predicted.

Typically, the **test dataset** should contain about **20%** of the entire dataset.

```python
[16]: from sklearn.model_selection import train_test_split

      # Load dataset again
      irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')

      X = irisdata_df.drop('species', axis=1)
      y = irisdata_df['species']

      X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

For training, do not use only the variables that correlate best with each other, but all of them.

Otherwise, the result of the prediction would be significantly worse. Maybe this is already an indication of **overfitting** of the ML model.

```python
[38]: # DO NOT USE THIS!!
      X_train, X_test, y_train, y_test = train_test_split(X[['sepal_length',
                                                             'sepal_width']],
                                                         y, test_size = 0.20)
```

8.2 Standardize feature values

Standardize the feature values by computing the **mean**, subtracting the mean from the data points, and then dividing by the **standard deviation**:

```python
[17]: scaler = StandardScaler()
      X_train = scaler.fit_transform(X_train)
      X_test = scaler.transform(X_test)

      #X_train
```

8.3 Train the SVC

In this step the SVC is trained with the training data. Training means to **fit** the SVC to the **training data**.

```python
[18]: # fit the model for the data
      classifier.fit(X_train, y_train)
```

```python
[18]: SVC()
```

8.4 Make predictions

In this step the aim is to **predict the species** using unlabeled test data.

```
[19]: y_pred = classifier.predict(X_test)
      #X_test
```

9 STEP 6: Evaluate model's performance

Subsequently to the training of the SVC model and the classification predictions made based on the test data, this step evaluates the **quality of the classification result** using known **metrics** such as the **accuracy score** as well as the **confusion matrix**.

9.1 Accuracy Score

In a multilabel classification (such as the Iris dataset), this **Accuracy classification score** computes the subset accuracy. For further details see sklearn.metrics.accuracy_score.

```
[20]: from sklearn.metrics import accuracy_score

      acc_score = accuracy_score(y_test, y_pred)

      print("Accuracy score: {:.2f} %".format(acc_score.mean()*100))
```

Accuracy score: 96.67 %

9.2 Classification Report

The classification report shows a representation of the main **classification metrics on a per-class basis**. This gives a deeper intuition of the classifier behavior over global accuracy which can mask functional weaknesses in one class of a multiclass problem (see Classification Report).

```
[21]: from sklearn.metrics import classification_report

      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	6
Iris-versicolor	1.00	0.90	0.95	10
Iris-virginica	0.93	1.00	0.97	14
accuracy			0.97	30
macro avg	0.98	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

9.3 Cross-validation score

The function `cross_val_score()` from the Scikit-learn package **trains and tests a model over multiple folds** of your dataset. This cross validation method gives a better **understanding of model performance** over the whole dataset instead of just a single train/test split (see Using cross_val_score in sklearn, simply explained).

```
[22]: from sklearn.model_selection import cross_val_score
```

```
accuracies = cross_val_score(estimator = classifier, X = X_train,
                             y = y_train, cv = 10)

print("Cross-validation score: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

```
Cross-validation score: 95.00 %
Standard Deviation: 5.53 %
```

9.4 Confusion matrix

The **confusion matrix** measures the quality of predictions from a classification model by looking at how many **predictions** are **True** and how many are **False** (see What the Confusion Matrix Measures?.

9.4.1 Textual confusion matrix

For checking the accuracy of the model, the **confusion matrix** can be used for the **cross validation**.

By using the function `sklearn.metrics.confusion_matrix()` a confusion matrix of the true iris class labels versus the predicted class labels is plotted.

[25]:
```
cm = metrics.confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[ 6  0  0]
 [ 0  9  1]
 [ 0  0 14]]
```

9.4.2 Colored confusion matrix

The function `sklearn.metrics.ConfusionMatrixDisplay()` plots a colored confusion matrix.

[26]:
```
sns.set_style("white")

# print colored confusion matrix
cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

cm_colored.figure_.suptitle("Colored Confusion Matrix")
cm_colored.figure_.set_figwidth(8)
cm_colored.figure_.set_figheight(7)

cm_colored.confusion_matrix

# save figure as PNG
plt.tight_layout()
plt.savefig('images/confusion_matrix.png', dpi=150, pad_inches=5)
plt.show()
```

10 STEP 7: Vary parameters of the ML model manually

This section was inspired by In Depth: Parameter tuning for SVC

In this section, the 4 SVC parameters `kernel`, `gamma`, `C` and `degree` will be introduced one by one. Furthermore, their influence on the classification result by varying these single parameters will be shown.

Disclaimer: In order to show the effects of varying the individual parameters in 2D graphs, only the best correlating variables `petal_length` and `petal_width` are used to train the SVC.

Colored Confusion Matrix

Figure 13: Checking the accuracy of the model by using the confusion matrix for cross-validation

10.1 Prepare dataset

```
[142]:  from sklearn.svm import SVC
        from sklearn.model_selection import train_test_split
        from sklearn.model_selection import cross_val_score
        import numpy as np

        # import iris dataset again
        irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')

        # encode the class column from class strings to integer equivalents
        irisdata_df_enc = irisdata_df.replace({"species":  {"Iris-setosa":0,
                                                             "Iris-versicolor":1,
                                                             "Iris-virginica":2}})

        #irisdata_df_enc
```

10.1.1 Prepare datasets for parameter variation and plotting

These datasets will be used for parameter variation and plotting only. In particular, for later **2D plotting** of the effects of parameter variation, only **2 variables** of the iris dataset can be used.

However, as seen in the previous section, this selection is very much at the expense of detection accuracy. Therefore, it is not useful to make predictions with this subset of data - it is not necessary to divide it into a training and a test dataset.

```
[143]:  # copy only 2 feature columns
        # and convert pandas dataframe to numpy array
        X_plot = irisdata_df_enc[['petal_length', 'petal_width']].to_numpy(copy=True)
        #X_plot = irisdata_df_enc[['sepal_length', 'sepal_width']].to_numpy(copy=True)
        #X_plot
```

```
[144]:  # convert pandas dataframe to numpy array
        # and get a flat 1D copy of 2D numpy array
        y_plot = irisdata_df_enc[['species']].to_numpy(copy=True).flatten()
        #y_plot
```

10.1.2 Prepare dataset for prediction and evaluation

To **evaluate the recognition accuracy** by parameter variation, the complete iris dataset with all variables must be used. To make predictions with test data, the dataset is again divided into a training and a test dataset.

```
[145]: X = irisdata_df.drop('species', axis=1)
       y = irisdata_df['species']

       X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

10.2 Plotting functions

This function helps to visualize the modifications by varying the individual SVC parameters:

```
[146]: def plotSVC(title, svc, X, y, xlabel, ylabel):
           # create a mesh to plot in
           x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
           y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

           # prevent division by zero
           if x_min == 0.0:
               x_min = 0.1

           h = (x_max / x_min)/1000
           xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

           plt.subplot(1, 1, 1)
           Z = svc.predict(np.c_[xx.ravel(), yy.ravel()])
           Z = Z.reshape(xx.shape)

           plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
           plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
           plt.xlabel(xlabel)
           plt.ylabel(ylabel)
           plt.xlim(xx.min(), xx.max())
           plt.title(title)
           plt.show()
```

This function cares for cross validation:

```
[168]: def crossValSVC(X_train, y_train, kernel='rbf', gamma='scale', C=1.0, degree=3):
           # train the SVC
           svc = svm.SVC(kernel=kernel,
                         gamma=gamma,
                         C=C,
                         degree=degree).fit(X_train, y_train)
           # calculate accuracies
           accuracies = cross_val_score(estimator = svc, X = X_train,
                                        y = y_train, cv = 10)

           accuracy = accuracies.mean()*100
           return accuracy
```

This function plots the variation of the SVC parameters against the prediction accuracy to show the effect of variation and its limits regarding the phenomenon **overfitting**:

```
[182]: def plotParamsAcc(param_list, acc_list, param_name, log_scale=False):
           fig, ax = plt.subplots(figsize=(10,6))
           title_str = 'Variation of {} parameter '.format(param_name) \
                       +'and its effect to prediction accuracy'
           plt.title(title_str)
           ax.plot(param_list, accuracy_list)
           if log_scale:
               # set the X axis scale to logarithmic
               ax.set_xscale('log')
           plt.xlabel(param_name)
           plt.ylabel('accuracy [%]')
           plt.grid()
           plt.show()
```

10.3 Vary `kernel` of SVC

The `kernel` parameter selects the type of hyperplane that is used to separate the data. Using `linear` (linear classifier) kernel will use a linear hyperplane (a line in the case of 2D data). The `rbf` (radial basis function kernel) and `poly` (polynomial kernel) kernel use non linear hyperplanes. The **default** is `kernel=rbf`.

```
[149]: kernels = ['linear', 'rbf', 'poly', 'sigmoid']

       xlabel = 'Petal length'
       ylabel = 'Petal width'

       for kernel in kernels:
           svc_plot = svm.SVC(kernel=kernel).fit(X_plot, y_plot)
           accuracy = crossValSVC(X_train, y_train, kernel=kernel)
           title_str = 'kernel: \''+str(kernel)+'\', '+'Acc. prediction: {:.2f}%'.
        ↪format(accuracy)
           plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```


Figure 14: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')

10.4 Vary `gamma` parameter

The `gamma` parameter is used for **non linear hyperplanes**. The higher the `gamma` float value it tries to exactly fit the training dataset. The **default** is `gamma='scale'`.

Figure 15: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')

Figure 16: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')

Figure 17: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')

```
[175]: gammas = [0.1, 1, 10, 100, 200]

xlabel = 'Petal length'
```

```
ylabel = 'Petal width'

for gamma in gammas:
    svc_plot = svm.SVC(kernel='rbf', gamma=gamma).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', gamma=gamma)
    title_str = 'gamma: \''+str(gamma)+'\', ' \
                +'Acc. prediction: {:.2f}%'.format(accuracy)
    plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```


Figure 18: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

Figure 19: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

Show the variation of the SVC parameter `gamma` against the **prediction accuracy**.

As we can see, increasing `gamma` leads to **overfitting** as the classifier tries to perfectly fit the training data.

[184]:
```
gammas = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1, 10, 100, 200]

accuracy_list = list()
for gamma in gammas:
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', gamma=gamma)
    accuracy_list.append(accuracy)

plotParamsAcc(gammas, accuracy_list, 'gamma', log_scale=True)
```

Figure 20: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

Figure 21: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel
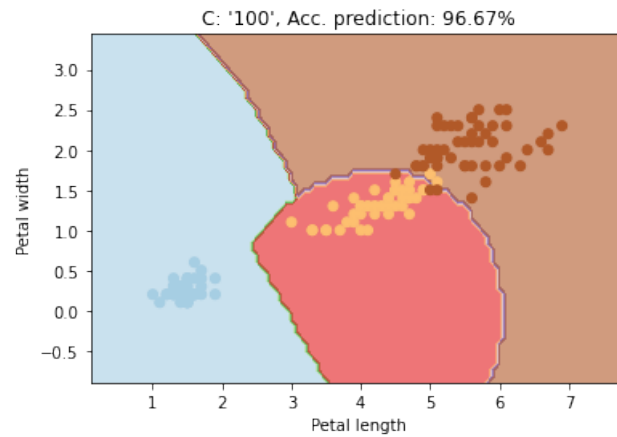

Figure 22: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

10.5 Vary `C` parameter

The `C` parameter is the **penalty** of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly. The **default** is `C=1.0`.

Figure 23: The plot shows the variation of the SVC parameter 'gamma' against the prediction accuracy

[172]:
```
cs = [0.1, 1, 5, 10, 100, 1000, 10000]

xlabel = 'Petal length'
ylabel = 'Petal width'

for c in cs:
    svc_plot = svm.SVC(kernel='rbf', C=c).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', C=c)
    title_str = 'C: \''+str(c)+'\', ' \
                +'Acc. prediction: {:.2f}%'.format(accuracy)
    plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```


Figure 24: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Show the variation of the SVC parameter `C` against the **prediction accuracy**.

But be careful: to high `C` values may lead to **overfitting** the training data.

Figure 25: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Figure 26: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Figure 27: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

```
[185]: cs = [0.1, 1, 5, 6, 7, 8, 10, 100, 1000, 10000]

accuracy_list = list()
```

Figure 28: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Figure 29: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Figure 30: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

```
for c in cs:
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', C=c)
    accuracy_list.append(accuracy)
```

```
plotParamsAcc(cs, accuracy_list, 'C', log_scale=True)
```
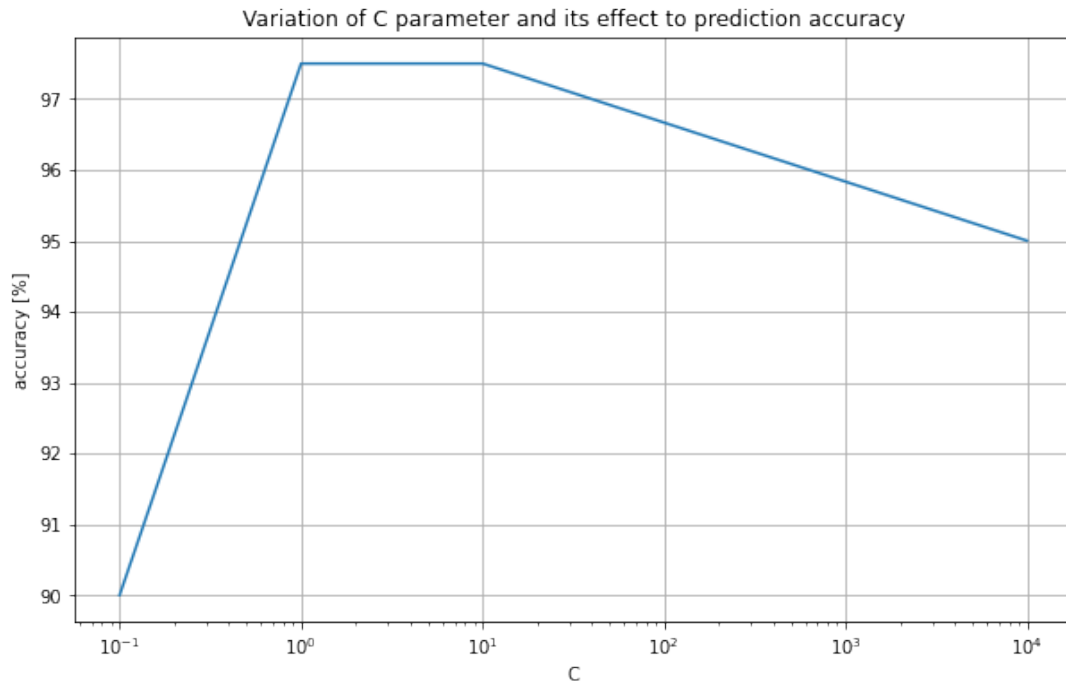

Figure 31: The plot shows the variation of the SVC parameter 'C' against the prediction accuracy

10.6 Vary `degree` parameter

The `degree` parameter is used when the `kernel` is set to `poly` and is ignored by all other kernels. It's basically the **degree of the polynomial** used to find the hyperplane to split the data. The **default** is degree=3.

Using `degree = 1` is the same as using a `linear` kernel. Also, increasing this parameters leads to **higher training times**.

```
[173]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

       xlabel = 'Petal length'
       ylabel = 'Petal width'

       for degree in degrees:
           svc_plot = svm.SVC(kernel='poly', degree=degree).fit(X_plot, y_plot)
           accuracy = crossValSVC(X_train, y_train, kernel='poly', degree=degree)
           title_str = 'degree: \''+str(degree)+'\', ' \
                       +'Acc. prediction: {:.2f}%'.format(accuracy)
           plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```

Show the variation of the SVC parameter `degree` against the **prediction accuracy**.

As we can see, increasing the `degree` of the polynomial hyperplane leads to **overfitting** the training data.

```
[186]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

       accuracy_list = list()
       for degree in degrees:
```

Figure 32: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 33: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 34: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel

```
accuracy = crossValSVC(X_train, y_train, kernel='poly', degree=degree)
accuracy_list.append(accuracy)
```

Figure 35: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 36: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 37: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel

```
plotParamsAcc(degrees, accuracy_list, 'degree', log_scale=False)
```

Figure 38: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 39: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 40: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel

11 STEP 8: Tune the ML model systematically

In the final step, two approaches to systematic hyper-parameter search are presented: **Grid Search** and **Randomized Search**. While the former exhaustively considers all parameter combinations for

Figure 41: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel


Figure 42: The plot shows the variation of the SVC parameter 'degree' against the prediction accuracy

given values, the latter selects a number of candidates from a parameter space with a particular random distribution.

Sources:

- 3.2. Tuning the hyper-parameters of an estimator
    - sklearn.model_selection.GridSearchCV
    - sklearn.model_selection.RandomizedSearchCV
- Introduction to hyperparameter tuning with scikit-learn and Python
    - Abalone Dataset
- Hyperparameter tuning using Grid Search and Random Search: A Conceptual Guide

Import the necessary packages:

```
[1]: # general packages
     from sklearn.preprocessing import StandardScaler
```

```python
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
import pandas as pd

# additional packages for grid search
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import GridSearchCV

# additional packages for randomized search
from sklearn.model_selection import RandomizedSearchCV
from sklearn.model_selection import RepeatedKFold

# import class MeasExecTimeOfProgram from python file MeasExecTimeOfProgramclass.py
from MeasExecTimeOfProgram_class import MeasExecTimeOfProgram
```

Set path and columns of the Iris dataset for import:

```python
[2]:  # specify the path of the dataset
      CSV_PATH = "./datasets/IRIS_flower_dataset_kaggle.csv"
```

Load dataset and split it into subsets for training and testing in the ratio 80% to 20%:

```python
[31]:  # load the dataset, separate the features and labels, and perform a
       # training and testing split using 80% of the data for training and
       # 20% for evaluation
       irisdata_df = pd.read_csv(CSV_PATH)

       X = irisdata_df.drop('species', axis=1)
       y = irisdata_df['species']

       X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

Standardize the feature values by computing the **mean**, subtracting the mean from the data points, and then dividing by the **standard deviation**:

```python
[32]:  scaler = StandardScaler()
       X_train = scaler.fit_transform(X_train)
       X_test = scaler.transform(X_test)

       #X_train
```

11.1 Finding a baseline

The aim of this sub-step is to establish a baseline on the Iris dataset by training a **Support Vector Classifier (SVC)** with no hyperparameter tuning.

Train the model with **no tuning of hyperparameters** to find the baseline for later improvements:

```python
[33]:  classifier = SVC(kernel = 'linear', random_state = 0)

       # initiate measuring execution time
       execTime = MeasExecTimeOfProgram()
       execTime.start()

       classifier.fit(X_train, y_train)

       # print time delta
       print('Execution time: {:.2f} s'.format(execTime.stop()/1000))
```

Execution time: 0.00 s

Evaluate our model using accuracy score:

```python
[34]: from sklearn.metrics import accuracy_score

y_pred = classifier.predict(X_test)

acc_score = accuracy_score(y_test, y_pred)

print("Accuracy score: {:.2f} %".format(acc_score.mean()*100))
```

Accuracy score: 96.67 %

```python
[35]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	9
Iris-versicolor	1.00	0.90	0.95	10
Iris-virginica	0.92	1.00	0.96	11
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

```python
[36]: sns.set_style("white")

# print colored confusion matrix
cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

cm_colored.figure_.suptitle("Colored Confusion Matrix")
cm_colored.figure_.set_figwidth(8)
cm_colored.figure_.set_figheight(7)

cm_colored.confusion_matrix

plt.tight_layout()
plt.show()
```

```python
[ ]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

```python
[11]: classifier.get_params()
```

```python
[11]: {'C': 1.0,
      'break_ties': False,
      'cache_size': 200,
      'class_weight': None,
      'coef0': 0.0,
      'decision_function_shape': 'ovr',
      'degree': 3,
      'gamma': 'scale',
      'kernel': 'linear',
```

Figure 43:

```
'max_iter': -1,
'probability': False,
'random_state': 0,
'shrinking': True,
'tol': 0.001,
'verbose': False}
```

11.2 Grid Search

Initialize the SVC model and define the **space of the hyperparameters** to perform the **grid-search** over:

```
[9]: classifier = SVC()

     kernels = ["linear", "rbf", "sigmoid", "poly"]
     gammas = [0.1, 1, 10, 100, 200]
     cs = [0.1, 1, 5, 10, 100, 1000, 10000]
     degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

     grid = dict(kernel=kernels, gamma=gammas, C=cs, degree=degrees)
```

Initialize a **cross-validation fold** and **perform a grid-search** to tune the hyperparameters:

```
[10]: cvFold = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

      gridSearch = GridSearchCV(estimator=classifier, param_grid=grid, n_jobs=-1,
                                cv=cvFold, scoring="accuracy")

      # initiate measuring execution time
      execTime = MeasExecTimeOfProgram()
      execTime.start()

      searchResults = gridSearch.fit(X_train, y_train)

      # print time delta
```

```
print('Execution time: {:.2f} s'.format(execTime.stop()/1000))
```

```
---------------------------------------------------------------------------
KeyboardInterrupt                         Traceback (most recent call last)
Input In [10], in <cell line: 10>()
      7 execTime = MeasExecTimeOfProgram()
      8 execTime.start()
---> 10 searchResults = gridSearch.fit(X_train, y_train)
     12 # print time delta
     13 print('Execution time: {:.2f} s'.format(execTime.stop()/1000))

File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/sklearn/model_selection/
 ↪_search.py:875, in BaseSearchCV.fit(self, X, y, groups, **fit_params)
    869     results = self._format_results(
    870         all_candidate_params, n_splits, all_out, all_more_results
    871     )
    873     return results
--> 875 self._run_search(evaluate_candidates)
    877 # multimetric is determined here because in the case of a callable
    878 # self.scoring the return type is only known after calling
    879 first_test_score = all_out[0]["test_scores"]

File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/sklearn/model_selection/
 ↪_search.py:1375, in GridSearchCV._run_search(self, evaluate_candidates)
   1373 def _run_search(self, evaluate_candidates):
   1374     """Search all candidates in param_grid"""
-> 1375     evaluate_candidates(ParameterGrid(self.param_grid))

File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/sklearn/model_selection/
 ↪_search.py:822, in BaseSearchCV.fit.<locals>.evaluate_candidates(candidate_params
 ↪cv, more_results)
    814 if self.verbose > 0:
    815     print(
    816         "Fitting {0} folds for each of {1} candidates,"
    817         " totalling {2} fits".format(
    818             n_splits, n_candidates, n_candidates * n_splits
    819         )
    820     )
--> 822 out = parallel(
    823     delayed(_fit_and_score)(
    824         clone(base_estimator),
    825         X,
    826         y,
    827         train=train,
    828         test=test,
    829         parameters=parameters,
    830         split_progress=(split_idx, n_splits),
    831         candidate_progress=(cand_idx, n_candidates),
    832         **fit_and_score_kwargs,
    833     )
    834     for (cand_idx, parameters), (split_idx, (train, test)) in product(
    835         enumerate(candidate_params), enumerate(cv.split(X, y, groups))
    836     )
    837 )
    839 if len(out) < 1:
    840     raise ValueError(
    841         "No fits were performed. "
    842         "Was the CV iterator empty? "
```

```
    843          "Were there no candidates?"
    844      )

 File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/joblib/parallel.py:1056, in
  ↪Parallel.__call__(self, iterable)
    1053      self._iterating = False
    1055 with self._backend.retrieval_context():
 -> 1056      self.retrieve()
    1057 # Make sure that we get a last message telling us we are done
    1058 elapsed_time = time.time() - self._start_time

 File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/joblib/parallel.py:935, in
  ↪Parallel.retrieve(self)
    933 try:
    934      if getattr(self._backend, 'supports_timeout', False):
 --> 935          self._output.extend(job.get(timeout=self.timeout))
    936      else:
    937          self._output.extend(job.get())

 File ~/miniconda3/envs/ipub/lib/python3.9/site-packages/joblib/_parallel_backends.py
  ↪542, in LokyBackend.wrap_future_result(future, timeout)
    539 """Wrapper for Future.result to implement the same behaviour as
    540 AsyncResults.get from multiprocessing."""
    541 try:
 --> 542      return future.result(timeout=timeout)
    543 except CfTimeoutError as e:
    544      raise TimeoutError from e

 File ~/miniconda3/envs/ipub/lib/python3.9/concurrent/futures/_base.py:441, in Future
  ↪result(self, timeout)
    438 elif self._state == FINISHED:
    439      return self.__get_result()
 --> 441 self._condition.wait(timeout)
    443 if self._state in [CANCELLED, CANCELLED_AND_NOTIFIED]:
    444      raise CancelledError()

 File ~/miniconda3/envs/ipub/lib/python3.9/threading.py:312, in Condition.wait(self,
  ↪timeout)
    310 try:    # restore state no matter what (e.g., KeyboardInterrupt)
    311      if timeout is None:
 --> 312          waiter.acquire()
    313          gotit = True
    314      else:

KeyboardInterrupt:
```

Extract the best model and evaluate it:

```python
from sklearn.metrics import accuracy_score

bestModel = searchResults.best_estimator_

y_pred = bestModel.predict(X_test)

acc_score = accuracy_score(y_test, y_pred)

print("Accuracy score: {:.2f} %".format(acc_score.mean()*100))
```

```
[ ]: from sklearn.metrics import classification_report

     print(classification_report(y_test, y_pred))
```

```
[ ]: bestModel.get_params()
```

11.3 Randomized Search

Initialize the SVC model and define the **space of the hyperparameters** to perform the **randomized-search** over:

```
[37]: classifier = SVC()

      kernels = ["linear", "rbf", "sigmoid", "poly"]
      gammas = [0.1, 1, 10, 100, 200]
      cs = [0.1, 1, 5, 10, 100, 1000, 10000]
      degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

      grid = dict(kernel=kernels, gamma=gammas, C=cs, degree=degrees)
```

Initialize a **cross-validation fold** and **perform a randomized-search** to tune the hyperparameters:

```
[38]: cvFold = RepeatedKFold(n_splits=10, n_repeats=3, random_state=1)

      randomSearch = RandomizedSearchCV(estimator=classifier, n_jobs=-1,
                                        cv=cvFold, param_distributions=grid,
                                        scoring="accuracy")

      # initiate measuring execution time
      execTime = MeasExecTimeOfProgram()
      execTime.start()

      searchResults = randomSearch.fit(X_train, y_train)

      # print time delta
      print('Execution time: {:.2f} s'.format(execTime.stop()/1000))
```

Execution time: 1.35 s

Extract the best model and evaluate it:

```
[39]: from sklearn.metrics import accuracy_score

      bestModel = searchResults.best_estimator_

      y_pred = bestModel.predict(X_test)

      acc_score = accuracy_score(y_test, y_pred)

      print("Accuracy score: {:.2f} %".format(acc_score.mean()*100))
```

Accuracy score: 96.67 %

```
[40]: from sklearn.metrics import classification_report

      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	9
Iris-versicolor	1.00	0.90	0.95	10
Iris-virginica	0.92	1.00	0.96	11
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

```python
[41]: sns.set_style("white")

      # print colored confusion matrix
      cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

      cm_colored.figure_.suptitle("Colored Confusion Matrix")
      cm_colored.figure_.set_figwidth(8)
      cm_colored.figure_.set_figheight(7)

      cm_colored.confusion_matrix

      plt.tight_layout()
      plt.show()
```


Figure 44:

```python
[42]: bestModel.get_params()
```

```python
[42]: {'C': 0.1,
       'break_ties': False,
       'cache_size': 200,
       'class_weight': None,
       'coef0': 0.0,
       'decision_function_shape': 'ovr',
       'degree': 1,
       'gamma': 200,
       'kernel': 'poly',
```

```
'max_iter': -1,
'probability': False,
'random_state': None,
'shrinking': True,
'tol': 0.001,
'verbose': False}
```

12 Summary and conclusions

[ ]: