

SVM_Iris_parameter_tuning

August 2, 2022

1 Introduction

This notebook was basically inspired by:

- [In Depth: Parameter tuning for SVC](#)
- [SVM Hyperparameter Tuning using GridSearchCV](#):

The goal of this notebook is to show the basic steps in machine learning and the influence of choosing the “right” the kernel of a **support vector classifier (SVC)**. Furthermore, the SVC parameters are described and their effect on the classification result is shown.

Following steps will be shown in next **chapters**:

- STEP 0: Get the data
- STEP 1: Exploring the data
- STEP 2: Prepare the data
- STEP 3: Classify by support vector classifier - SVC
- STEP 4: Evaluate the results - metrics
- STEP 5: Vary parameters

2 Load globally used libraries and set plot parameters

```
[40]: import time

from IPython.display import HTML

import pandas as pd
import matplotlib.pyplot as plt
from sklearn import svm, metrics
import seaborn as sns
%matplotlib inline
```

3 STEP 0: Get the data

Since this is intended to be an introduction to the world of machine learning (ML), this step does NOT deal with the design of an application suitable for ML and the acquisition of valid measurement data.

In order to get to know the typical work steps and ML tools, the use of **well-known and well-researched data sets** is clearly **recommended**.

In the further course, the famous [Iris flower data sets](#) will be used. It can be downloaded on [Iris Flower Dataset | Kaggle](#). Furthermore, the dataset is included in Python in the machine learning package [Scikit-learn](#), so that users can access it without having to find a special source for it.

```
[1]: # import some data to play with
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

4 STEP 1: Exploring the data

4.1 Goals of exploration

The objectives of the exploration of the dataset are as follows:

1. Clarify the **origins history**:
 - Where did the data come from? => Contact persons and licensing permissions?
 - Who obtained the data and with which (measurement) methods? => Did systematic errors occur during the acquisition?
 - What were they originally intended for? => Can they be used for my application?
2. Overview of the internal **structure and organisation** of the data:
 - Which columns are there? => With which methods can they be read in (e.g. import of CSV files)?
 - What do they contain for (physical) measured variables? => Which technical or physical correlations exist?
 - Which data formats or types are there? => Do they have to be converted?
 - In which value ranges do the measurement data vary? => Are normalizations necessary?
3. Identify **anomalies** in the data sets:
 - Do the data have **gaps** or **duplicates**? => Does the data set needs to be cleaned?
 - Are there obvious erroneous entries or measurement outliers? => Does (statistical) filtering have to be carried out?
4. Avoidance of **tendencies due to bias**:
 - Are all possible classes included in the dataset and equally distributed? => Does the data set need to be enriched with additional data for balance?
5. Find a first rough **idea of which correlations** could be in the data set

4.2 Clarify the origins history

The **Iris flower data sets** is a multivariate data set introduced by the British statistician and biologist *Ronald Fisher* in his paper “The use of multiple measurements in taxonomic problems as an example of linear discriminant analysis” (1936). It is sometimes called *Anderson’s Iris data set* because Edgar Anderson collected the data to quantify the morphologic variation of Iris flowers of three related species (source: [Iris flower data set](#)).

The dataset is published in Public Domain with a [CC0-License](#).

This dataset became a typical test case for many statistical classification techniques in machine learning such as **support vector machines**.

[..] measurements of the flowers of fifty plants each of the two species *Iris setosa* and *I. versicolor*, found **growing together in the same colony** and measured by Dr E. Anderson [..] (source: R. A. Fisher (1936). “The use of multiple measurements in taxonomic problems”. [Annals of Eugenics](#))

[..] *Iris virginica*, differs from the two other samples in **not being taken from the same natural colony** [..] (source: ibidem)

4.3 Overview of the internal structure and organisation of the data

The data set consists of 50 samples from each of three species of Iris (*Iris setosa*, *Iris virginica* and *Iris versicolor*), so there are 150 total samples. Four features were measured from each sample: the length and the width of the [sepals](#) and [petals](#), in centimetres.

Here is a principle illustration of a flower with sepal and petal:

```
[21]: display(HTML("<figure><img src='../images/Mature_flower_diagram.svg' \
    ↪width='800px'> \
    ↪    <figcaption>Principle illustration of a flower with sepal and \
    ↪petal (source: <a href='https://en.wikipedia.org/wiki/File: \
    ↪Mature_flower_diagram.svg'>Mature_flower_diagram.svg</a></figcaption> \
    ↪    </figure>"))
```

<IPython.core.display.HTML object>

Here are pictures of the three different Iris species (*Iris setosa*, *Iris virginica* and *Iris versicolor*). Given the dimensions of the flower, it will be possible to predict the class of the flower.

```
[22]: display(HTML("<table> \
    ↪    <tr> \
    ↪        <td><figure><img src='../images/Iris_setosa_640px.jpg' \
    ↪width='320px'> \
    ↪            <figcaption><i>Iris setosa</i> (source: <a href='https:/ \
    ↪commons.wikimedia.org/wiki/File:Irissetosa1.jpg'>Irissetosa1.jpg</a></ \
    ↪figcaption> \
    ↪        </figure></td> \
    ↪        <td><figure><img src='../images/Iris_versicolor_640px.jpg' \
    ↪width='320px'> \
    ↪            <figcaption><i>Iris versicolor</i> (source: <a \
    ↪href='https://en.wikipedia.org/wiki/File:Iris_versicolor_3.jpg'>Iris \
    ↪versicolor 3.jpg</a></figcaption> \
    ↪        </figure></td> \
    ↪        <td><figure><img src='../images/Iris_virginica_590px.jpg' \
    ↪width='295px'> \
    ↪            <figcaption><i>Iris virginica</i> (source: <a \
    ↪href='https://en.wikipedia.org/wiki/File:Iris_virginica.jpg'>Iris virginica. \
    ↪jpg</a></figcaption> \
    ↪        </figure></td> \
    ↪    </tr> \
    ↪    </table>"))
```

<IPython.core.display.HTML object>

4.3.1 Inspect structure of dataframe

Print first or last 5 rows of dataframe:

```
[3]: irisdata_df.head()
```

```
[3]:   sepal_length  sepal_width  petal_length  petal_width  species
0         5.1         3.5         1.4         0.2  Iris-setosa
1         4.9         3.0         1.4         0.2  Iris-setosa
2         4.7         3.2         1.3         0.2  Iris-setosa
3         4.6         3.1         1.5         0.2  Iris-setosa
4         5.0         3.6         1.4         0.2  Iris-setosa
```

```
[4]: irisdata_df.tail()
```

```
[4]:   sepal_length  sepal_width  petal_length  petal_width  species
145         6.7         3.0         5.2         2.3  Iris-virginica
146         6.3         2.5         5.0         1.9  Iris-virginica
147         6.5         3.0         5.2         2.0  Iris-virginica
148         6.2         3.4         5.4         2.3  Iris-virginica
149         5.9         3.0         5.1         1.8  Iris-virginica
```

While printing a dataframe - only an abbreviated view of the dataframe is shown :(
Default setting in the pandas library makes it to display only 5 lines from head and from tail.

```
[6]: irisdata_df
```

```
[6]:   sepal_length  sepal_width  petal_length  petal_width  species
0         5.1         3.5         1.4         0.2  Iris-setosa
1         4.9         3.0         1.4         0.2  Iris-setosa
2         4.7         3.2         1.3         0.2  Iris-setosa
3         4.6         3.1         1.5         0.2  Iris-setosa
4         5.0         3.6         1.4         0.2  Iris-setosa
..         ...         ...         ...         ...         ...
145         6.7         3.0         5.2         2.3  Iris-virginica
146         6.3         2.5         5.0         1.9  Iris-virginica
147         6.5         3.0         5.2         2.0  Iris-virginica
148         6.2         3.4         5.4         2.3  Iris-virginica
149         5.9         3.0         5.1         1.8  Iris-virginica
```

[150 rows x 5 columns]

To print all rows of a dataframe, the option `display.max_rows` has to set to `None` in pandas:

```
[7]: pd.set_option('display.max_rows', None)
irisdata_df
```

```

[7]:      sepal_length  sepal_width  petal_length  petal_width      species
0          5.1          3.5          1.4          0.2      Iris-setosa
1          4.9          3.0          1.4          0.2      Iris-setosa
2          4.7          3.2          1.3          0.2      Iris-setosa
3          4.6          3.1          1.5          0.2      Iris-setosa
4          5.0          3.6          1.4          0.2      Iris-setosa
5          5.4          3.9          1.7          0.4      Iris-setosa
6          4.6          3.4          1.4          0.3      Iris-setosa
7          5.0          3.4          1.5          0.2      Iris-setosa
8          4.4          2.9          1.4          0.2      Iris-setosa
9          4.9          3.1          1.5          0.1      Iris-setosa
10         5.4          3.7          1.5          0.2      Iris-setosa
11         4.8          3.4          1.6          0.2      Iris-setosa
12         4.8          3.0          1.4          0.1      Iris-setosa
13         4.3          3.0          1.1          0.1      Iris-setosa
14         5.8          4.0          1.2          0.2      Iris-setosa
15         5.7          4.4          1.5          0.4      Iris-setosa
16         5.4          3.9          1.3          0.4      Iris-setosa
17         5.1          3.5          1.4          0.3      Iris-setosa
18         5.7          3.8          1.7          0.3      Iris-setosa
19         5.1          3.8          1.5          0.3      Iris-setosa
20         5.4          3.4          1.7          0.2      Iris-setosa
21         5.1          3.7          1.5          0.4      Iris-setosa
22         4.6          3.6          1.0          0.2      Iris-setosa
23         5.1          3.3          1.7          0.5      Iris-setosa
24         4.8          3.4          1.9          0.2      Iris-setosa
25         5.0          3.0          1.6          0.2      Iris-setosa
26         5.0          3.4          1.6          0.4      Iris-setosa
27         5.2          3.5          1.5          0.2      Iris-setosa
28         5.2          3.4          1.4          0.2      Iris-setosa
29         4.7          3.2          1.6          0.2      Iris-setosa
30         4.8          3.1          1.6          0.2      Iris-setosa
31         5.4          3.4          1.5          0.4      Iris-setosa
32         5.2          4.1          1.5          0.1      Iris-setosa
33         5.5          4.2          1.4          0.2      Iris-setosa
34         4.9          3.1          1.5          0.1      Iris-setosa
35         5.0          3.2          1.2          0.2      Iris-setosa
36         5.5          3.5          1.3          0.2      Iris-setosa
37         4.9          3.1          1.5          0.1      Iris-setosa
38         4.4          3.0          1.3          0.2      Iris-setosa
39         5.1          3.4          1.5          0.2      Iris-setosa
40         5.0          3.5          1.3          0.3      Iris-setosa
41         4.5          2.3          1.3          0.3      Iris-setosa
42         4.4          3.2          1.3          0.2      Iris-setosa
43         5.0          3.5          1.6          0.6      Iris-setosa
44         5.1          3.8          1.9          0.4      Iris-setosa
45         4.8          3.0          1.4          0.3      Iris-setosa

```

46	5.1	3.8	1.6	0.2	Iris-setosa
47	4.6	3.2	1.4	0.2	Iris-setosa
48	5.3	3.7	1.5	0.2	Iris-setosa
49	5.0	3.3	1.4	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
51	6.4	3.2	4.5	1.5	Iris-versicolor
52	6.9	3.1	4.9	1.5	Iris-versicolor
53	5.5	2.3	4.0	1.3	Iris-versicolor
54	6.5	2.8	4.6	1.5	Iris-versicolor
55	5.7	2.8	4.5	1.3	Iris-versicolor
56	6.3	3.3	4.7	1.6	Iris-versicolor
57	4.9	2.4	3.3	1.0	Iris-versicolor
58	6.6	2.9	4.6	1.3	Iris-versicolor
59	5.2	2.7	3.9	1.4	Iris-versicolor
60	5.0	2.0	3.5	1.0	Iris-versicolor
61	5.9	3.0	4.2	1.5	Iris-versicolor
62	6.0	2.2	4.0	1.0	Iris-versicolor
63	6.1	2.9	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
65	6.7	3.1	4.4	1.4	Iris-versicolor
66	5.6	3.0	4.5	1.5	Iris-versicolor
67	5.8	2.7	4.1	1.0	Iris-versicolor
68	6.2	2.2	4.5	1.5	Iris-versicolor
69	5.6	2.5	3.9	1.1	Iris-versicolor
70	5.9	3.2	4.8	1.8	Iris-versicolor
71	6.1	2.8	4.0	1.3	Iris-versicolor
72	6.3	2.5	4.9	1.5	Iris-versicolor
73	6.1	2.8	4.7	1.2	Iris-versicolor
74	6.4	2.9	4.3	1.3	Iris-versicolor
75	6.6	3.0	4.4	1.4	Iris-versicolor
76	6.8	2.8	4.8	1.4	Iris-versicolor
77	6.7	3.0	5.0	1.7	Iris-versicolor
78	6.0	2.9	4.5	1.5	Iris-versicolor
79	5.7	2.6	3.5	1.0	Iris-versicolor
80	5.5	2.4	3.8	1.1	Iris-versicolor
81	5.5	2.4	3.7	1.0	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
83	6.0	2.7	5.1	1.6	Iris-versicolor
84	5.4	3.0	4.5	1.5	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
90	5.5	2.6	4.4	1.2	Iris-versicolor
91	6.1	3.0	4.6	1.4	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor

93	5.0	2.3	3.3	1.0	Iris-versicolor
94	5.6	2.7	4.2	1.3	Iris-versicolor
95	5.7	3.0	4.2	1.2	Iris-versicolor
96	5.7	2.9	4.2	1.3	Iris-versicolor
97	6.2	2.9	4.3	1.3	Iris-versicolor
98	5.1	2.5	3.0	1.1	Iris-versicolor
99	5.7	2.8	4.1	1.3	Iris-versicolor
100	6.3	3.3	6.0	2.5	Iris-virginica
101	5.8	2.7	5.1	1.9	Iris-virginica
102	7.1	3.0	5.9	2.1	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
104	6.5	3.0	5.8	2.2	Iris-virginica
105	7.6	3.0	6.6	2.1	Iris-virginica
106	4.9	2.5	4.5	1.7	Iris-virginica
107	7.3	2.9	6.3	1.8	Iris-virginica
108	6.7	2.5	5.8	1.8	Iris-virginica
109	7.2	3.6	6.1	2.5	Iris-virginica
110	6.5	3.2	5.1	2.0	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica
113	5.7	2.5	5.0	2.0	Iris-virginica
114	5.8	2.8	5.1	2.4	Iris-virginica
115	6.4	3.2	5.3	2.3	Iris-virginica
116	6.5	3.0	5.5	1.8	Iris-virginica
117	7.7	3.8	6.7	2.2	Iris-virginica
118	7.7	2.6	6.9	2.3	Iris-virginica
119	6.0	2.2	5.0	1.5	Iris-virginica
120	6.9	3.2	5.7	2.3	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
122	7.7	2.8	6.7	2.0	Iris-virginica
123	6.3	2.7	4.9	1.8	Iris-virginica
124	6.7	3.3	5.7	2.1	Iris-virginica
125	7.2	3.2	6.0	1.8	Iris-virginica
126	6.2	2.8	4.8	1.8	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica
128	6.4	2.8	5.6	2.1	Iris-virginica
129	7.2	3.0	5.8	1.6	Iris-virginica
130	7.4	2.8	6.1	1.9	Iris-virginica
131	7.9	3.8	6.4	2.0	Iris-virginica
132	6.4	2.8	5.6	2.2	Iris-virginica
133	6.3	2.8	5.1	1.5	Iris-virginica
134	6.1	2.6	5.6	1.4	Iris-virginica
135	7.7	3.0	6.1	2.3	Iris-virginica
136	6.3	3.4	5.6	2.4	Iris-virginica
137	6.4	3.1	5.5	1.8	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
139	6.9	3.1	5.4	2.1	Iris-virginica

140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

4.3.2 Get data types

```
[8]: irisdata_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 5.3+ KB
```

```
[9]: irisdata_df.describe()
```

```
[9]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

4.3.3 Get data ranges with Boxplots

Boxplots can be used to explore the data ranges in the data set. These also provide information about **outliers**.

```
[36]: sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.0})
sns.set_style("whitegrid")
```



```

#sns.set_style("white")

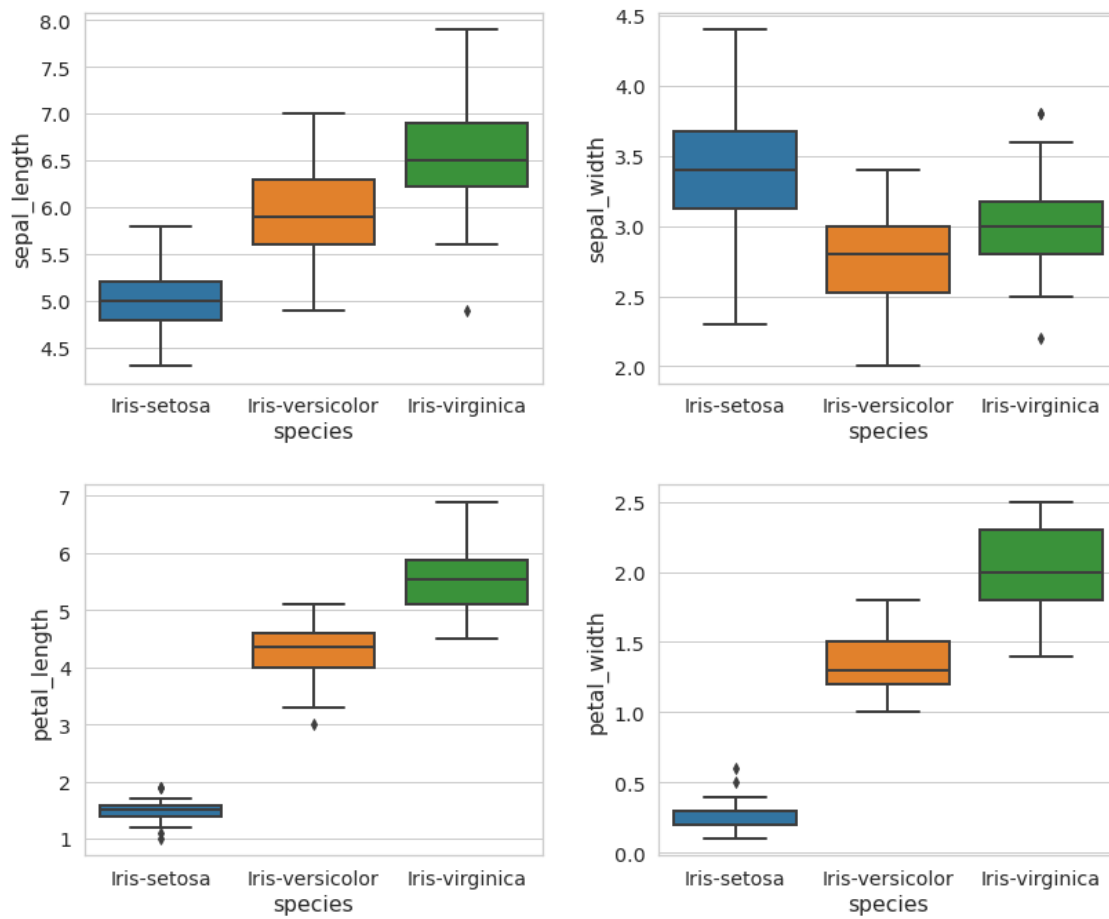
fig, axs = plt.subplots(2, 2, figsize=(12, 10))

fn = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
box1 = sns.boxplot(x = 'species', y = 'sepal_length', data = irisdata_df, order_
    ↪ cn, ax = axs[0,0])
box2 = sns.boxplot(x = 'species', y = 'sepal_width', data = irisdata_df, order_
    ↪ cn, ax = axs[0,1])
box3 = sns.boxplot(x = 'species', y = 'petal_length', data = irisdata_df, order_
    ↪ cn, ax = axs[1,0])
box4 = sns.boxplot(x = 'species', y = 'petal_width', data = irisdata_df, order_
    ↪ cn, ax = axs[1,1])

# add some spacing between subplots
fig.tight_layout(pad=2.0)

plt.show()

```



4.4 Identify anomalies in the data sets

4.4.1 Find gaps in dataset

This section was inspired by [Working with Missing Data in Pandas](#).

Checking for missing values using `isnull()` In order to check for missing values in Pandas DataFrame, we use the function `isnull()`. This function returns a dataframe of Boolean values which are True for NaN values.

```
[37]: pd.set_option('display.max_rows', 40)
      pd.set_option('display.min_rows', 30)
```

```
[38]: irisdata_df.isnull()
```

```
[38]:      sepal_length  sepal_width  petal_length  petal_width  species
0             False           False           False           False  False
1             False           False           False           False  False
2             False           False           False           False  False
3             False           False           False           False  False
4             False           False           False           False  False
5             False           False           False           False  False
6             False           False           False           False  False
7             False           False           False           False  False
8             False           False           False           False  False
9             False           False           False           False  False
10            False           False           False           False  False
11            False           False           False           False  False
12            False           False           False           False  False
13            False           False           False           False  False
14            False           False           False           False  False
..            ...             ...             ...             ...
135           False           False           False           False  False
136           False           False           False           False  False
137           False           False           False           False  False
138           False           False           False           False  False
139           False           False           False           False  False
140           False           False           False           False  False
141           False           False           False           False  False
142           False           False           False           False  False
143           False           False           False           False  False
144           False           False           False           False  False
145           False           False           False           False  False
146           False           False           False           False  False
147           False           False           False           False  False
148           False           False           False           False  False
```

149 False False False False False

[150 rows x 5 columns]

Show only the gaps:

```
[5]: irisdata_df_gaps = irisdata_df[irisdata_df.isnull().any(axis=1)]
irisdata_df_gaps
```

```
[5]: Empty DataFrame
Columns: [sepal_length, sepal_width, petal_length, petal_width, species]
Index: []
```

Fine - this dataset seems to be complete :)

So let's look for something else for exercise: [employees.csv](#)

```
[39]: # import data to dataframe from csv file
employees_df = pd.read_csv("./datasets/employees_edit.csv")

employees_df
```

```
[39]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	

998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management		Team
0	True		Marketing
1	True		NaN
2	False		Finance
3	True		Finance
4	True		Client Services
5	False		Legal
6	True		Product
7	NaN		Finance
8	True		Engineering
9	True	Business Development	
10	True		NaN
11	True		Legal
12	True	Human Resources	
13	False		Sales
14	True		Finance
...
989	False		Legal
990	False		Marketing
991	True		Finance
992	False	Human Resources	
993	False		Legal
994	True	Client Services	
995	True		Marketing
996	True		Finance
997	True		Engineering
998	True		Marketing
999	False	Distribution	
1000	False		Finance
1001	False		Product
1002	False	Business Development	
1003	True		Sales

[1004 rows x 8 columns]

Show only the gaps from this gappy dataset again:

```
[40]: employees_df_gaps = employees_df[employees_df.isnull().any(axis=1)]
employees_df_gaps
```

[40]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
20	Lois	NaN	4/22/1995	7:18 PM	64714	4934.00	
22	Joshua	NaN	3/8/2012	1:58 AM	90816	18816.00	
23	NaN	Male	6/14/2012	4:19 PM	125792	5042.00	
25	NaN	Male	10/8/2012	1:12 AM	37076	18576.00	
27	Scott	NaN	7/11/1991	6:58 PM	122367	5218.00	
31	Joyce	NaN	2/20/2005	2:40 PM	88657	12752.00	
32	NaN	Male	8/21/1998	2:27 PM	122340	6417.00	
39	NaN	Male	1/29/2016	2:33 AM	122173	7797.00	
41	Christine	NaN	6/28/2015	1:08 AM	66582	11308.00	
49	Chris	NaN	1/24/1980	12:13 PM	113590	3055.00	
51	NaN	NaN	12/17/2011	8:29 AM	41126	14009.00	
53	Alan	NaN	3/3/2014	1:28 PM	40341	17578.00	
..		
916	Joe	Male	12/8/1998	10:28 AM	126120	1.02	
927	Irene	NaN	2/28/1991	10:23 PM	135369	4.38	
929	NaN	Female	8/23/2000	4:19 PM	95866	19388.00	
941	Aaron	NaN	1/22/1986	7:39 PM	63126	18424.00	
942	Mark	NaN	9/9/2006	12:27 PM	44836	2657.00	
943	Ralph	NaN	7/28/1995	6:53 PM	70635	2147.00	
949	Gerald	NaN	4/15/1989	12:44 PM	93712	17426.00	
950	NaN	Female	9/15/1985	1:50 AM	133472	16941.00	
951	NaN	Male	7/30/2012	3:07 PM	107351	5329.00	
955	NaN	Female	9/14/2010	5:19 AM	143638	9662.00	
965	Antonio	NaN	6/18/1989	9:37 PM	103050	3.05	
976	Victor	NaN	7/28/2006	2:49 PM	76381	11159.00	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	

	Senior Management	Team
1	True	NaN
7	NaN	Finance
10	True	NaN
20	True	Legal
22	True	Client Services
23	NaN	NaN
25	NaN	Client Services
27	False	Legal
31	False	Product
32	NaN	NaN
39	NaN	Client Services
41	True	Business Development
49	False	Sales

```

51          NaN          Sales
53         True          Finance
..          ...          ...
916        False          NaN
927        False  Business Development
929         NaN          Sales
941        False  Client Services
942        False  Client Services
943        False  Client Services
949         True          Distribution
950         NaN          Distribution
951         NaN          Marketing
955         NaN          NaN
965        False          Legal
976         True          Sales
989        False          Legal
993        False          Legal
999        False  Distribution

```

[237 rows x 8 columns]

Fill the missing values with fillna() Now we are going to fill all the null (NaN) values in Gender column with “No Gender”.

Attention: We are doing that directly in this dataframe with `inplace = True` - we don’t make a deep copy!

```

[41]: # filling a null values using fillna()
employees_df["Gender"].fillna("No Gender", inplace = True)
employees_df

```

```

[41]:
   First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
0    Douglas    Male    8/6/1993         12:42 PM   97308    6945.00
1     Thomas    Male    3/31/1996         6:53 AM    61933         4.17
2      Maria   Female    4/23/1993        11:17 AM  130590   11858.00
3      Jerry    Male    3/4/2005         1:00 PM   138705         9.34
4      Larry    Male    1/24/1998         4:47 PM   101004    1389.00
5     Dennis    Male    4/18/1987         1:35 AM   115163   10125.00
6       Ruby   Female    8/17/1987         4:20 PM    65476   10012.00
7         NaN   Female    7/20/2015        10:43 AM    45906   11598.00
8     Angela   Female   11/22/2005         6:29 AM    95570   18523.00
9    Frances   Female    8/8/2002         6:51 AM   139852    7524.00
10    Louise   Female    8/12/1980         9:01 AM    63241   15132.00
11     Julie   Female   10/26/1997         3:19 PM   102508   12637.00
12   Brandon    Male    12/1/1980         1:08 AM   112807   17492.00
13     Gary     Male    1/27/2008        11:40 PM   109831    5831.00
14  Kimberly   Female    1/14/1999         7:13 AM    41426   14543.00

```

...
989	Stephen	No Gender	7/10/1983		8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982		7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014		5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004		9:34 AM	47638	11209.00
993	Justin	No Gender	2/10/1991		4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987		1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002		5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011		8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997		3:53 PM	56450	19.04
998	George	Male	6/21/2013		5:47 PM	98874	4479.00
999	Henry	No Gender	11/23/2014		6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984		6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product

```

1002                False Business Development
1003                True      Sales

```

```
[1004 rows x 8 columns]
```

Dropping missing values using dropna() In order to drop null values from a dataframe, we use `dropna()` function. This function drops rows or columns of datasets with NaN values in different ways.

Default is to drop rows with at least 1 null value (NaN). Giving the parameter `how = 'all'` the function drops rows with all data missing or contain null values (NaN).

```
[42]: # making a new dataframe with dropped NaN values
employees_df_dropped = employees_df.dropna(axis = 0, how = 'any')
employees_df_dropped
```

```
[42]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
15	Lillian	Female	6/5/2016	6:09 AM	59414	1256.00	
16	Jeremy	Male	9/21/2010	5:56 AM	90370	7369.00	
17	Shawn	Male	12/7/1986	7:45 PM	111737	6414.00	
...	
989	Stephen	No Gender	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	No Gender	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	
998	George	Male	6/21/2013	5:47 PM	98874	4479.00	
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	

1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00
------	--------	------	-----------	---------	--------	----------

	Senior Management	Team
0	True	Marketing
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
8	True	Engineering
9	True	Business Development
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
15	False	Product
16	False	Human Resources
17	False	Product
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[903 rows x 8 columns]

Finally we compare the sizes of dataframes so that we learn how many rows had at least 1 Null value.

```
[43]: print("Old data frame length:", len(employees_df))
      print("New data frame length:", len(employees_df_dropped))
      print("Number of rows with at least 1 NaN value: ",
            ↪(len(employees_df)-len(employees_df_dropped)))
```

```
Old data frame length: 1004
New data frame length: 903
Number of rows with at least 1 NaN value: 101
```

4.4.2 Find and remove duplicates in dataset

This section was inspired by: - [How to Find Duplicates in Pandas DataFrame \(With Examples\)](#) - [How to Drop Duplicate Rows in a Pandas DataFrame](#)

Checking for duplicate values using duplicated() In order to check for duplicate values in Pandas DataFrame, we use a function `duplicated()`. This function can be used in two ways: - find duplicate rows across **all columns** with `duplicateRows = df[df.duplicated()]` - find duplicate rows across **specific columns** `duplicateRows = df[df.duplicated(subset=['col1', 'col2'])]`

Find duplicate rows across **all columns**:

```
[44]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[45]: # find duplicate rows across all columns
duplicateRows = employees_df[employees_df.duplicated()]
duplicateRows
```

```
[45]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
112      Karen    Female  11/30/1999          7:46 AM  102488  17653.0
127      Linda    Female  5/25/2000          5:45 PM  119009  12506.0
296    Brandon      NaN  11/3/1997          8:17 PM  121333  15295.0
580   Nicholas    Male   3/1/2013          9:26 PM  101036   2826.0

      Senior Management      Team
112              True      Product
127              True  Business Development
296              False  Business Development
580              True   Human Resources
```

```
[46]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(keep='last')]
duplicateRows
```

```
[46]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
55      Karen    Female  11/30/1999          7:46 AM  102488  17653.0
92      Linda    Female  5/25/2000          5:45 PM  119009  12506.0
153    Brandon      NaN  11/3/1997          8:17 PM  121333  15295.0
442   Nicholas    Male   3/1/2013          9:26 PM  101036   2826.0

      Senior Management      Team
55              True      Product
92              True  Business Development
153              False  Business Development
442              True   Human Resources
```

Find duplicate rows across **specific columns**:

```
[47]: # identify duplicate rows across 'First Name' and 'Last Login Time' columns
duplicateRows = employees_df[employees_df.duplicated(subset=['First Name',
↳ 'Last Login Time'])]
duplicateRows
```

```
[47]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
112      Karen  Female  11/30/1999      7:46 AM  102488  17653.0
127      Linda  Female  5/25/2000      5:45 PM  119009  12506.0
296    Brandon   NaN  11/3/1997      8:17 PM  121333  15295.0
577      NaN  Female  1/13/2009      1:01 PM  118736   7421.0
580  Nicholas   Male  3/1/2013      9:26 PM  101036   2826.0
632      NaN   NaN  9/2/1988      12:49 PM  147309   1702.0
881      NaN   Male  9/5/1980      7:36 AM  114896  13823.0
929      NaN  Female  8/23/2000      4:19 PM   95866  19388.0
934     Nancy  Female  9/10/2001     11:57 PM   85213   2386.0
973     Linda  Female  2/4/2010      8:49 PM   44486  17308.0
```

	Senior Management	Team
112	True	Product
127	True	Business Development
296	False	Business Development
577	NaN	Client Services
580	True	Human Resources
632	NaN	Distribution
881	NaN	Client Services
929	NaN	Sales
934	True	Marketing
973	True	Engineering

```
[48]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(subset=['First Name',
↳ 'Last Login Time'], keep='last')]
duplicateRows
```

```
[48]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
23      NaN   Male  6/14/2012      4:19 PM  125792   5042.00
37     Linda  Female  10/19/1981      8:49 PM   57427   9557.00
55     Karen  Female  11/30/1999      7:46 AM  102488  17653.00
66     Nancy  Female  12/15/2012     11:57 PM  125250   2672.00
92     Linda  Female  5/25/2000      5:45 PM  119009  12506.00
153   Brandon   NaN  11/3/1997      8:17 PM  121333  15295.00
222      NaN  Female  6/17/1991     12:49 PM   71945    5.56
269      NaN   Male  2/4/2005      1:01 PM   40451  16044.00
442  Nicholas   Male  3/1/2013      9:26 PM  101036   2826.00
778      NaN  Female  6/18/2000      7:36 AM  106428  10867.00
```

	Senior Management	Team
--	-------------------	------

23	NaN	NaN
37	True	Client Services
55	True	Product
66	True	Business Development
92	True	Business Development
153	False	Business Development
222	NaN	Marketing
269	NaN	Distribution
442	True	Human Resources
778	NaN	NaN

Dropping duplicate values using `drop_duplicates()` In order to drop duplicate values from a dataframe, we use `drop_duplicates()` function.

This function can be used in two ways: - remove duplicate rows across **all columns** with `df.drop_duplicates()` - find duplicate rows across **specific columns** `df.drop_duplicates(subset=['col1', 'col2'])`

Attention: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

Remove duplicate rows across **all columns**:

```
[49]: # remove duplicate rows across all columns
employees_df.drop_duplicates(inplace=True)
employees_df
```

```
[49]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	

993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[1000 rows x 8 columns]

Remove duplicate rows across **specific columns**:

```
[50]: # remove duplicate rows across 'First Name' and 'Last Login Time' columns
employees_df.drop_duplicates(subset=['First Name', 'Last Login Time'],
                             keep='last', inplace=True)
employees_df
```

```
[50]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	
998	George	Male	6/21/2013	5:47 PM	98874	4479.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal

6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[994 rows x 8 columns]

4.5 Avoidance of tendencies due to bias

The description of the Iris dataset says, that it consists of **50 samples** from **each of three species** of Iris (Iris setosa, Iris virginica and Iris versicolor), so there are **150 total samples**.

But how to prove it?

4.5.1 Count occurrences of unique values

To prove whether all possible classes included in the dataset and equally distributed, you can use the function `df.value_counts`.

Following parameters can be used for fine tuning: - `dropna=False` causes that NaN values are included - `normalize=True`: relative frequencies of the unique values are returned - `ascending=False`: sort resulting classes descending

```
[51]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[6]: # count unique values without missing values in a column, ordered descending
      ↪and normalized
irisdata_df['species'].value_counts(ascending=False, dropna=False,
      ↪normalize=True)
```

```
[6]: Iris-setosa      0.333333
      Iris-versicolor 0.333333
      Iris-virginica  0.333333
      Name: species, dtype: float64
```

```
[59]: # count unique values and missing values in a column, ordered descending and
      ↪not absolute values
employees_df['Team'].value_counts(ascending=False, dropna=False,
      ↪normalize=False)
```

```
[59]: Client Services      106
      Business Development 103
      Finance             102
      Marketing           98
      Product             96
      Sales               94
      Engineering         92
      Human Resources      92
      Distribution        90
      Legal               88
      NaN                 43
      Name: Team, dtype: int64
```

4.5.2 Display Histogram


This section was inspired by: [Pandas Histogram – DataFrame.hist\(\)](#).

Histograms represent **frequency distributions** graphically. This requires the separation of the data into classes (so-called **bins**).

These classes are represented in the histogram as rectangles of equal or variable width. The height of each rectangle then represents the (relative or absolute) **frequency density**.


```
[37]: employees_df.hist(column=['Salary'])
```

```
[37]: array([[<AxesSubplot:title={'center': 'Salary'}>]], dtype=object)
```

```
[38]: employees_df.hist(column='Salary', by='Gender')
```

```
[38]: array([<AxesSubplot:title={'center':'Female'}>,  
        <AxesSubplot:title={'center':'Male'}>], dtype=object)
```



4.6 First idea of correlations in data set

To get a rough idea of the **dependencies** and **correlations** in the data set, it can be helpful to visualize the whole dataset in a **correlation heatmap**. They show in a glance which variables are correlated, to what degree and in which direction.

Later, 2 particularly well correlated variables are selected from the data set and plotted in a **scatterplot**.

4.6.1 Visualise data with correlation heatmap

This section was inspired by [How to Create a Seaborn Correlation Heatmap in Python?](#).

Correlation matrices are an **essential tool of exploratory data analysis**. Correlation heatmaps contain the same information in a visually appealing way. What more: they show in a glance which variables are correlated, to what degree, in which direction, and alerts us to potential multicollinearity problems (source: ibidem).

Simple correlation matrix Because **string values can never be correlated**, the class names (species) have to be converted first:

```
[91]: # encoding the class column
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":
↪0, "Iris-versicolor":1, "Iris-virginica":2}})
irisdata_df_enc
```

```
[91]:      sepal_length  sepal_width  petal_length  petal_width  species
0              5.1           3.5           1.4           0.2         0
1              4.9           3.0           1.4           0.2         0
2              4.7           3.2           1.3           0.2         0
3              4.6           3.1           1.5           0.2         0
4              5.0           3.6           1.4           0.2         0
..              ...           ...           ...           ...         ...
145             6.7           3.0           5.2           2.3         2
146             6.3           2.5           5.0           1.9         2
147             6.5           3.0           5.2           2.0         2
148             6.2           3.4           5.4           2.3         2
149             5.9           3.0           5.1           1.8         2
```

[150 rows x 5 columns]

```
[92]: irisdata_df_enc.corr()
```

```
[92]:      sepal_length  sepal_width  petal_length  petal_width  species
sepal_length      1.000000   -0.109369    0.871754    0.817954  0.782561
sepal_width      -0.109369    1.000000   -0.420516   -0.356544 -0.419446
petal_length      0.871754   -0.420516    1.000000    0.962757  0.949043
petal_width      0.817954   -0.356544    0.962757    1.000000  0.956464
species           0.782561   -0.419446    0.949043    0.956464  1.000000
```


Correlation heatmap Choose the color sets from [color map](#).

```
[93]: # increase the size of the heatmap
plt.figure(figsize=(16, 6))

# store heatmap object in a variable to easily access it
# when you want to include more features (such as title)
# set the range of values to be displayed on the colormap from -1 to 1,
# and set 'annotation=True' to display the correlation values on the heatmap
heatmap = sns.heatmap(irisdata_df_enc.corr(), vmin=-1, vmax=1, annot=True, ↵
↪cmap='PRGn_r')

# give a title to the heatmap
# 'pad=12' defines the distance of the title from the top of the heatmap
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
```

```
[93]: Text(0.5, 1.0, 'Correlation Heatmap')
```



Triangle correlation heatmap When looking at the correlation heatmaps above, you would not lose any information by **cutting** away half of it **along the diagonal** line marked by 1-s.

The **numpy** function `np.triu()` can be used to isolate the upper triangle of a matrix while turning all the values in the lower triangle into 0.

```
[94]: import numpy as np

np.triu(np.ones_like(irisdata_df_enc.corr()))
```

```
[94]: array([[1., 1., 1., 1., 1.],
            [0., 1., 1., 1., 1.],
            [0., 0., 1., 1., 1.],
            [0., 0., 0., 1., 1.],
            [0., 0., 0., 0., 1.]])
```

Use this mask to cut the heatmap along the diagonal:


```
[95]: plt.figure(figsize=(16, 6))

# define the mask to set the values in the upper triangle to 'True'
mask = np.triu(np.ones_like(irisdata_df_enc.corr()), dtype=bool))

heatmap = sns.heatmap(irisdata_df_enc.corr(), mask=mask, vmin=-1, vmax=1,
    ↪annot=True, cmap='PRGn_r')

heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18},
    ↪pad=16)
```

```
[95]: Text(0.5, 1.0, 'Triangle Correlation Heatmap')
```



As a result from the **heatmaps** we can see, that the shape of the **petals** are the **most correlated columns** (0.96) with the **type of flowers** (species classes).

Somewhat lower correlates **sepal length** with **petal length** (0.87).

4.6.2 Visualise data with scatter plot


In the following, [Seaborn](#) is applied which is a library for making statistical graphics in Python. It is built on top of matplotlib and closely integrated with pandas data structures.

To investigate whether there are dependencies (e.g. correlations) in `irisdata_df` between individual variables in the data set, it is advisable to plot them in a **scatter plot**.

```
[16]: # There are five preset seaborn themes: darkgrid, whitegrid, dark, white, and
      ↪ ticks.
sns.set_style("whitegrid")
# set scale of fonts
sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.5})

# 'sepal_length', 'petal_length' are iris feature data
# 'height' used to define height of graph
# 'hue' stores the class/label of iris dataset
sns.FacetGrid(irisdata_df, hue="species",
              height = 7).map(plt.scatter,
                              'petal_width',
                              'petal_length').add_legend()
```

```
[16]: <seaborn.axisgrid.FacetGrid at 0xa7c69810>
```



4.6.3 Visualise data with pairs plot

For systematic investigation of dependencies, all variables (each against each) are plotted in separate scatter plots.

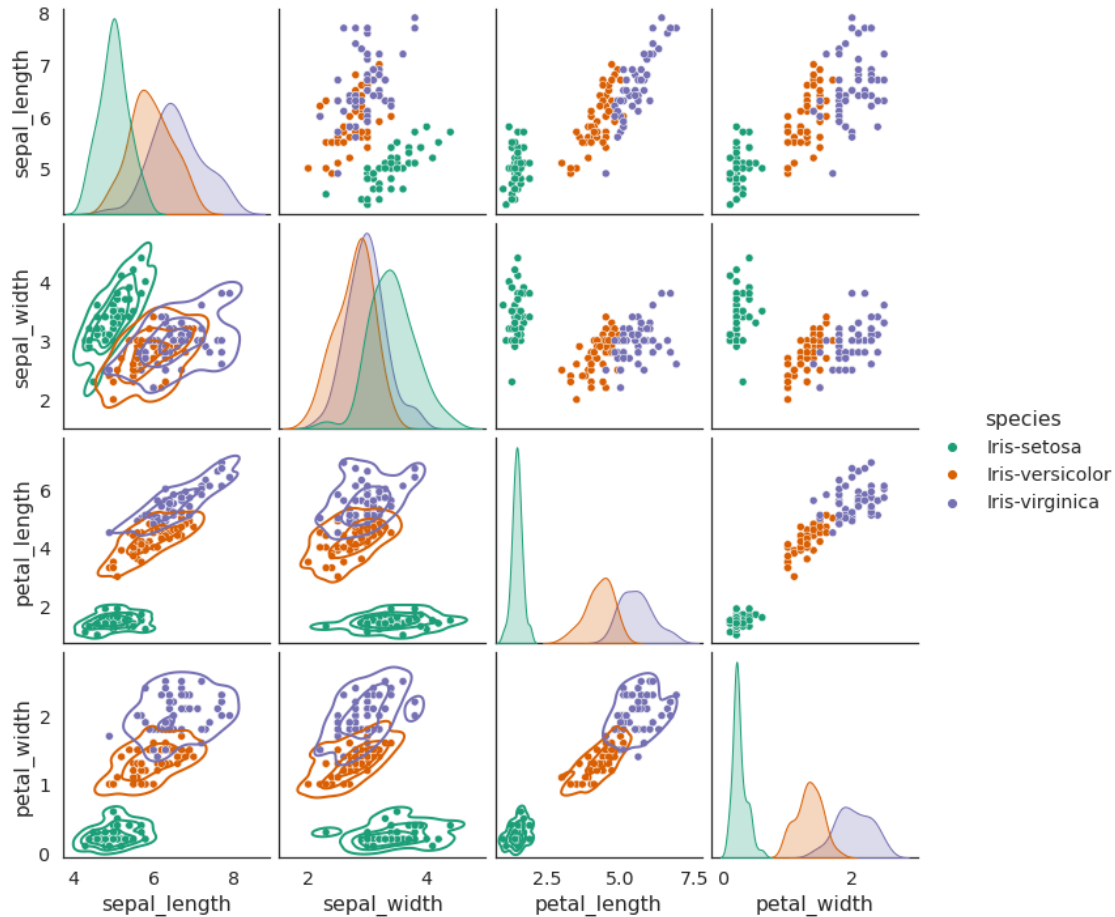
With this so called **pairs plot** it is possible to see both **relationships** between two variables and **distribution** of single variables.

This function will create a grid of Axes such that **each numeric variable** in `irisdata_df` will be shared in the y-axis across a single row and in the x-axis across a single column.

```
[60]: sns.set_style("white")
g = sns.pairplot(irisdata_df, diag_kind="kde", hue='species', palette='Dark2',
               height=2.5)

g.map_lower(sns.kdeplot, levels=4, color=".2")
```

```
[60]: <seaborn.axisgrid.PairGrid at 0x618ba670>
```



5 STEP 2: Prepare the data

Through the intensive exploration of the data in Step 1 ([STEP 1: Exploring the data](#)), we know that special **preparation** of the data is **not necessary**. The values are **complete** and **without gaps** and there are **no duplicates**. The values are in similar ranges, which **does not require normalization** of the data.

Furthermore, we know that the **classes** are very **evenly distributed** and thus bias tendencies should be avoided.

6 STEP 3: Classify by support vector classifier - SVC

6.1 Operating principal

Support Vectors Classifier tries to **find the best hyperplane to separate** the different classes by maximizing the distance between sample points and the hyperplane (source: [In Depth: Parameter tuning for SVC](#)).

Following graphic shows the operating principal of SVC: the hyperplane $H1$ does not separate the

classes. H_2 does, but only with a small margin. H_3 separates them with the maximal margin (source: [Support-vector machine](#)).

```
[7]: display(HTML("<figure><img src='../images/SVM_separating_hyperplanes.svg' \n
    \n
    <figcaption>SVC separete the data in classes by finding the \n
    \n
    best hyperplane (source: <a href='https://en.wikipedia.org/wiki/File:\n
    \n
    Svm_separating_hyperplanes_(SVG).svg'>Svm separating hyperplanes (SVG).svg</\n
    \n
    a></figcaption> \n
    \n
    </figure>"))
```

<IPython.core.display.HTML object>

6.2 Split the dataset

In the next very important step, the dataset is split into **2 subsets**: a **training dataset** and a **test dataset**. As the names suggest, the training dataset is used to train the ML algorithm. The test data set is then used to check the quality of the trained ML algorithm (here the **recognition rate**). For this purpose, the **class labels** are **removed** from the training data set - after all, these are to be predicted.

Typically, the **test dataset** set should contain **20%** of the entire dataset.

```
[43]: from sklearn.model_selection import train_test_split

X = irisdata_df.drop('species', axis=1)
y = irisdata_df['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

For training, do not use only the variables that correlate best with each other, but all of them.

Otherwise, the result of the prediction would be significantly worse. Maybe this is already an indication of **overfitting** of the ML model.

```
[42]: # do not use this!
X_train, X_test, y_train, y_test = train_test_split(X[['sepal_length', \n
    \n
    'sepal_width']], y, test_size = 0.20)
```

6.3 Create the SVM model

In this step we create the SVC model and fit it to our training data.

```
[44]: from sklearn.svm import SVC

classifier = SVC(kernel = 'linear', random_state = 0)

# fit the model for the data
classifier.fit(X_train, y_train)
```



```
[44]: SVC(kernel='linear', random_state=0)
```

6.4 Make predictions

```
[45]: y_pred = classifier.predict(X_test)
      #X_test
```

7 STEP 4: Evaluate the results - metrics

And finally for checking the accuracy of the model, the **confusion matrix** is used for the **cross validation**.

By using the function `sklearn.metrics.confusion_matrix()` a confusion matrix of the true digit values versus the predicted digit values is plotted.

7.1 Textual confusion matrix

```
[46]: cm = metrics.confusion_matrix(y_test, y_pred)
      print(cm)
```

```
[[ 8  0  0]
 [ 0  7  0]
 [ 0  0 15]]
```

7.2 Colored confusion matrix


The function `sklearn.metrics.ConfusionMatrixDisplay()` plots a colored confusion matrix.

```
[58]: sns.set_style("white")

      # print colored confusion matrix
      cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

      #cm_colored.figure_.suptitle("Confusion Matrix")
      cm_colored.figure_.set_figwidth(7)
      cm_colored.figure_.set_figheight(6)

      cm_colored.confusion_matrix
      plt.show()
```



```
[156]: from sklearn.model_selection import cross_val_score

accuracies = cross_val_score(estimator = classifier, X = X_train, y = y_train,
                               cv = 10)
print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

Accuracy: 96.67 %
Standard Deviation: 5.53 %

8 STEP 5: Vary parameters

This section was inspired by [In Depth: Parameter tuning for SVC](#)

In this section, the 4 SVC parameters `kernel`, `gamma`, `C` and `degree` will be introduced one by one. Furthermore, their influence on the classification result by varying these single parameters will be shown.

Disclaimer: In order to show the effects of varying the individual parameters in 2D graphs, only the best correlating variables `petal_length` and `petal_width` are used to train the SVC.

8.1 Prepare dataset

```
[98]: # import iris dataset again
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')

# encode the class column from class strings to integer equivalents
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":
    ↳0, "Iris-versicolor":1, "Iris-virginica":2}})
irisdata_df_enc
```

```
[98]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
..
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

[150 rows x 5 columns]

```
[116]: # copy only 2 feature columns
# and convert pandas dataframe to numpy array
X = irisdata_df_enc[['petal_length', 'petal_width']].to_numpy(copy=True)
#X = irisdata_df_enc[['sepal_length', 'sepal_width']].to_numpy(copy=True)
#X
```

```
[118]: # convert pandas dataframe to numpy array
# and get a flat 1D copy of 2D numpy array
y = irisdata_df_enc[['species']].to_numpy(copy=True).flatten()
#y
```

8.2 Plotting function

This function helps to visualize the modifications by varying the individual SVC parameters.

```
[101]: def plotSVC(title, xlabel, ylabel):
    # create a mesh to plot in
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

    # prevent division by zero
    if x_min == 0.0:
```

```

x_min = 0.1

h = (x_max / x_min)/1000
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

plt.subplot(1, 1, 1)
Z = svc.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
plt.xlabel(xlabel)
plt.ylabel(ylablel)
plt.xlim(xx.min(), xx.max())
plt.title(title)
plt.show()

```

8.3 Vary kernel parameter

The `kernel` parameter selects the type of hyperplane that is used to separate the data. Using `linear` ([linear classifier](#)) kernel will use a linear hyperplane (a line in the case of 2D data). The `rbf` ([radial basis function kernel](#)) and `poly` ([polynomial kernel](#)) kernel use non linear hyperplanes.



```


[102]: kernels = ['linear', 'rbf', 'poly']

xlabel = 'Petal length'
ylabel = 'Petal width'

for kernel in kernels:
    svc = svm.SVC(kernel=kernel).fit(X, y)
    plotSVC('kernel = ' + str(kernel), xlabel, ylabel)

```





8.4 Vary gamma parameter



The `gamma` parameter is used for non linear hyperplanes. The higher the `gamma` value it tries to exactly fit the training data set.



As we can see, increasing `gamma` leads to **overfitting** as the classifier tries to perfectly fit the training data.


```
[106]: gammas = [0.1, 1, 10, 100, 200]

xlabel = 'Petal length'
ylabel = 'Petal width'

for gamma in gammas:
    svc = svm.SVC(kernel='rbf', gamma=gamma).fit(X, y)
    plotSVC('gamma = ' + str(gamma), xlabel, ylabel)
```







8.5 Vary C parameter



The **C** parameter is the **penalty** of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly.



But be careful: to high **C** values may lead to **overfitting** the training data.


```
[108]: cs = [0.1, 1, 10, 100, 1000, 10000]


xlabel = 'Petal length'
ylabel = 'Petal width'

for c in cs:
    svc = svm.SVC(kernel='rbf', C=c).fit(X, y)
    plotSVC('C = ' + str(c), xlabel, ylabel)
```









8.6 Vary degree parameter



The `degree` parameter is used when the `kernel` is set to `poly`. It's basically the **degree of the polynomial** used to find the hyperplane to split the data.



Using `degree = 1` is the same as using a `linear` kernel. Also, increasing this parameters leads to **higher training times**.


```
[113]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]



xlabel = 'Petal length'
ylabel = 'Petal width'

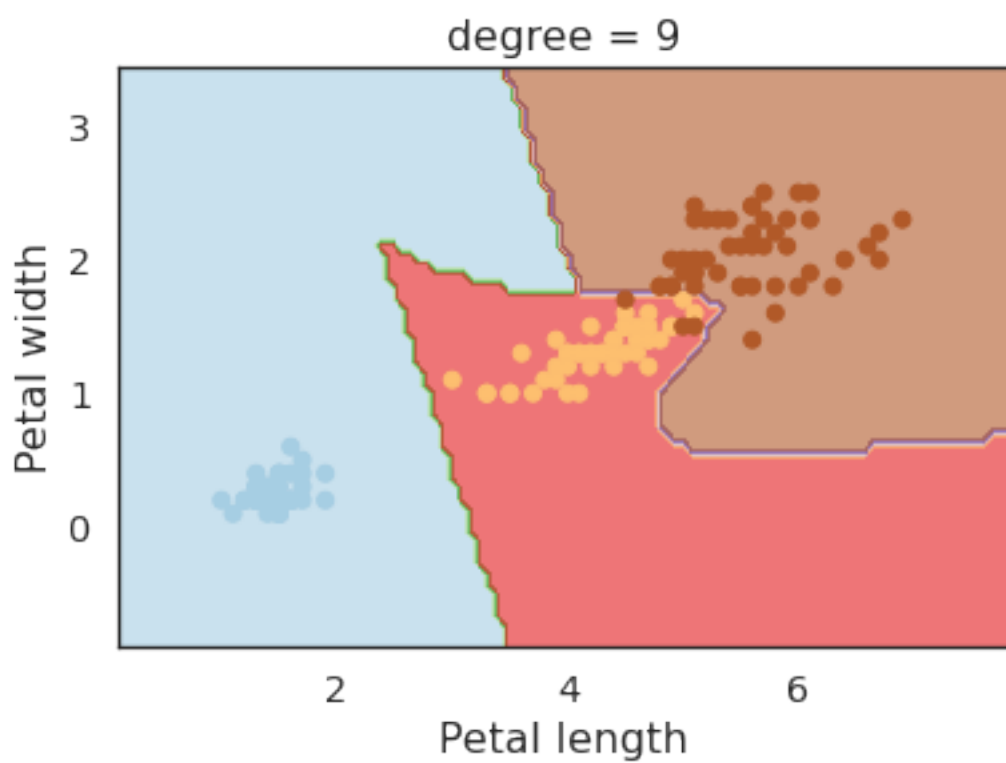

for degree in degrees:
    svc = svm.SVC(kernel='poly', degree=degree).fit(X, y)
    plotSVC('degree = ' + str(degree), xlabel, ylabel)
```














[]: