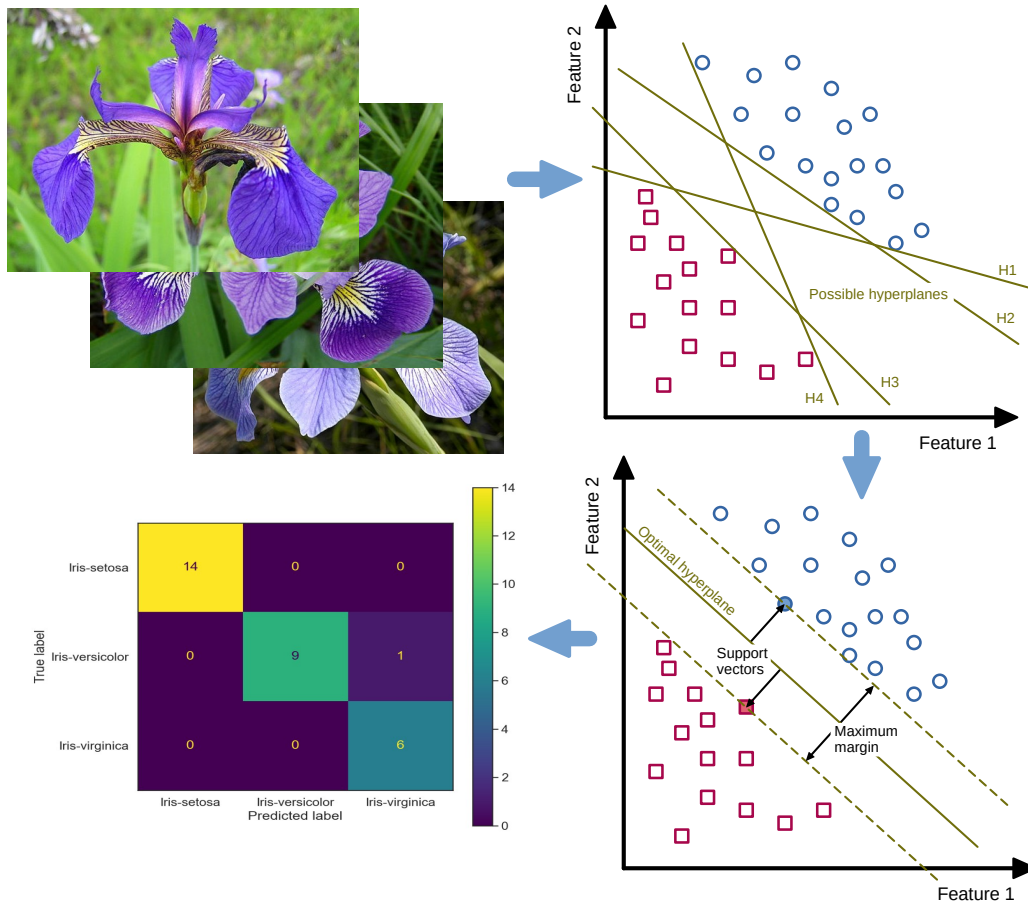


Getting started with Support Vector Classifiers (SVC) - A systematic step-by-step approach

Dipl.-Ing. Björn Kasper (kasper.bjoern@bgetem.de)

Test and Certification Body for Electrical Engineering at BG ETEM

August 4, 2022



Anyone who wants to seriously deal with the hypothetical topic of our time “Artificial Intelligence (AI)” or “Machine Learning (ML)” cannot avoid dealing with the basic ML algorithms, corresponding software tools, libraries and programming systems. However, someone who opens the door for the first time to this equally very exciting and arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. Here, it is a good idea to consult introductory and systematic tutorials. Therefore, this Getting Started tutorial systematically demonstrates the typical ML work process step-by-step using the very powerful and performant “Support Vector Classifier (SVC)” and the widely known and very beginner-friendly “Iris Dataset”. Furthermore, the selection of the “correct” SVC kernel and its parameters are described and their effect on the classification result is shown.



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/) (CC BY-SA 4.0).

Contents

1	Introduction	2
2	Load globally used libraries and set plot parameters	3
3	STEP 0: Get the data	3
4	STEP 1: Exploring the data	4
4.1	Goals of exploration	4
4.2	Clarify the origins history	4
4.3	Overview of the internal structure and organisation of the data	4
4.3.1	Inspect structure of dataframe	5
4.3.2	Get data types	9
4.3.3	Get data ranges with Boxplots	9
4.4	Identify anomalies in the data sets	10
4.4.1	Find gaps in dataset	10
4.4.2	Find and remove duplicates in dataset	17
4.5	Avoidance of tendencies due to bias	21
4.5.1	Count occurrences of unique values	22
4.5.2	Display Histogram	22
4.6	First idea of correlations in data set	24
4.6.1	Visualise data with correlation heatmap	24
4.6.2	Visualise data with scatter plot	26
4.6.3	Visualise data with pairs plot	27
5	STEP 2: Prepare the data	28
6	STEP 3: Classify by support vector classifier - SVC	28
6.1	Operating principal	28
6.2	Split the dataset	29
6.3	Create the SVM model	29
6.4	Make predictions	30
7	STEP 4: Evaluate the results - metrics	30
7.1	Textual confusion matrix	30
7.2	Colored confusion matrix	30
8	STEP 5: Vary parameters	31
8.1	Prepare dataset	31
8.2	Plotting function	32
8.3	Vary kernel parameter	33
8.4	Vary gamma parameter	35
8.5	Vary C parameter	35
8.6	Vary degree parameter	38

1 Introduction

Anyone who wants to seriously deal with the hypothetical topic of our time **Artificial Intelligence (AI)** or **Machine Learning (ML)** cannot avoid dealing with the basic ML algorithms, corresponding software tools, libraries and programming systems.

However, someone who opens the door for the first time to this equally very exciting and arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. Here, it is a good idea to consult introductory and systematic tutorials.

Therefore, this Getting Started tutorial systematically demonstrates the typical ML work process step-by-step using the very powerful and performant **Support Vector Classifier (SVC)** and the widely known and very beginner-friendly **Iris Dataset**.

This tutorial will be presented as part of a workshop at the DGUV symposium **Artificial Intelligence**, probably in November 2022 in Dresden. The workshop addresses interested ML novices.

For the target audience in the workshop, the SVC algorithm was intentionally chosen to show that there are many other very powerful and performant ML algorithms apart from the **deep neural networks** that are very present in the media. On the other hand, a necessary and comprehensible introduction to neural networks for newcomers would not be possible within the time frame given for the workshop.

Furthermore, this tutorial does *not* address the generation or acquisition of ML-ready datasets. A newcomer to ML will (or should) first try to familiarize himself with ML algorithms, tools, libraries and programming systems. Only then it makes sense to explore one's own environment with respect to ML-suitable applications and to acquire suitable data sets from them.

Therefore, the tutorial will demonstrate the usage of selected ML tools in the form of Python libraries as well as the systematic approach to the widely known and very beginner-friendly **Iris dataset**.

Furthermore, the selection of the “correct” SVC kernel and its parameters are described and their effect on the classification result is shown.

The following steps of the systematic ML process are covered in the next main sections:

- STEP 0: Get the data
- STEP 1: Exploring the data
- STEP 2: Prepare the data
- STEP 3: Classify by support vector classifier - SVC
- STEP 4: Evaluate the results - metrics
- STEP 5: Vary parameters

2 Load globally used libraries and set plot parameters

```
[1]: import time

from IPython.display import HTML

import pandas as pd
import matplotlib.pyplot as plt
from sklearn import svm, metrics
import seaborn as sns
%matplotlib inline
```

3 STEP 0: Get the data

Since this is intended to be an introduction to the world of machine learning (ML), this step does NOT deal with the design of an application suitable for ML and the acquisition of valid measurement data.

In order to get to know the typical work steps and ML tools, the use of **well-known and well-researched data sets** is clearly **recommended**.

In the further course, the famous [Iris flower data sets](#) will be used. It can be downloaded on [Iris Flower Dataset | Kaggle](#). Furthermore, the dataset is included in Python in the machine learning package [Scikit-learn](#), so that users can access it without having to find a special source for it.

```
[2]: # import some data to play with
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

4 STEP 1: Exploring the data

4.1 Goals of exploration

The objectives of the exploration of the dataset are as follows:

1. Clarify the **origins history**:
 - Where did the data come from? => Contact persons and licensing permissions?
 - Who obtained the data and with which (measurement) methods? => Did systematic errors occur during the acquisition?
 - What were they originally intended for? => Can they be used for my application?
2. Overview of the internal **structure and organisation** of the data:
 - Which columns are there? => With which methods can they be read in (e.g. import of CSV files)?
 - What do they contain for (physical) measured variables? => Which technical or physical correlations exist?
 - Which data formats or types are there? => Do they have to be converted?
 - In which value ranges do the measurement data vary? => Are normalizations necessary?
3. Identify **anomalies** in the data sets:
 - Do the data have **gaps** or **duplicates**? => Does the data set needs to be cleaned?
 - Are there obvious erroneous entries or measurement outliers? => Does (statistical) filtering have to be carried out?
4. Avoidance of **tendencies due to bias**:
 - Are all possible classes included in the dataset and equally distributed? => Does the data set need to be enriched with additional data for balance?
5. Find a first rough **idea of which correlations** could be in the data set

4.2 Clarify the origins history

The ***Iris* flower data sets** is a multivariate data set introduced by the British statistician and biologist *Ronald Fisher* in his paper “The use of multiple measurements in taxonomic problems as an example of linear discriminant analysis” (1936). It is sometimes called *Anderson’s Iris data set* because Edgar Anderson collected the data to quantify the morphologic variation of *Iris* flowers of three related species (source: [Iris flower data set](#)).

The dataset is published in Public Domain with a [CC0-License](#).

This dataset became a typical test case for many statistical classification techniques in machine learning such as **support vector machines**.

[..] measurements of the flowers of fifty plants each of the two species *Iris setosa* and *I. versicolor*, found **growing together in the same colony** and measured by Dr E. Anderson [..] (source: R. A. Fisher (1936). “The use of multiple measurements in taxonomic problems”. [Annals of Eugenics](#))

[..] *Iris virginica*, differs from the two other samples in **not being taken from the same natural colony** [..] (source: ibidem)

4.3 Overview of the internal structure and organisation of the data

The data set consists of 50 samples from each of three species of *Iris* (*Iris setosa*, *Iris virginica* and *Iris versicolor*), so there are 150 total samples. Four features were measured from each sample: the length and the width of the [sepals](#) and [petals](#), in centimetres.

Here is a principle illustration of a flower with sepal and petal:




Figure 1: Principle illustration of a flower with sepal and petal (source: [Mature_flower_diagram.svg](#), license: public domain)

Here are pictures of the three different Iris species (*Iris setosa*, *Iris virginica* and *Iris versicolor*). Given the dimensions of the flower, it will be possible to predict the class of the flower.



Figure 2: left: *Iris setosa* (source: [Irissetosa1.jpg](#), license: public domain); middle: *Iris versicolor* (source: [Iris_versicolor_3.jpg](#), license: CC-SA 3.0); right: *Iris virginica* (source: [Iris_virginica.jpg](#), license: CC-SA 2.0)

4.3.1 Inspect structure of dataframe

Print first or last 5 rows of dataframe:

```
[3]: irisdata_df.head()
```

```
[3]:   sepal_length  sepal_width  petal_length  petal_width  species
0         5.1         3.5         1.4         0.2  Iris-setosa
1         4.9         3.0         1.4         0.2  Iris-setosa
2         4.7         3.2         1.3         0.2  Iris-setosa
3         4.6         3.1         1.5         0.2  Iris-setosa
4         5.0         3.6         1.4         0.2  Iris-setosa
```

```
[4]: irisdata_df.tail()
```

```
[4]:   sepal_length  sepal_width  petal_length  petal_width  species
145         6.7         3.0         5.2         2.3  Iris-virginica
```

146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

While printing a dataframe - only an abbreviated view of the dataframe is shown :(
Default setting in the pandas library makes it to display only 5 lines from head and from tail.

```
[5]: irisdata_df
```

```
[5]:      sepal_length  sepal_width  petal_length  petal_width      species
0           5.1           3.5           1.4           0.2  Iris-setosa
1           4.9           3.0           1.4           0.2  Iris-setosa
2           4.7           3.2           1.3           0.2  Iris-setosa
3           4.6           3.1           1.5           0.2  Iris-setosa
4           5.0           3.6           1.4           0.2  Iris-setosa
..          ...           ...           ...           ...          ...
145          6.7           3.0           5.2           2.3  Iris-virginica
146          6.3           2.5           5.0           1.9  Iris-virginica
147          6.5           3.0           5.2           2.0  Iris-virginica
148          6.2           3.4           5.4           2.3  Iris-virginica
149          5.9           3.0           5.1           1.8  Iris-virginica
```

```
[150 rows x 5 columns]
```

To print all rows of a dataframe, the option `display.max_rows` has to set to `None` in pandas:

```
[6]: pd.set_option('display.max_rows', None)
irisdata_df
```

```
[6]:      sepal_length  sepal_width  petal_length  petal_width      species
0           5.1           3.5           1.4           0.2  Iris-setosa
1           4.9           3.0           1.4           0.2  Iris-setosa
2           4.7           3.2           1.3           0.2  Iris-setosa
3           4.6           3.1           1.5           0.2  Iris-setosa
4           5.0           3.6           1.4           0.2  Iris-setosa
5           5.4           3.9           1.7           0.4  Iris-setosa
6           4.6           3.4           1.4           0.3  Iris-setosa
7           5.0           3.4           1.5           0.2  Iris-setosa
8           4.4           2.9           1.4           0.2  Iris-setosa
9           4.9           3.1           1.5           0.1  Iris-setosa
10          5.4           3.7           1.5           0.2  Iris-setosa
11          4.8           3.4           1.6           0.2  Iris-setosa
12          4.8           3.0           1.4           0.1  Iris-setosa
13          4.3           3.0           1.1           0.1  Iris-setosa
14          5.8           4.0           1.2           0.2  Iris-setosa
15          5.7           4.4           1.5           0.4  Iris-setosa
16          5.4           3.9           1.3           0.4  Iris-setosa
17          5.1           3.5           1.4           0.3  Iris-setosa
18          5.7           3.8           1.7           0.3  Iris-setosa
19          5.1           3.8           1.5           0.3  Iris-setosa
20          5.4           3.4           1.7           0.2  Iris-setosa
21          5.1           3.7           1.5           0.4  Iris-setosa
22          4.6           3.6           1.0           0.2  Iris-setosa
23          5.1           3.3           1.7           0.5  Iris-setosa
24          4.8           3.4           1.9           0.2  Iris-setosa
25          5.0           3.0           1.6           0.2  Iris-setosa
26          5.0           3.4           1.6           0.4  Iris-setosa
```

27	5.2	3.5	1.5	0.2	Iris-setosa
28	5.2	3.4	1.4	0.2	Iris-setosa
29	4.7	3.2	1.6	0.2	Iris-setosa
30	4.8	3.1	1.6	0.2	Iris-setosa
31	5.4	3.4	1.5	0.4	Iris-setosa
32	5.2	4.1	1.5	0.1	Iris-setosa
33	5.5	4.2	1.4	0.2	Iris-setosa
34	4.9	3.1	1.5	0.1	Iris-setosa
35	5.0	3.2	1.2	0.2	Iris-setosa
36	5.5	3.5	1.3	0.2	Iris-setosa
37	4.9	3.1	1.5	0.1	Iris-setosa
38	4.4	3.0	1.3	0.2	Iris-setosa
39	5.1	3.4	1.5	0.2	Iris-setosa
40	5.0	3.5	1.3	0.3	Iris-setosa
41	4.5	2.3	1.3	0.3	Iris-setosa
42	4.4	3.2	1.3	0.2	Iris-setosa
43	5.0	3.5	1.6	0.6	Iris-setosa
44	5.1	3.8	1.9	0.4	Iris-setosa
45	4.8	3.0	1.4	0.3	Iris-setosa
46	5.1	3.8	1.6	0.2	Iris-setosa
47	4.6	3.2	1.4	0.2	Iris-setosa
48	5.3	3.7	1.5	0.2	Iris-setosa
49	5.0	3.3	1.4	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
51	6.4	3.2	4.5	1.5	Iris-versicolor
52	6.9	3.1	4.9	1.5	Iris-versicolor
53	5.5	2.3	4.0	1.3	Iris-versicolor
54	6.5	2.8	4.6	1.5	Iris-versicolor
55	5.7	2.8	4.5	1.3	Iris-versicolor
56	6.3	3.3	4.7	1.6	Iris-versicolor
57	4.9	2.4	3.3	1.0	Iris-versicolor
58	6.6	2.9	4.6	1.3	Iris-versicolor
59	5.2	2.7	3.9	1.4	Iris-versicolor
60	5.0	2.0	3.5	1.0	Iris-versicolor
61	5.9	3.0	4.2	1.5	Iris-versicolor
62	6.0	2.2	4.0	1.0	Iris-versicolor
63	6.1	2.9	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
65	6.7	3.1	4.4	1.4	Iris-versicolor
66	5.6	3.0	4.5	1.5	Iris-versicolor
67	5.8	2.7	4.1	1.0	Iris-versicolor
68	6.2	2.2	4.5	1.5	Iris-versicolor
69	5.6	2.5	3.9	1.1	Iris-versicolor
70	5.9	3.2	4.8	1.8	Iris-versicolor
71	6.1	2.8	4.0	1.3	Iris-versicolor
72	6.3	2.5	4.9	1.5	Iris-versicolor
73	6.1	2.8	4.7	1.2	Iris-versicolor
74	6.4	2.9	4.3	1.3	Iris-versicolor
75	6.6	3.0	4.4	1.4	Iris-versicolor
76	6.8	2.8	4.8	1.4	Iris-versicolor
77	6.7	3.0	5.0	1.7	Iris-versicolor
78	6.0	2.9	4.5	1.5	Iris-versicolor
79	5.7	2.6	3.5	1.0	Iris-versicolor
80	5.5	2.4	3.8	1.1	Iris-versicolor
81	5.5	2.4	3.7	1.0	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
83	6.0	2.7	5.1	1.6	Iris-versicolor

84	5.4	3.0	4.5	1.5	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
90	5.5	2.6	4.4	1.2	Iris-versicolor
91	6.1	3.0	4.6	1.4	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor
93	5.0	2.3	3.3	1.0	Iris-versicolor
94	5.6	2.7	4.2	1.3	Iris-versicolor
95	5.7	3.0	4.2	1.2	Iris-versicolor
96	5.7	2.9	4.2	1.3	Iris-versicolor
97	6.2	2.9	4.3	1.3	Iris-versicolor
98	5.1	2.5	3.0	1.1	Iris-versicolor
99	5.7	2.8	4.1	1.3	Iris-versicolor
100	6.3	3.3	6.0	2.5	Iris-virginica
101	5.8	2.7	5.1	1.9	Iris-virginica
102	7.1	3.0	5.9	2.1	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
104	6.5	3.0	5.8	2.2	Iris-virginica
105	7.6	3.0	6.6	2.1	Iris-virginica
106	4.9	2.5	4.5	1.7	Iris-virginica
107	7.3	2.9	6.3	1.8	Iris-virginica
108	6.7	2.5	5.8	1.8	Iris-virginica
109	7.2	3.6	6.1	2.5	Iris-virginica
110	6.5	3.2	5.1	2.0	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica
113	5.7	2.5	5.0	2.0	Iris-virginica
114	5.8	2.8	5.1	2.4	Iris-virginica
115	6.4	3.2	5.3	2.3	Iris-virginica
116	6.5	3.0	5.5	1.8	Iris-virginica
117	7.7	3.8	6.7	2.2	Iris-virginica
118	7.7	2.6	6.9	2.3	Iris-virginica
119	6.0	2.2	5.0	1.5	Iris-virginica
120	6.9	3.2	5.7	2.3	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
122	7.7	2.8	6.7	2.0	Iris-virginica
123	6.3	2.7	4.9	1.8	Iris-virginica
124	6.7	3.3	5.7	2.1	Iris-virginica
125	7.2	3.2	6.0	1.8	Iris-virginica
126	6.2	2.8	4.8	1.8	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica
128	6.4	2.8	5.6	2.1	Iris-virginica
129	7.2	3.0	5.8	1.6	Iris-virginica
130	7.4	2.8	6.1	1.9	Iris-virginica
131	7.9	3.8	6.4	2.0	Iris-virginica
132	6.4	2.8	5.6	2.2	Iris-virginica
133	6.3	2.8	5.1	1.5	Iris-virginica
134	6.1	2.6	5.6	1.4	Iris-virginica
135	7.7	3.0	6.1	2.3	Iris-virginica
136	6.3	3.4	5.6	2.4	Iris-virginica
137	6.4	3.1	5.5	1.8	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
139	6.9	3.1	5.4	2.1	Iris-virginica
140	6.7	3.1	5.6	2.4	Iris-virginica

141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

4.3.2 Get data types

```
[7]: irisdata_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
[8]: irisdata_df.describe()
```

```
[8]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

4.3.3 Get data ranges with Boxplots

Boxplots can be used to explore the data ranges in the dataset. These also provide information about outliers.

```
[9]: sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.0})
sns.set_style("whitegrid")
#sns.set_style("white")

fig, axs = plt.subplots(2, 2, figsize=(12, 10))

fn = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
box1 = sns.boxplot(x = 'species', y = 'sepal_length',
                   data = irisdata_df, order = cn, ax = axs[0,0])
box2 = sns.boxplot(x = 'species', y = 'sepal_width',
                   data = irisdata_df, order = cn, ax = axs[0,1])
```

```

box3 = sns.boxplot(x = 'species', y = 'petal_length',
                  data = irisdata_df, order = cn, ax = axs[1,0])
box4 = sns.boxplot(x = 'species', y = 'petal_width',
                  data = irisdata_df, order = cn, ax = axs[1,1])

# add some spacing between subplots
fig.tight_layout(pad=2.0)

plt.show()

```

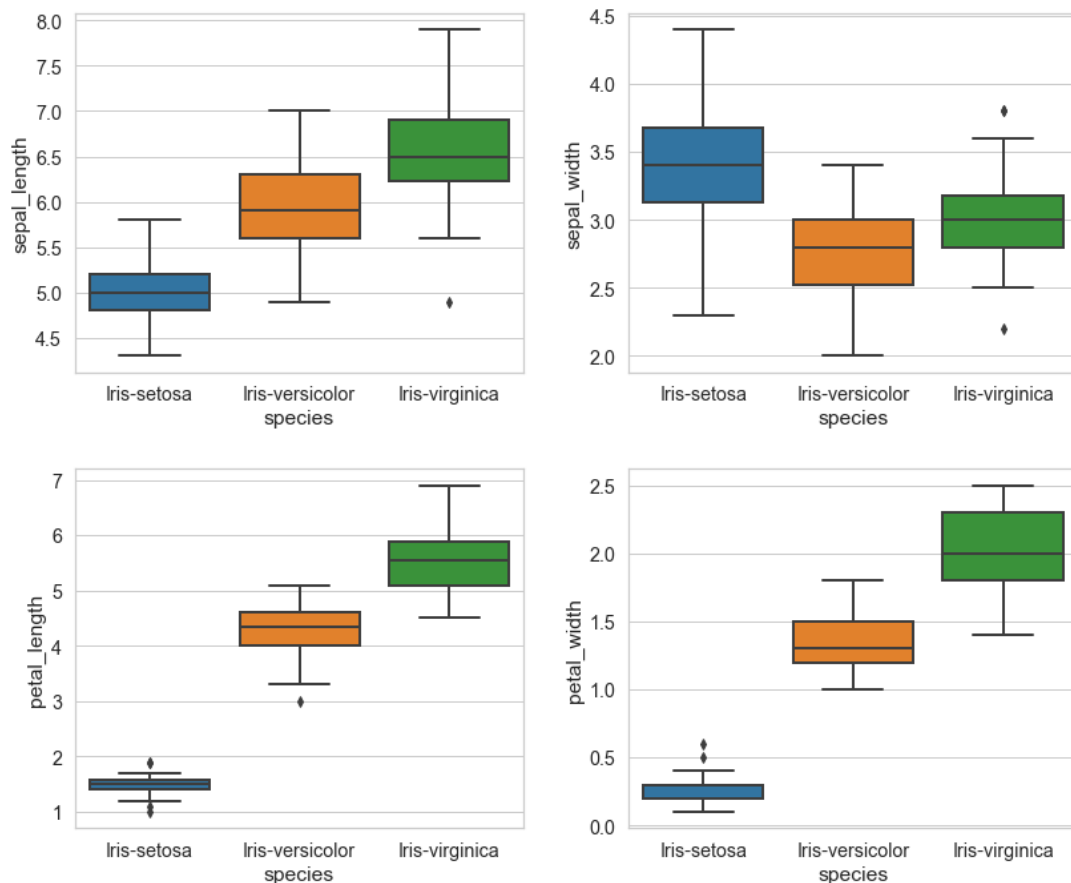


Figure 3: Boxplots used to explore the data ranges in the Iris dataset

4.4 Identify anomalies in the data sets

4.4.1 Find gaps in dataset

This section was inspired by [Working with Missing Data in Pandas](#).

Checking for missing values using `isnull()` In order to check for missing values in Pandas DataFrame, we use the function `isnull()`. This function returns a dataframe of Boolean values which are True for NaN values.

```

[10]: pd.set_option('display.max_rows', 40)
      pd.set_option('display.min_rows', 30)

```

```

[11]: irisdata_df.isnull()

```

```
[11]:      sepal_length  sepal_width  petal_length  petal_width  species
0          False          False          False          False    False
1          False          False          False          False    False
2          False          False          False          False    False
3          False          False          False          False    False
4          False          False          False          False    False
5          False          False          False          False    False
6          False          False          False          False    False
7          False          False          False          False    False
8          False          False          False          False    False
9          False          False          False          False    False
10         False          False          False          False    False
11         False          False          False          False    False
12         False          False          False          False    False
13         False          False          False          False    False
14         False          False          False          False    False
..         ...           ...           ...           ...           ...
135        False          False          False          False    False
136        False          False          False          False    False
137        False          False          False          False    False
138        False          False          False          False    False
139        False          False          False          False    False
140        False          False          False          False    False
141        False          False          False          False    False
142        False          False          False          False    False
143        False          False          False          False    False
144        False          False          False          False    False
145        False          False          False          False    False
146        False          False          False          False    False
147        False          False          False          False    False
148        False          False          False          False    False
149        False          False          False          False    False
```

[150 rows x 5 columns]

Show only the gaps:

```
[12]: irisdata_df_gaps = irisdata_df[irisdata_df.isnull().any(axis=1)]
irisdata_df_gaps
```

```
[12]: Empty DataFrame
Columns: [sepal_length, sepal_width, petal_length, petal_width, species]
Index: []
```

Fine - this dataset seems to be complete :)

So let's look for something else for exercise: [employees.csv](#)

```
[13]: # import data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")

employees_df
```

```
[13]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus % \
0      Douglas    Male    8/6/1993      12:42 PM    97308    6945.00
1      Thomas    Male    3/31/1996      6:53 AM    61933      4.17
2      Maria    Female    4/23/1993      11:17 AM   130590   11858.00
3      Jerry     Male    3/4/2005      1:00 PM   138705      9.34
```

4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
...
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance

```

1001          False          Product
1002          False Business Development
1003           True           Sales

```

```
[1004 rows x 8 columns]
```

Show only the gaps from this gappy dataset again:

```
[14]: employees_df_gaps = employees_df[employees_df.isnull().any(axis=1)]
employees_df_gaps
```

```
[14]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
20	Lois	NaN	4/22/1995	7:18 PM	64714	4934.00	
22	Joshua	NaN	3/8/2012	1:58 AM	90816	18816.00	
23	NaN	Male	6/14/2012	4:19 PM	125792	5042.00	
25	NaN	Male	10/8/2012	1:12 AM	37076	18576.00	
27	Scott	NaN	7/11/1991	6:58 PM	122367	5218.00	
31	Joyce	NaN	2/20/2005	2:40 PM	88657	12752.00	
32	NaN	Male	8/21/1998	2:27 PM	122340	6417.00	
39	NaN	Male	1/29/2016	2:33 AM	122173	7797.00	
41	Christine	NaN	6/28/2015	1:08 AM	66582	11308.00	
49	Chris	NaN	1/24/1980	12:13 PM	113590	3055.00	
51	NaN	NaN	12/17/2011	8:29 AM	41126	14009.00	
53	Alan	NaN	3/3/2014	1:28 PM	40341	17578.00	
..	
916	Joe	Male	12/8/1998	10:28 AM	126120	1.02	
927	Irene	NaN	2/28/1991	10:23 PM	135369	4.38	
929	NaN	Female	8/23/2000	4:19 PM	95866	19388.00	
941	Aaron	NaN	1/22/1986	7:39 PM	63126	18424.00	
942	Mark	NaN	9/9/2006	12:27 PM	44836	2657.00	
943	Ralph	NaN	7/28/1995	6:53 PM	70635	2147.00	
949	Gerald	NaN	4/15/1989	12:44 PM	93712	17426.00	
950	NaN	Female	9/15/1985	1:50 AM	133472	16941.00	
951	NaN	Male	7/30/2012	3:07 PM	107351	5329.00	
955	NaN	Female	9/14/2010	5:19 AM	143638	9662.00	
965	Antonio	NaN	6/18/1989	9:37 PM	103050	3.05	
976	Victor	NaN	7/28/2006	2:49 PM	76381	11159.00	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	

	Senior Management	Team
1	True	NaN
7	NaN	Finance
10	True	NaN
20	True	Legal
22	True	Client Services
23	NaN	NaN
25	NaN	Client Services
27	False	Legal
31	False	Product
32	NaN	NaN
39	NaN	Client Services
41	True	Business Development
49	False	Sales

```

51          NaN          Sales
53          True          Finance
..          ...          ...
916         False         NaN
927         False  Business Development
929          NaN          Sales
941         False  Client Services
942         False  Client Services
943         False  Client Services
949          True    Distribution
950          NaN    Distribution
951          NaN    Marketing
955          NaN          NaN
965         False    Legal
976          True    Sales
989         False    Legal
993         False    Legal
999         False  Distribution

```

[237 rows x 8 columns]

Fill the missing values with fillna() Now we are going to fill all the null (NaN) values in Gender column with "No Gender".

Attention: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

```

[15]: # filling a null values using fillna()
employees_df["Gender"].fillna("No Gender", inplace = True)
employees_df

```

```

[15]:   First Name  Gender  Start Date  Last Login Time  Salary  Bonus % \
0    Douglas    Male    8/6/1993      12:42 PM    97308    6945.00
1     Thomas    Male    3/31/1996      6:53 AM    61933      4.17
2      Maria  Female    4/23/1993     11:17 AM   130590   11858.00
3      Jerry    Male    3/4/2005      1:00 PM   138705      9.34
4      Larry    Male    1/24/1998      4:47 PM   101004    1389.00
5     Dennis    Male    4/18/1987      1:35 AM   115163   10125.00
6       Ruby  Female    8/17/1987      4:20 PM    65476   10012.00
7         NaN  Female    7/20/2015     10:43 AM    45906   11598.00
8     Angela  Female   11/22/2005      6:29 AM    95570   18523.00
9    Frances  Female    8/8/2002      6:51 AM   139852    7524.00
10    Louise  Female    8/12/1980      9:01 AM    63241   15132.00
11     Julie  Female   10/26/1997      3:19 PM   102508   12637.00
12   Brandon    Male    12/1/1980      1:08 AM   112807   17492.00
13     Gary    Male    1/27/2008     11:40 PM   109831    5831.00
14  Kimberly  Female    1/14/1999      7:13 AM    41426   14543.00
...     ...     ...     ...     ...     ...
989   Stephen  No Gender    7/10/1983      8:10 PM    85668    1909.00
990    Donna  Female   11/26/1982      7:04 AM    82871   17999.00
991   Gloria  Female   12/8/2014      5:08 AM   136709   10331.00
992   Alice  Female   10/5/2004      9:34 AM    47638   11209.00
993   Justin  No Gender    2/10/1991      4:58 PM    38344    3794.00
994   Robin  Female    7/24/1987      1:35 PM   100765   10982.00
995    Rose  Female    8/25/2002      5:12 AM   134505   11051.00
996  Anthony    Male   10/16/2011      8:35 AM   112769   11625.00
997    Tina  Female    5/15/1997      3:53 PM    56450     19.04

```

998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[1004 rows x 8 columns]

Dropping missing values using dropna() In order to drop null values from a dataframe, we use dropna() function. This function drops rows or columns of datasets with NaN values in different ways.

Default is to drop rows with at least 1 null value (NaN). Giving the parameter **how** = 'all' the function drops rows with all data missing or contain null values (NaN).

```
[16]: # making a new dataframe with dropped NaN values
employees_df_dropped = employees_df.dropna(axis = 0, how = 'any')
employees_df_dropped
```

```
[16]:   First Name  Gender  Start Date  Last Login Time  Salary  Bonus % \
0    Douglas    Male    8/6/1993      12:42 PM    97308    6945.00
2     Maria   Female    4/23/1993      11:17 AM   130590   11858.00
3     Jerry    Male    3/4/2005       1:00 PM   138705     9.34
4     Larry    Male    1/24/1998       4:47 PM   101004   1389.00
```

5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
15	Lillian	Female	6/5/2016	6:09 AM	59414	1256.00
16	Jeremy	Male	9/21/2010	5:56 AM	90370	7369.00
17	Shawn	Male	12/7/1986	7:45 PM	111737	6414.00
...
989	Stephen	No Gender	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	No Gender	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
8	True	Engineering
9	True	Business Development
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
15	False	Product
16	False	Human Resources
17	False	Product
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance


```

1001          False          Product
1002          False Business Development
1003           True           Sales

```

```
[903 rows x 8 columns]
```

Finally we compare the sizes of dataframes so that we learn how many rows had at least 1 Null value.

```
[17]: print("Old data frame length:", len(employees_df))
      print("New data frame length:", len(employees_df_dropped))
      print("Number of rows with at least 1 NaN value: ",
            (len(employees_df)-len(employees_df_dropped)))
```

```
Old data frame length: 1004
```

```
New data frame length: 903
```

```
Number of rows with at least 1 NaN value: 101
```

4.4.2 Find and remove duplicates in dataset

This section was inspired by: - [How to Find Duplicates in Pandas DataFrame \(With Examples\)](#) - [How to Drop Duplicate Rows in a Pandas DataFrame](#)

Checking for duplicate values using duplicated() In order to check for duplicate values in Pandas DataFrame, we use a function `duplicated()`. This function can be used in two ways: - find duplicate rows across **all columns** with `duplicateRows = df[df.duplicated()]` - find duplicate rows across **specific columns** `duplicateRows = df[df.duplicated(subset=['col1', 'col2'])]`

Find duplicate rows across **all columns**:

```
[18]: # import (again) data to dataframe from csv file
      employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[19]: # find duplicate rows across all columns
      duplicateRows = employees_df[employees_df.duplicated()]
      duplicateRows
```

```
[19]:   First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
112      Karen  Female  11/30/1999      7:46 AM  102488  17653.0
127      Linda  Female   5/25/2000      5:45 PM  119009  12506.0
296    Brandon   NaN    11/3/1997      8:17 PM  121333  15295.0
580   Nicholas   Male    3/1/2013      9:26 PM  101036   2826.0
```

```

      Senior Management          Team
112          True          Product
127          True Business Development
296          False Business Development
580          True   Human Resources

```

```
[20]: # argument keep='last' displays the first duplicate rows instead of the last
      duplicateRows = employees_df[employees_df.duplicated(keep='last')]
      duplicateRows
```

```
[20]:   First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
55      Karen  Female  11/30/1999      7:46 AM  102488  17653.0
92      Linda  Female   5/25/2000      5:45 PM  119009  12506.0
153    Brandon   NaN    11/3/1997      8:17 PM  121333  15295.0
442   Nicholas   Male    3/1/2013      9:26 PM  101036   2826.0
```

	Senior Management	Team
55	True	Product
92	True	Business Development
153	False	Business Development
442	True	Human Resources

Find duplicate rows across **specific columns**:

```
[21]: # identify duplicate rows across 'First Name' and 'Last Login Time' columns
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'])]
duplicateRows
```

```
[21]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
112	Karen	Female	11/30/1999	7:46 AM	102488	17653.0	
127	Linda	Female	5/25/2000	5:45 PM	119009	12506.0	
296	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.0	
577	NaN	Female	1/13/2009	1:01 PM	118736	7421.0	
580	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.0	
632	NaN	NaN	9/2/1988	12:49 PM	147309	1702.0	
881	NaN	Male	9/5/1980	7:36 AM	114896	13823.0	
929	NaN	Female	8/23/2000	4:19 PM	95866	19388.0	
934	Nancy	Female	9/10/2001	11:57 PM	85213	2386.0	
973	Linda	Female	2/4/2010	8:49 PM	44486	17308.0	

	Senior Management	Team
112	True	Product
127	True	Business Development
296	False	Business Development
577	NaN	Client Services
580	True	Human Resources
632	NaN	Distribution
881	NaN	Client Services
929	NaN	Sales
934	True	Marketing
973	True	Engineering

```
[22]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'], keep='last')]
duplicateRows
```

```
[22]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
23	NaN	Male	6/14/2012	4:19 PM	125792	5042.00	
37	Linda	Female	10/19/1981	8:49 PM	57427	9557.00	
55	Karen	Female	11/30/1999	7:46 AM	102488	17653.00	
66	Nancy	Female	12/15/2012	11:57 PM	125250	2672.00	
92	Linda	Female	5/25/2000	5:45 PM	119009	12506.00	
153	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.00	
222	NaN	Female	6/17/1991	12:49 PM	71945	5.56	
269	NaN	Male	2/4/2005	1:01 PM	40451	16044.00	
442	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.00	
778	NaN	Female	6/18/2000	7:36 AM	106428	10867.00	

	Senior Management	Team
23	NaN	NaN

```

37          True      Client Services
55          True      Product
66          True Business Development
92          True Business Development
153         False Business Development
222         NaN      Marketing
269         NaN      Distribution
442         True      Human Resources
778         NaN      NaN

```

Dropping duplicate values using drop_duplicates() In order to drop duplicate values from a dataframe, we use drop_duplicates() function.

This function can be used in two ways: - remove duplicate rows across **all columns** with df.drop_duplicates() - find duplicate rows across **specific columns** df.drop_duplicates(subset=['col1', 'col2'])

Attention: We are doing that directly in this dataframe with inplace = True - we don't make a deep copy!

Remove duplicate rows across **all columns**:

```

[23]: # remove duplicate rows across all columns
employees_df.drop_duplicates(inplace=True)
employees_df

```

```

[23]:
   First Name  Gender  Start Date  Last Login Time  Salary  Bonus % \
0    Douglas    Male    8/6/1993      12:42 PM    97308    6945.00
1     Thomas    Male    3/31/1996      6:53 AM    61933      4.17
2      Maria  Female    4/23/1993      11:17 AM   130590   11858.00
3      Jerry    Male    3/4/2005      1:00 PM   138705     9.34
4      Larry    Male    1/24/1998      4:47 PM   101004    1389.00
5     Dennis    Male    4/18/1987      1:35 AM   115163   10125.00
6       Ruby  Female    8/17/1987      4:20 PM    65476   10012.00
7        NaN  Female    7/20/2015     10:43 AM    45906   11598.00
8     Angela  Female   11/22/2005      6:29 AM    95570   18523.00
9    Frances  Female    8/8/2002      6:51 AM   139852    7524.00
10    Louise  Female    8/12/1980      9:01 AM    63241   15132.00
11     Julie  Female   10/26/1997      3:19 PM   102508   12637.00
12   Brandon    Male   12/1/1980      1:08 AM   112807   17492.00
13     Gary    Male   1/27/2008     11:40 PM   109831    5831.00
14  Kimberly  Female   1/14/1999      7:13 AM    41426   14543.00
...
989   Stephen   NaN    7/10/1983      8:10 PM    85668    1909.00
990    Donna  Female   11/26/1982      7:04 AM    82871   17999.00
991   Gloria  Female   12/8/2014      5:08 AM   136709   10331.00
992    Alice  Female   10/5/2004      9:34 AM    47638   11209.00
993   Justin   NaN    2/10/1991      4:58 PM    38344    3794.00
994    Robin  Female    7/24/1987      1:35 PM   100765   10982.00
995     Rose  Female    8/25/2002      5:12 AM   134505   11051.00
996  Anthony    Male   10/16/2011      8:35 AM   112769   11625.00
997     Tina  Female    5/15/1997      3:53 PM    56450     19.04
998   George    Male    6/21/2013      5:47 PM    98874    4479.00
999    Henry    NaN   11/23/2014      6:09 AM   132483   16655.00
1000  Phillip    Male   1/31/1984      6:30 AM    42392   19675.00
1001  Russell    Male   5/20/2013     12:39 PM    96914    1421.00
1002   Larry    Male   4/20/2013      4:45 PM    60500   11985.00
1003  Albert    Male   5/15/2012      6:24 PM   129949   10169.00

```

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[1000 rows x 8 columns]

Remove duplicate rows across **specific columns**:

```
[24]: # remove duplicate rows across 'First Name' and 'Last Login Time' columns
employees_df.drop_duplicates(
    subset=['First Name', 'Last Login Time'], keep='last', inplace=True)
employees_df
```

```
[24]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	

14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
...
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[994 rows x 8 columns]

4.5 Avoidance of tendencies due to bias

The description of the Iris dataset says, that it consists of **50 samples** from **each of three species** of Iris (Iris setosa, Iris virginica and Iris versicolor), so there are **150 total samples**.

But how to prove it?

4.5.1 Count occurrences of unique values

To prove whether all possible classes included in the dataset and equally distributed, you can use the function `df.value_counts`.

Following parameters can be used for fine tuning: - `dropna=False` causes that NaN values are included - `normalize=True`: relative frequencies of the unique values are returned - `ascending=False`: sort resulting classes descending

```
[25]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[26]: # count unique values without missing values in a column,
# ordered descending and normalized
irisdata_df['species'].value_counts(ascending=False, dropna=False, normalize=True)
```

```
[26]: Iris-setosa      0.333333
Iris-versicolor    0.333333
Iris-virginica     0.333333
Name: species, dtype: float64
```

```
[27]: # count unique values and missing values in a column,
# ordered descending and not absolute values
employees_df['Team'].value_counts(ascending=False, dropna=False, normalize=False)
```

```
[27]: Client Services      106
Business Development    103
Finance                 102
Marketing                98
Product                 96
Sales                   94
Engineering              92
Human Resources          92
Distribution             90
Legal                    88
NaN                      43
Name: Team, dtype: int64
```

4.5.2 Display Histogram

This section was inspired by: [Pandas Histogram – DataFrame.hist\(\)](#).

Histograms represent **frequency distributions** graphically. This requires the separation of the data into classes (so-called **bins**).

These classes are represented in the histogram as rectangles of equal or variable width. The height of each rectangle then represents the (relative or absolute) **frequency density**.

```
[28]: employees_df.hist(column=['Salary'])
plt.show()
```

```
[29]: employees_df.hist(column='Salary', by='Gender')
plt.show()
```



Figure 4:




Figure 5:

4.6 First idea of correlations in data set

To get a rough idea of the **dependencies** and **correlations** in the data set, it can be helpful to visualize the whole dataset in a **correlation heatmap**. They show in a glance which variables are correlated, to what degree and in which direction.

Later, 2 particularly well correlated variables are selected from the data set and plotted in a **scatterplot**.

4.6.1 Visualise data with correlation heatmap

This section was inspired by [How to Create a Seaborn Correlation Heatmap in Python?](#).

Correlation matrices are an **essential tool of exploratory data analysis**. Correlation heatmaps contain the same information in a visually appealing way. What more: they show in a glance which variables are correlated, to what degree, in which direction, and alerts us to potential multicollinearity problems (source: ibidem).

Simple correlation matrix Because **string values can never be correlated**, the class names (species) have to be converted first:

```
[30]: # encoding the class column
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
                                                    "Iris-virginica":2}})

irisdata_df_enc
```

```
[30]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
5	5.4	3.9	1.7	0.4	0
6	4.6	3.4	1.4	0.3	0
7	5.0	3.4	1.5	0.2	0
8	4.4	2.9	1.4	0.2	0
9	4.9	3.1	1.5	0.1	0
10	5.4	3.7	1.5	0.2	0
11	4.8	3.4	1.6	0.2	0
12	4.8	3.0	1.4	0.1	0
13	4.3	3.0	1.1	0.1	0
14	5.8	4.0	1.2	0.2	0
..
135	7.7	3.0	6.1	2.3	2
136	6.3	3.4	5.6	2.4	2
137	6.4	3.1	5.5	1.8	2
138	6.0	3.0	4.8	1.8	2
139	6.9	3.1	5.4	2.1	2
140	6.7	3.1	5.6	2.4	2
141	6.9	3.1	5.1	2.3	2
142	5.8	2.7	5.1	1.9	2
143	6.8	3.2	5.9	2.3	2
144	6.7	3.3	5.7	2.5	2
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

[150 rows x 5 columns]

```
[31]: irisdata_df_enc.corr()
```

```
[31]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
sepal_length	1.000000	-0.109369	0.871754	0.817954	0.782561
sepal_width	-0.109369	1.000000	-0.420516	-0.356544	-0.419446
petal_length	0.871754	-0.420516	1.000000	0.962757	0.949043
petal_width	0.817954	-0.356544	0.962757	1.000000	0.956464
species	0.782561	-0.419446	0.949043	0.956464	1.000000

Correlation heatmap Choose the color sets from [color map](#).

```
[32]: # increase the size of the heatmap
plt.figure(figsize=(16, 6))

# store heatmap object in a variable to easily access it
# when you want to include more features (such as title)
# set the range of values to be displayed on the colormap from -1 to 1,
# and set 'annotation=True' to display the correlation values on the heatmap
heatmap = sns.heatmap(irisdata_df_enc.corr(), vmin=-1, vmax=1,
                      annot=True, cmap='PRGn_r')

# give a title to the heatmap
# 'pad=12' defines the distance of the title from the top of the heatmap
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
plt.show()
```




Figure 6:

Triangle correlation heatmap When looking at the correlation heatmaps above, you would not lose any information by **cutting** away half of it **along the diagonal** line marked by 1-s.

The **numpy** function `np.triu()` can be used to isolate the upper triangle of a matrix while turning all the values in the lower triangle into 0.

```
[33]: import numpy as np

np.triu(np.ones_like(irisdata_df_enc.corr()))
```

```
[33]: array([[1., 1., 1., 1., 1.],
           [0., 1., 1., 1., 1.],
           [0., 0., 1., 1., 1.],
           [0., 0., 0., 1., 1.],
           [0., 0., 0., 0., 1.]])
```

Use this mask to cut the heatmap along the diagonal:

```
[34]: plt.figure(figsize=(16, 6))

# define the mask to set the values in the upper triangle to 'True'
mask = np.triu(np.ones_like(irisdata_df_enc.corr(), dtype=bool))

heatmap = sns.heatmap(irisdata_df_enc.corr(), mask=mask,
                      vmin=-1, vmax=1, annot=True, cmap='PRGn_r')

heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
plt.show()
```




Figure 7:

As a result from the **heatmaps** we can see, that the shape of the **petals** are the **most correlated columns** (0.96) with the **type of flowers** (species classes).

Somewhat lower correlates **sepal length** with **petal length** (0.87).

4.6.2 Visualise data with scatter plot

In the following, **Seaborn** is applied which is a library for making statistical graphics in Python. It is built on top of matplotlib and closely integrated with pandas data structures.

To investigate whether there are dependencies (e.g. correlations) in **irisdata_df** between individual variables in the data set, it is advisable to plot them in a **scatter plot**.

```
[69]: # There are five preset seaborn themes: darkgrid, whitegrid, dark, white, and ticks.
sns.set_style("whitegrid")
# set scale of fonts
sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.5})

# 'sepal_length', 'petal_length' are iris feature data
# 'height' used to define height of graph
# 'hue' stores the class/label of iris dataset
sns.FacetGrid(irisdata_df, hue="species",
```

```

height = 7).map(plt.scatter,
                'petal_width',
                'petal_length').add_legend()

plt.title('Scatterplot of petal length and width')
plt.show()

```




Figure 8:

4.6.3 Visualise data with pairs plot

For systematic investigation of dependencies, all variables (each against each) are plotted in separate scatter plots.

With this so called **pairs plot** it is possible to see both **relationships** between two variables and **distribution** of single variables.

This function will create a grid of Axes such that **each numeric variable** in `irisdata_df` will be shared in the y-axis across a single row and in the x-axis across a single column.

```

[36]: sns.set_style("white")
      g = sns.pairplot(irisdata_df, diag_kind="kde", hue='species',
                      palette='Dark2', height=2.5)

      g.map_lower(sns.kdeplot, levels=4, color=".2")

      plt.show()

```




Figure 9:

5 STEP 2: Prepare the data

Through the intensive exploration of the data in Step 1 ([STEP 1: Exploring the data](#)), we know that special **preparation** of the data is **not necessary**. The values are **complete** and **without gaps** and there are **no duplicates**. The values are in similar ranges, which **does not require normalization** of the data.

Furthermore, we know that the **classes** are very **evenly distributed** and thus bias tendencies should be avoided.

6 STEP 3: Classify by support vector classifier - SVC

6.1 Operating principal

Support Vectors Classifier tries to **find the best hyperplane to separate** the different classes by maximizing the distance between sample points and the hyperplane (source: [In Depth: Parameter tuning for SVC](#)).

The figure ?? shows the operating principal of the SVC algorithm: the hyperplanes $H1$ till $H4$ (left graphic) do separate the classes. A good separation is achieved by the hyperplane that has the largest distance to the nearest training-data point of any class (so-called functional margin), since in general the larger the margin, the lower the generalization error of the classifier (source: [Support-vector machine](#)).

The right graphic shows the optimal hyperplane characterized by maximising the margin between the classes. The perpendicular distance of the closest data points to the hyperplane determines their position and orientation. These perpendicular distances are the **support vectors** of the hyperplane - this is how the algorithm got its name.




Figure 10: Support Vectors Classifiers (SVC) separate the data points in classes by finding the best hyperplane

6.2 Split the dataset

In the next very important step, the dataset is split into **2 subsets**: a **training dataset** and a **test dataset**. As the names suggest, the training dataset is used to train the ML algorithm. The test data set is then used to check the quality of the trained ML algorithm (here the **recognition rate**). For this purpose, the **class labels** are **removed** from the training data set - after all, these are to be predicted.

Typically, the **test dataset** should contain about **20%** of the entire dataset.

```
[52]: from sklearn.model_selection import train_test_split

X = irisdata_df.drop('species', axis=1)
y = irisdata_df['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

For training, do not use only the variables that correlate best with each other, but all of them.

Otherwise, the result of the prediction would be significantly worse. Maybe this is already an indication of **overfitting** of the ML model.

```
[38]: # DO NOT USE THIS!!
X_train, X_test, y_train, y_test = train_test_split(X[['sepal_length',
                                                    'sepal_width']],
                                                    y, test_size = 0.20)
```

6.3 Create the SVM model

In this step we create the SVC model and fit it to our training data.

```
[53]: from sklearn.svm import SVC
classifier = SVC(kernel = 'linear', random_state = 0)

# fit the model for the data
classifier.fit(X_train, y_train)
```

```
[53]: SVC(kernel='linear', random_state=0)
```

6.4 Make predictions

```
[54]: y_pred = classifier.predict(X_test)
      #X_test
```

7 STEP 4: Evaluate the results - metrics

And finally for checking the accuracy of the model, the **confusion matrix** is used for the **cross validation**.

By using the function `sklearn.metrics.confusion_matrix()` a confusion matrix of the true digit values versus the predicted digit values is plotted.

7.1 Textual confusion matrix

```
[55]: cm = metrics.confusion_matrix(y_test, y_pred)
      print(cm)
```

```
[[14  0  0]
 [ 0  9  1]
 [ 0  0  6]]
```

7.2 Colored confusion matrix

The function `sklearn.metrics.ConfusionMatrixDisplay()` plots a colored confusion matrix.

```
[67]: sns.set_style("white")

      # print colored confusion matrix
      cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

      #cm_colored.figure_.suptitle("Confusion Matrix")
      cm_colored.figure_.set_figwidth(8)
      cm_colored.figure_.set_figheight(7)

      cm_colored.confusion_matrix

      # save figure as PNG
      plt.tight_layout()
      plt.savefig('images/confusion_matrix.png', dpi=150, pad_inches=5)
      plt.show()
```

```
[43]: from sklearn.model_selection import cross_val_score

      accuracies = cross_val_score(estimator = classifier, X = X_train,
                                   y = y_train, cv = 10)

      print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
      print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

Accuracy: 79.17 %

Standard Deviation: 6.72 %




Figure 11:

8 STEP 5: Vary parameters

This section was inspired by [In Depth: Parameter tuning for SVC](#)

In this section, the 4 SVC parameters `kernel`, `gamma`, `C` and `degree` will be introduced one by one. Furthermore, their influence on the classification result by varying these single parameters will be shown.

Disclaimer: In order to show the effects of varying the individual parameters in 2D graphs, only the best correlating variables `petal_length` and `petal_width` are used to train the SVC.

8.1 Prepare dataset

```
[44]: # import iris dataset again
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')

# encode the class column from class strings to integer equivalents
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
                                                    "Iris-virginica":2}})

irisdata_df_enc
```

```
[44]:   sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2         0
1           4.9           3.0           1.4           0.2         0
2           4.7           3.2           1.3           0.2         0
3           4.6           3.1           1.5           0.2         0
```

4	5.0	3.6	1.4	0.2	0
5	5.4	3.9	1.7	0.4	0
6	4.6	3.4	1.4	0.3	0
7	5.0	3.4	1.5	0.2	0
8	4.4	2.9	1.4	0.2	0
9	4.9	3.1	1.5	0.1	0
10	5.4	3.7	1.5	0.2	0
11	4.8	3.4	1.6	0.2	0
12	4.8	3.0	1.4	0.1	0
13	4.3	3.0	1.1	0.1	0
14	5.8	4.0	1.2	0.2	0
..
135	7.7	3.0	6.1	2.3	2
136	6.3	3.4	5.6	2.4	2
137	6.4	3.1	5.5	1.8	2
138	6.0	3.0	4.8	1.8	2
139	6.9	3.1	5.4	2.1	2
140	6.7	3.1	5.6	2.4	2
141	6.9	3.1	5.1	2.3	2
142	5.8	2.7	5.1	1.9	2
143	6.8	3.2	5.9	2.3	2
144	6.7	3.3	5.7	2.5	2
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

[150 rows x 5 columns]

```
[45]: # copy only 2 feature columns
# and convert pandas dataframe to numpy array
X = irisdata_df_enc[['petal_length', 'petal_width']].to_numpy(copy=True)
#X = irisdata_df_enc[['sepal_length', 'sepal_width']].to_numpy(copy=True)
#X
```

```
[46]: # convert pandas dataframe to numpy array
# and get a flat 1D copy of 2D numpy array
y = irisdata_df_enc[['species']].to_numpy(copy=True).flatten()
#y
```

8.2 Plotting function

This function helps to visualize the modifications by varying the individual SVC parameters.

```
[47]: def plotSVC(title, xlabel, ylabel):
# create a mesh to plot in
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

# prevent division by zero
if x_min == 0.0:
    x_min = 0.1

h = (x_max / x_min)/1000
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
```



```
plt.subplot(1, 1, 1)
Z = svc.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
plt.xlabel(xlabel)
plt.ylabel(ylabel)
plt.xlim(xx.min(), xx.max())
plt.title(title)
plt.show()
```

8.3 Vary kernel parameter

The kernel parameter selects the type of hyperplane that is used to separate the data. Using **linear** ([linear classifier](#)) kernel will use a linear hyperplane (a line in the case of 2D data). The **rbf** ([radial basis function kernel](#)) and **poly** ([polynomial kernel](#)) kernel use non linear hyperplanes.

```
[48]: kernels = ['linear', 'rbf', 'poly']

xlabel = 'Petal length'
ylabel = 'Petal width'

for kernel in kernels:
    svc = svm.SVC(kernel=kernel).fit(X, y)
    plotSVC('kernel = ' + str(kernel), xlabel, ylabel)
```




Figure 12:




Figure 13:




Figure 14:

8.4 Vary gamma parameter

The **gamma** parameter is used for non linear hyperplanes. The higher the **gamma** value it tries to exactly fit the training data set.

As we can see, increasing **gamma** leads to **overfitting** as the classifier tries to perfectly fit the training data.

```
[49]: gammas = [0.1, 1, 10, 100, 200]

xlabel = 'Petal length'
ylabel = 'Petal width'

for gamma in gammas:
    svc = svm.SVC(kernel='rbf', gamma=gamma).fit(X, y)
    plotSVC('gamma = ' + str(gamma), xlabel, ylabel)
```




Figure 15:

8.5 Vary C parameter

The **C** parameter is the **penalty** of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly.

But be careful: to high **C** values may lead to **overfitting** the training data.

```
[50]: cs = [0.1, 1, 10, 100, 1000, 10000]

xlabel = 'Petal length'
ylabel = 'Petal width'

for c in cs:
    svc = svm.SVC(kernel='rbf', C=c).fit(X, y)
```




Figure 16:




Figure 17:

```
plotSVC('C = ' + str(c), xlabel, ylabel)
```




Figure 18:




Figure 19:




Figure 20:




Figure 21:

8.6 Vary degree parameter

The **degree** parameter is used when the **kernel** is set to **poly**. It's basically the **degree of the polynomial** used to find the hyperplane to split the data.




Figure 22:




Figure 23:

Using `degree = 1` is the same as using a `linear` kernel. Also, increasing this parameters leads to **higher training times**.




Figure 24:




Figure 25:

```
[51]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```



```
xlabel = 'Petal length'
ylabel = 'Petal width'

for degree in degrees:
    svc = svm.SVC(kernel='poly', degree=degree).fit(X, y)
    plotSVC('degree = ' + str(degree), xlabel, ylabel)
```




Figure 26:




Figure 27:




Figure 28:




Figure 29:




Figure 30:




Figure 31:




Figure 32:




Figure 33:




Figure 34:




Figure 35: