


Getting started with ML and Support Vector Classifiers (SVC)

Björn Kasper (kasper.bjoern@bgetem.de)

August 2, 2022



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/) (CC BY-SA 4.0).

This is a test abstract.

Contents

1	Introduction	2
2	Load globally used libraries and set plot parameters	3
3	STEP 0: Get the data	3
4	STEP 1: Exploring the data	3
4.1	Goals of exploration	3
4.2	Clarify the origins history	4
4.3	Overview of the internal structure and organisation of the data	4
4.3.1	Inspect structure of dataframe	5
4.3.2	Get data types	8
4.3.3	Get data ranges with Boxplots	9
4.4	Identify anomalies in the data sets	10
4.4.1	Find gaps in dataset	10
4.4.2	Find and remove duplicates in dataset	14
4.5	Avoidance of tendencies due to bias	18
4.5.1	Count occurrences of unique values	18
4.5.2	Display Histogram	19
4.6	First idea of correlations in data set	20
4.6.1	Visualise data with correlation heatmap	20
4.6.2	Visualise data with scatter plot	22
4.6.3	Visualise data with pairs plot	23
5	STEP 2: Prepare the data	24
6	STEP 3: Classify by support vector classifier - SVC	24
6.1	Operating principal	24
6.2	Split the dataset	25
6.3	Create the SVM model	25
6.4	Make predictions	26
7	STEP 4: Evaluate the results - metrics	26
7.1	Textual confusion matrix	26
7.2	Colored confusion matrix	26
8	STEP 5: Vary parameters	27
8.1	Prepare dataset	27
8.2	Plotting function	28
8.3	Vary kernel parameter	29
8.4	Vary gamma parameter	31
8.5	Vary C parameter	34
8.6	Vary degree parameter	37

1 Introduction

This notebook was basically inspired by:

- [In Depth: Parameter tuning for SVC](#)
- [SVM Hyperparameter Tuning using GridSearchCV](#):

The goal of this notebook is to show the basic steps in machine learning and the influence of choosing the “right” the kernel of a **support vector classifier (SVC)**. Furthermore, the SVC parameters are described and their effect on the classification result is shown.

Following steps will be shown in next **chapters**:

- STEP 0: Get the data
- STEP 1: Exploring the data
- STEP 2: Prepare the data
- STEP 3: Classify by support vector classifier - SVC
- STEP 4: Evaluate the results - metrics
- STEP 5: Vary parameters

2 Load globally used libraries and set plot parameters

```
[40]: import time

from IPython.display import HTML

import pandas as pd
import matplotlib.pyplot as plt
from sklearn import svm, metrics
import seaborn as sns
%matplotlib inline
```

3 STEP 0: Get the data

Since this is intended to be an introduction to the world of machine learning (ML), this step does NOT deal with the design of an application suitable for ML and the acquisition of valid measurement data.

In order to get to know the typical work steps and ML tools, the use of **well-known and well-researched data sets** is clearly **recommended**.

In the further course, the famous [Iris flower data sets](#) will be used. It can be downloaded on [Iris Flower Dataset | Kaggle](#). Furthermore, the dataset is included in Python in the machine learning package [Scikit-learn](#), so that users can access it without having to find a special source for it.

```
[1]: # import some data to play with
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

4 STEP 1: Exploring the data

4.1 Goals of exploration

The objectives of the exploration of the dataset are as follows:

1. Clarify the **origins history**:
 - Where did the data come from? => Contact persons and licensing permissions?
 - Who obtained the data and with which (measurement) methods? => Did systematic errors occur during the acquisition?
 - What were they originally intended for? => Can they be used for my application?
2. Overview of the internal **structure and organisation** of the data:
 - Which columns are there? => With which methods can they be read in (e.g. import of CSV files)?
 - What do they contain for (physical) measured variables? => Which technical or physical correlations exist?
 - Which data formats or types are there? => Do they have to be converted?
 - In which value ranges do the measurement data vary? => Are normalizations necessary?
3. Identify **anomalies** in the data sets:
 - Do the data have **gaps** or **duplicates**? => Does the data set needs to be cleaned?

- Are there obvious erroneous entries or measurement outliers? => Does (statistical) filtering have to be carried out?
4. Avoidance of **tendencies due to bias**:
 - Are all possible classes included in the dataset and equally distributed? => Does the data set need to be enriched with additional data for balance?
 5. Find a first rough **idea of which correlations** could be in the data set

4.2 Clarify the origins history

The ***Iris* flower data sets** is a multivariate data set introduced by the British statistician and biologist *Ronald Fisher* in his paper “The use of multiple measurements in taxonomic problems as an example of linear discriminant analysis” (1936). It is sometimes called *Anderson’s Iris data set* because Edgar Anderson collected the data to quantify the morphologic variation of Iris flowers of three related species (source: [Iris flower data set](#)).

The dataset is published in Public Domain with a [CC0-License](#).

This dataset became a typical test case for many statistical classification techniques in machine learning such as **support vector machines**.

[..] measurements of the flowers of fifty plants each of the two species *Iris setosa* and *I. versicolor*, found **growing together in the same colony** and measured by Dr E. Anderson [..] (source: R. A. Fisher (1936). “The use of multiple measurements in taxonomic problems”. [Annals of Eugenics](#))

[..] *Iris virginica*, differs from the two other samples in **not being taken from the same natural colony** [..] (source: ibidem)

4.3 Overview of the internal structure and organisation of the data

The data set consists of 50 samples from each of three species of Iris (*Iris setosa*, *Iris virginica* and *Iris versicolor*), so there are 150 total samples. Four features were measured from each sample: the length and the width of the [sepals](#) and [petals](#), in centimetres.

Here is a principle illustration of a flower with sepal and petal:

```
[21]: display(HTML("<figure><img src='./images/Mature_flower_diagram.svg' width='800px'> \
    <figcaption>Principle illustration of a flower with sepal and
    ↪petal (source: <a href='https://en.wikipedia.org/wiki/File:Mature_flower_diagram.
    ↪svg'>Mature_flower_diagram.svg</a></figcaption> \
    </figure>"))
```

<IPython.core.display.HTML object>

Here are pictures of the three different Iris species (*Iris setosa*, *Iris virginica* and *Iris versicolor*). Given the dimensions of the flower, it will be possible to predict the class of the flower.

```
[22]: display(HTML("<table> \
    <tr> \
    <td><figure><img src='./images/Iris_setosa_640px.jpg' \
    ↪width='320px'> \
    <figcaption><i>Iris setosa</i> (source: <a href='https://
    ↪commons.wikimedia.org/wiki/File:Irissetosa1.jpg'>Irissetosa1.jpg</a></
    ↪figcaption> \
    </figure></td> \
    <td><figure><img src='./images/Iris_versicolor_640px.jpg' \
    ↪width='320px'> \
    <figcaption><i>Iris versicolor</i> (source: <a href='https:/
    ↪en.wikipedia.org/wiki/File:Iris_versicolor_3.jpg'>Iris versicolor 3.jpg</a></
    ↪figcaption> \
    </figure></td> \
```

```

<td><figure><img src='../images/Iris_virginica_590px.jpg'
width='295px'> \
    <figcaption><i>Iris virginica</i> (source: <a href='https://
en.wikipedia.org/wiki/File:Iris_virginica.jpg'>Iris virginica.jpg</a></
figcaption> \
    </figure></td> \
</tr> \
</table>"))

```

<IPython.core.display.HTML object>

4.3.1 Inspect structure of dataframe

Print first or last 5 rows of dataframe:

```
[3]: irisdata_df.head()
```

```

[3]:   sepal_length  sepal_width  petal_length  petal_width  species
0         5.1         3.5         1.4         0.2  Iris-setosa
1         4.9         3.0         1.4         0.2  Iris-setosa
2         4.7         3.2         1.3         0.2  Iris-setosa
3         4.6         3.1         1.5         0.2  Iris-setosa
4         5.0         3.6         1.4         0.2  Iris-setosa

```

```
[4]: irisdata_df.tail()
```

```

[4]:   sepal_length  sepal_width  petal_length  petal_width  species
145         6.7         3.0         5.2         2.3  Iris-virginica
146         6.3         2.5         5.0         1.9  Iris-virginica
147         6.5         3.0         5.2         2.0  Iris-virginica
148         6.2         3.4         5.4         2.3  Iris-virginica
149         5.9         3.0         5.1         1.8  Iris-virginica

```

While printing a dataframe - only an abbreviated view of the dataframe is shown :(
Default setting in the pandas library makes it to display only 5 lines from head and from tail.

```
[6]: irisdata_df
```

```

[6]:   sepal_length  sepal_width  petal_length  petal_width  species
0         5.1         3.5         1.4         0.2  Iris-setosa
1         4.9         3.0         1.4         0.2  Iris-setosa
2         4.7         3.2         1.3         0.2  Iris-setosa
3         4.6         3.1         1.5         0.2  Iris-setosa
4         5.0         3.6         1.4         0.2  Iris-setosa
..         ...         ...         ...         ...         ...
145         6.7         3.0         5.2         2.3  Iris-virginica
146         6.3         2.5         5.0         1.9  Iris-virginica
147         6.5         3.0         5.2         2.0  Iris-virginica
148         6.2         3.4         5.4         2.3  Iris-virginica
149         5.9         3.0         5.1         1.8  Iris-virginica

```

[150 rows x 5 columns]

To print all rows of a dataframe, the option `display.max_rows` has to set to `None` in pandas:

```

[7]: pd.set_option('display.max_rows', None)
irisdata_df

```

```
[7]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa
10	5.4	3.7	1.5	0.2	Iris-setosa
11	4.8	3.4	1.6	0.2	Iris-setosa
12	4.8	3.0	1.4	0.1	Iris-setosa
13	4.3	3.0	1.1	0.1	Iris-setosa
14	5.8	4.0	1.2	0.2	Iris-setosa
15	5.7	4.4	1.5	0.4	Iris-setosa
16	5.4	3.9	1.3	0.4	Iris-setosa
17	5.1	3.5	1.4	0.3	Iris-setosa
18	5.7	3.8	1.7	0.3	Iris-setosa
19	5.1	3.8	1.5	0.3	Iris-setosa
20	5.4	3.4	1.7	0.2	Iris-setosa
21	5.1	3.7	1.5	0.4	Iris-setosa
22	4.6	3.6	1.0	0.2	Iris-setosa
23	5.1	3.3	1.7	0.5	Iris-setosa
24	4.8	3.4	1.9	0.2	Iris-setosa
25	5.0	3.0	1.6	0.2	Iris-setosa
26	5.0	3.4	1.6	0.4	Iris-setosa
27	5.2	3.5	1.5	0.2	Iris-setosa
28	5.2	3.4	1.4	0.2	Iris-setosa
29	4.7	3.2	1.6	0.2	Iris-setosa
30	4.8	3.1	1.6	0.2	Iris-setosa
31	5.4	3.4	1.5	0.4	Iris-setosa
32	5.2	4.1	1.5	0.1	Iris-setosa
33	5.5	4.2	1.4	0.2	Iris-setosa
34	4.9	3.1	1.5	0.1	Iris-setosa
35	5.0	3.2	1.2	0.2	Iris-setosa
36	5.5	3.5	1.3	0.2	Iris-setosa
37	4.9	3.1	1.5	0.1	Iris-setosa
38	4.4	3.0	1.3	0.2	Iris-setosa
39	5.1	3.4	1.5	0.2	Iris-setosa
40	5.0	3.5	1.3	0.3	Iris-setosa
41	4.5	2.3	1.3	0.3	Iris-setosa
42	4.4	3.2	1.3	0.2	Iris-setosa
43	5.0	3.5	1.6	0.6	Iris-setosa
44	5.1	3.8	1.9	0.4	Iris-setosa
45	4.8	3.0	1.4	0.3	Iris-setosa
46	5.1	3.8	1.6	0.2	Iris-setosa
47	4.6	3.2	1.4	0.2	Iris-setosa
48	5.3	3.7	1.5	0.2	Iris-setosa
49	5.0	3.3	1.4	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
51	6.4	3.2	4.5	1.5	Iris-versicolor
52	6.9	3.1	4.9	1.5	Iris-versicolor
53	5.5	2.3	4.0	1.3	Iris-versicolor
54	6.5	2.8	4.6	1.5	Iris-versicolor
55	5.7	2.8	4.5	1.3	Iris-versicolor

56	6.3	3.3	4.7	1.6	Iris-versicolor
57	4.9	2.4	3.3	1.0	Iris-versicolor
58	6.6	2.9	4.6	1.3	Iris-versicolor
59	5.2	2.7	3.9	1.4	Iris-versicolor
60	5.0	2.0	3.5	1.0	Iris-versicolor
61	5.9	3.0	4.2	1.5	Iris-versicolor
62	6.0	2.2	4.0	1.0	Iris-versicolor
63	6.1	2.9	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
65	6.7	3.1	4.4	1.4	Iris-versicolor
66	5.6	3.0	4.5	1.5	Iris-versicolor
67	5.8	2.7	4.1	1.0	Iris-versicolor
68	6.2	2.2	4.5	1.5	Iris-versicolor
69	5.6	2.5	3.9	1.1	Iris-versicolor
70	5.9	3.2	4.8	1.8	Iris-versicolor
71	6.1	2.8	4.0	1.3	Iris-versicolor
72	6.3	2.5	4.9	1.5	Iris-versicolor
73	6.1	2.8	4.7	1.2	Iris-versicolor
74	6.4	2.9	4.3	1.3	Iris-versicolor
75	6.6	3.0	4.4	1.4	Iris-versicolor
76	6.8	2.8	4.8	1.4	Iris-versicolor
77	6.7	3.0	5.0	1.7	Iris-versicolor
78	6.0	2.9	4.5	1.5	Iris-versicolor
79	5.7	2.6	3.5	1.0	Iris-versicolor
80	5.5	2.4	3.8	1.1	Iris-versicolor
81	5.5	2.4	3.7	1.0	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
83	6.0	2.7	5.1	1.6	Iris-versicolor
84	5.4	3.0	4.5	1.5	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
90	5.5	2.6	4.4	1.2	Iris-versicolor
91	6.1	3.0	4.6	1.4	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor
93	5.0	2.3	3.3	1.0	Iris-versicolor
94	5.6	2.7	4.2	1.3	Iris-versicolor
95	5.7	3.0	4.2	1.2	Iris-versicolor
96	5.7	2.9	4.2	1.3	Iris-versicolor
97	6.2	2.9	4.3	1.3	Iris-versicolor
98	5.1	2.5	3.0	1.1	Iris-versicolor
99	5.7	2.8	4.1	1.3	Iris-versicolor
100	6.3	3.3	6.0	2.5	Iris-virginica
101	5.8	2.7	5.1	1.9	Iris-virginica
102	7.1	3.0	5.9	2.1	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
104	6.5	3.0	5.8	2.2	Iris-virginica
105	7.6	3.0	6.6	2.1	Iris-virginica
106	4.9	2.5	4.5	1.7	Iris-virginica
107	7.3	2.9	6.3	1.8	Iris-virginica
108	6.7	2.5	5.8	1.8	Iris-virginica
109	7.2	3.6	6.1	2.5	Iris-virginica
110	6.5	3.2	5.1	2.0	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica

113	5.7	2.5	5.0	2.0	Iris-virginica
114	5.8	2.8	5.1	2.4	Iris-virginica
115	6.4	3.2	5.3	2.3	Iris-virginica
116	6.5	3.0	5.5	1.8	Iris-virginica
117	7.7	3.8	6.7	2.2	Iris-virginica
118	7.7	2.6	6.9	2.3	Iris-virginica
119	6.0	2.2	5.0	1.5	Iris-virginica
120	6.9	3.2	5.7	2.3	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
122	7.7	2.8	6.7	2.0	Iris-virginica
123	6.3	2.7	4.9	1.8	Iris-virginica
124	6.7	3.3	5.7	2.1	Iris-virginica
125	7.2	3.2	6.0	1.8	Iris-virginica
126	6.2	2.8	4.8	1.8	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica
128	6.4	2.8	5.6	2.1	Iris-virginica
129	7.2	3.0	5.8	1.6	Iris-virginica
130	7.4	2.8	6.1	1.9	Iris-virginica
131	7.9	3.8	6.4	2.0	Iris-virginica
132	6.4	2.8	5.6	2.2	Iris-virginica
133	6.3	2.8	5.1	1.5	Iris-virginica
134	6.1	2.6	5.6	1.4	Iris-virginica
135	7.7	3.0	6.1	2.3	Iris-virginica
136	6.3	3.4	5.6	2.4	Iris-virginica
137	6.4	3.1	5.5	1.8	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
139	6.9	3.1	5.4	2.1	Iris-virginica
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

4.3.2 Get data types

```
[8]: irisdata_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 5.3+ KB
```

```
[9]: irisdata_df.describe()
```



```
[9]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

4.3.3 Get data ranges with Boxplots

Boxplots can be used to explore the data ranges in the data set. These also provide information about outliers.

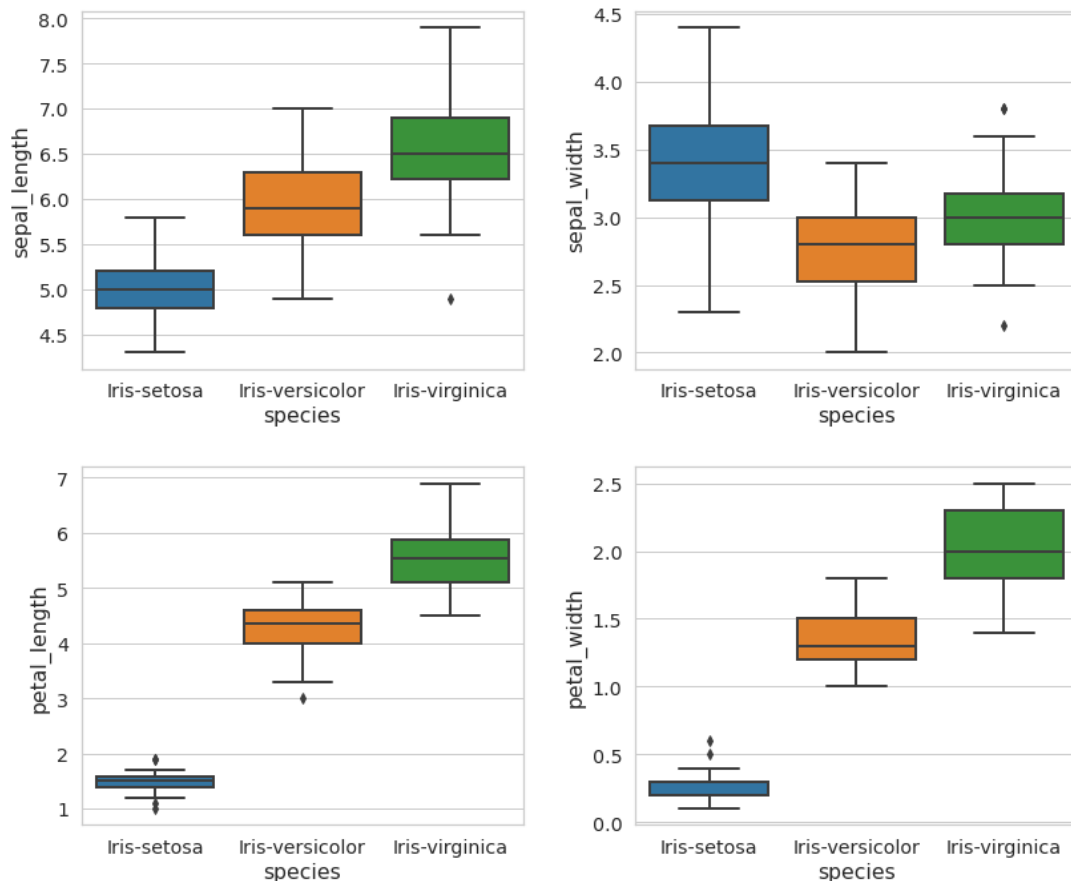
```
[157]: sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.0})
sns.set_style("whitegrid")
#sns.set_style("white")

fig, axs = plt.subplots(2, 2, figsize=(12, 10))

fn = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
box1 = sns.boxplot(x = 'species', y = 'sepal_length',
                   data = irisdata_df, order = cn, ax = axs[0,0])
box2 = sns.boxplot(x = 'species', y = 'sepal_width',
                   data = irisdata_df, order = cn, ax = axs[0,1])
box3 = sns.boxplot(x = 'species', y = 'petal_length',
                   data = irisdata_df, order = cn, ax = axs[1,0])
box4 = sns.boxplot(x = 'species', y = 'petal_width',
                   data = irisdata_df, order = cn, ax = axs[1,1])

# add some spacing between subplots
fig.tight_layout(pad=2.0)

plt.show()
```



4.4 Identify anomalies in the data sets

4.4.1 Find gaps in dataset

This section was inspired by [Working with Missing Data in Pandas](#).

Checking for missing values using `isnull()` In order to check for missing values in Pandas DataFrame, we use the function `isnull()`. This function returns a dataframe of Boolean values which are True for NaN values.

```
[37]: pd.set_option('display.max_rows', 40)
      pd.set_option('display.min_rows', 30)
```

```
[38]: irisdata_df.isnull()
```

```
[38]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
0	False	False	False	False	False
1	False	False	False	False	False
2	False	False	False	False	False
3	False	False	False	False	False
4	False	False	False	False	False
5	False	False	False	False	False
6	False	False	False	False	False
7	False	False	False	False	False
8	False	False	False	False	False
9	False	False	False	False	False

```

10      False      False      False      False      False
11      False      False      False      False      False
12      False      False      False      False      False
13      False      False      False      False      False
14      False      False      False      False      False
..      ...      ...      ...      ...      ...
135     False      False      False      False      False
136     False      False      False      False      False
137     False      False      False      False      False
138     False      False      False      False      False
139     False      False      False      False      False
140     False      False      False      False      False
141     False      False      False      False      False
142     False      False      False      False      False
143     False      False      False      False      False
144     False      False      False      False      False
145     False      False      False      False      False
146     False      False      False      False      False
147     False      False      False      False      False
148     False      False      False      False      False
149     False      False      False      False      False

```

[150 rows x 5 columns]

Show only the gaps:

```
[5]: irisdata_df_gaps = irisdata_df[irisdata_df.isnull().any(axis=1)]
irisdata_df_gaps
```

```
[5]: Empty DataFrame
Columns: [sepal_length, sepal_width, petal_length, petal_width, species]
Index: []
```

Fine - this dataset seems to be complete :)

So let's look for something else for exercise: [employees.csv](#)

```
[159]: # import data to dataframe from csv file
employees_df = pd.read_csv("./datasets/employees_edit.csv")

employees_df
```

```
[159]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
...	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

	Senior Management	Team
0	True	Marketing
1	True	NaN

```

2          False          Finance
3           True          Finance
4           True  Client Services
...
999        False      Distribution
1000       False          Finance
1001       False          Product
1002       False  Business Development
1003        True           Sales

```

[1004 rows x 8 columns]

Show only the gaps from this gappy dataset again:

```
[160]: employees_df_gaps = employees_df[employees_df.isnull().any(axis=1)]
employees_df_gaps
```

```
[160]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
20	Lois	NaN	4/22/1995	7:18 PM	64714	4934.00	
22	Joshua	NaN	3/8/2012	1:58 AM	90816	18816.00	
..	
965	Antonio	NaN	6/18/1989	9:37 PM	103050	3.05	
976	Victor	NaN	7/28/2006	2:49 PM	76381	11159.00	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	

	Senior Management	Team
1	True	NaN
7	NaN	Finance
10	True	NaN
20	True	Legal
22	True	Client Services
..
965	False	Legal
976	True	Sales
989	False	Legal
993	False	Legal
999	False	Distribution

[237 rows x 8 columns]

Fill the missing values with fillna() Now we are going to fill all the null (NaN) values in Gender column with “No Gender”.

Attention: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

```
[161]: # filling a null values using fillna()
employees_df["Gender"].fillna("No Gender", inplace = True)
employees_df
```

```
[161]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	

```

1      Thomas      Male  3/31/1996      6:53 AM  61933      4.17
2      Maria      Female 4/23/1993      11:17 AM  130590  11858.00
3      Jerry      Male  3/4/2005      1:00 PM  138705      9.34
4      Larry      Male  1/24/1998      4:47 PM  101004  1389.00
...
999    Henry      No Gender 11/23/2014      6:09 AM  132483  16655.00
1000   Phillip    Male  1/31/1984      6:30 AM  42392  19675.00
1001   Russell    Male  5/20/2013      12:39 PM  96914  1421.00
1002   Larry      Male  4/20/2013      4:45 PM  60500  11985.00
1003   Albert     Male  5/15/2012      6:24 PM  129949  10169.00

```

```

      Senior Management      Team
0              True      Marketing
1              True      NaN
2             False      Finance
3              True      Finance
4              True  Client Services
...
999             False      Distribution
1000             False      Finance
1001             False      Product
1002             False  Business Development
1003             True      Sales

```

[1004 rows x 8 columns]

Dropping missing values using dropna() In order to drop null values from a dataframe, we use dropna() function. This function drops rows or columns of datasets with NaN values in different ways.

Default is to drop rows with at least 1 null value (NaN). Giving the parameter how = 'all' the function drops rows with all data missing or contain null values (NaN).

```

[162]: # making a new dataframe with dropped NaN values
employees_df_dropped = employees_df.dropna(axis = 0, how = 'any')
employees_df_dropped

```

```

[162]:      First Name      Gender  Start Date  Last Login Time  Salary  Bonus % \
0      Douglas      Male  8/6/1993      12:42 PM  97308  6945.00
2      Maria      Female 4/23/1993      11:17 AM  130590  11858.00
3      Jerry      Male  3/4/2005      1:00 PM  138705      9.34
4      Larry      Male  1/24/1998      4:47 PM  101004  1389.00
5      Dennis      Male  4/18/1987      1:35 AM  115163  10125.00
...
999    Henry      No Gender 11/23/2014      6:09 AM  132483  16655.00
1000   Phillip    Male  1/31/1984      6:30 AM  42392  19675.00
1001   Russell    Male  5/20/2013      12:39 PM  96914  1421.00
1002   Larry      Male  4/20/2013      4:45 PM  60500  11985.00
1003   Albert     Male  5/15/2012      6:24 PM  129949  10169.00

```

```

      Senior Management      Team
0              True      Marketing
2             False      Finance
3              True      Finance
4              True  Client Services
5             False      Legal
...
999             False      Distribution

```

1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[903 rows x 8 columns]

Finally we compare the sizes of dataframes so that we learn how many rows had at least 1 Null value.

```
[163]: print("Old data frame length:", len(employees_df))
print("New data frame length:", len(employees_df_dropped))
print("Number of rows with at least 1 NaN value: ",
      (len(employees_df)-len(employees_df_dropped)))
```

Old data frame length: 1004

New data frame length: 903

Number of rows with at least 1 NaN value: 101

4.4.2 Find and remove duplicates in dataset

This section was inspired by: - [How to Find Duplicates in Pandas DataFrame \(With Examples\)](#) - [How to Drop Duplicate Rows in a Pandas DataFrame](#)

Checking for duplicate values using duplicated() In order to check for duplicate values in Pandas DataFrame, we use a function `duplicated()`. This function can be used in two ways: - find duplicate rows across **all columns** with `duplicateRows = df[df.duplicated()]` - find duplicate rows across **specific columns** `duplicateRows = df[df.duplicated(subset=['col1', 'col2'])]`

Find duplicate rows across **all columns**:

```
[44]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[45]: # find duplicate rows across all columns
duplicateRows = employees_df[employees_df.duplicated()]
duplicateRows
```

```
[45]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
112	Karen	Female	11/30/1999	7:46 AM	102488	17653.0	
127	Linda	Female	5/25/2000	5:45 PM	119009	12506.0	
296	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.0	
580	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.0	

	Senior Management	Team
112	True	Product
127	True	Business Development
296	False	Business Development
580	True	Human Resources

```
[46]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(keep='last')]
duplicateRows
```

```
[46]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
55	Karen	Female	11/30/1999	7:46 AM	102488	17653.0	
92	Linda	Female	5/25/2000	5:45 PM	119009	12506.0	
153	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.0	

442	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.0
	Senior Management				Team	
55		True			Product	
92		True	Business Development			
153		False	Business Development			
442		True	Human Resources			

Find duplicate rows across **specific columns**:

```
[164]: # identify duplicate rows across 'First Name' and 'Last Login Time' columns
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'])]
duplicateRows
```

```
[164]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
112	Karen	Female	11/30/1999	7:46 AM	102488	17653.0	
127	Linda	Female	5/25/2000	5:45 PM	119009	12506.0	
296	Brandon	No Gender	11/3/1997	8:17 PM	121333	15295.0	
577	NaN	Female	1/13/2009	1:01 PM	118736	7421.0	
580	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.0	
632	NaN	No Gender	9/2/1988	12:49 PM	147309	1702.0	
881	NaN	Male	9/5/1980	7:36 AM	114896	13823.0	
929	NaN	Female	8/23/2000	4:19 PM	95866	19388.0	
934	Nancy	Female	9/10/2001	11:57 PM	85213	2386.0	
973	Linda	Female	2/4/2010	8:49 PM	44486	17308.0	

	Senior Management				Team
112		True			Product
127		True	Business Development		
296		False	Business Development		
577		NaN	Client Services		
580		True	Human Resources		
632		NaN	Distribution		
881		NaN	Client Services		
929		NaN	Sales		
934		True	Marketing		
973		True	Engineering		

```
[165]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'], keep='last')]
duplicateRows
```

```
[165]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
23	NaN	Male	6/14/2012	4:19 PM	125792	5042.00	
37	Linda	Female	10/19/1981	8:49 PM	57427	9557.00	
55	Karen	Female	11/30/1999	7:46 AM	102488	17653.00	
66	Nancy	Female	12/15/2012	11:57 PM	125250	2672.00	
92	Linda	Female	5/25/2000	5:45 PM	119009	12506.00	
153	Brandon	No Gender	11/3/1997	8:17 PM	121333	15295.00	
222	NaN	Female	6/17/1991	12:49 PM	71945	5.56	
269	NaN	Male	2/4/2005	1:01 PM	40451	16044.00	
442	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.00	
778	NaN	Female	6/18/2000	7:36 AM	106428	10867.00	

	Senior Management				Team
--	-------------------	--	--	--	------

23	NaN	NaN
37	True	Client Services
55	True	Product
66	True	Business Development
92	True	Business Development
153	False	Business Development
222	NaN	Marketing
269	NaN	Distribution
442	True	Human Resources
778	NaN	NaN

Dropping duplicate values using drop_duplicates() In order to drop duplicate values from a dataframe, we use drop_duplicates() function.

This function can be used in two ways: - remove duplicate rows across **all columns** with df.drop_duplicates() - find duplicate rows across **specific columns** df.drop_duplicates(subset=['col1', 'col2'])

Attention: We are doing that directly in this dataframe with inplace = True - we don't make a deep copy!

Remove duplicate rows across **all columns**:

```
[49]: # remove duplicate rows across all columns
employees_df.drop_duplicates(inplace=True)
employees_df
```

```
[49]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	
998	George	Male	6/21/2013	5:47 PM	98874	4479.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	

1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00
------	--------	------	-----------	---------	--------	----------

	Senior Management		Team
0	True		Marketing
1	True		NaN
2	False		Finance
3	True		Finance
4	True	Client Services	
5	False		Legal
6	True		Product
7	NaN		Finance
8	True	Engineering	
9	True	Business Development	
10	True		NaN
11	True		Legal
12	True	Human Resources	
13	False		Sales
14	True		Finance
...
989	False		Legal
990	False		Marketing
991	True		Finance
992	False	Human Resources	
993	False		Legal
994	True	Client Services	
995	True		Marketing
996	True		Finance
997	True	Engineering	
998	True		Marketing
999	False	Distribution	
1000	False		Finance
1001	False		Product
1002	False	Business Development	
1003	True		Sales

[1000 rows x 8 columns]

Remove duplicate rows across **specific columns**:

```
[166]: # remove duplicate rows across 'First Name' and 'Last Login Time' columns
employees_df.drop_duplicates(
    subset=['First Name', 'Last Login Time'], keep='last', inplace=True)
employees_df
```

```
[166]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
...	
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

Senior Management	Team
-------------------	------

```

0          True          Marketing
1          True           NaN
2         False          Finance
3          True          Finance
4          True    Client Services
...
999         False    Distribution
1000        False          Finance
1001        False          Product
1002        False    Business Development
1003         True           Sales

```

```
[994 rows x 8 columns]
```

4.5 Avoidance of tendencies due to bias

The description of the Iris dataset says, that it consists of **50 samples** from **each of three species** of Iris (Iris setosa, Iris virginica and Iris versicolor), so there are **150 total samples**.

But how to prove it?

4.5.1 Count occurrences of unique values

To prove whether all possible classes included in the dataset and equally distributed, you can use the function `df.value_counts`.

Following parameters can be used for fine tuning: - `dropna=False` causes that NaN values are included - `normalize=True`: relative frequencies of the unique values are returned - `ascending=False`: sort resulting classes descending

```
[167]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[168]: # count unique values without missing values in a column,
# ordered descending and normalized
irisdata_df['species'].value_counts(ascending=False, dropna=False, normalize=True)
```

```
[168]: Iris-setosa      0.333333
Iris-versicolor    0.333333
Iris-virginica     0.333333
Name: species, dtype: float64
```

```
[169]: # count unique values and missing values in a column,
# ordered descending and not absolute values
employees_df['Team'].value_counts(ascending=False, dropna=False, normalize=False)
```

```
[169]: Client Services      106
Business Development    103
Finance                 102
Marketing                98
Product                 96
Sales                   94
Engineering              92
Human Resources          92
Distribution             90
Legal                   88
NaN                     43
```

Name: Team, dtype: int64


4.5.2 Display Histogram

This section was inspired by: [Pandas Histogram – DataFrame.hist\(\)](#).


Histograms represent **frequency distributions** graphically. This requires the separation of the data into classes (so-called **bins**).

These classes are represented in the histogram as rectangles of equal or variable width. The height of each rectangle then represents the (relative or absolute) **frequency density**.

```
[176]: employees_df.hist(column=['Salary'])  
plt.show()
```



```
[175]: employees_df.hist(column='Salary', by='Gender')  
plt.show()
```



4.6 First idea of correlations in data set

To get a rough idea of the **dependencies** and **correlations** in the data set, it can be helpful to visualize the whole dataset in a **correlation heatmap**. They show in a glance which variables are correlated, to what degree and in which direction.

Later, 2 particularly well correlated variables are selected from the data set and plotted in a **scatterplot**.

4.6.1 Visualise data with correlation heatmap

This section was inspired by [How to Create a Seaborn Correlation Heatmap in Python?](#).

Correlation matrices are an **essential tool of exploratory data analysis**. Correlation heatmaps contain the same information in a visually appealing way. What more: they show in a glance which variables are correlated, to what degree, in which direction, and alerts us to potential multicollinearity problems (source: ibidem).

Simple correlation matrix Because **string values can never be correlated**, the class names (species) have to be converted first:

```
[170]: # encoding the class column
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
                                                    "Iris-virginica":2}})

irisdata_df_enc
```

```
[170]:   sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2         0
```

```

1      4.9      3.0      1.4      0.2      0
2      4.7      3.2      1.3      0.2      0
3      4.6      3.1      1.5      0.2      0
4      5.0      3.6      1.4      0.2      0
..      ...      ...      ...      ...      ...
145    6.7      3.0      5.2      2.3      2
146    6.3      2.5      5.0      1.9      2
147    6.5      3.0      5.2      2.0      2
148    6.2      3.4      5.4      2.3      2
149    5.9      3.0      5.1      1.8      2

```

[150 rows x 5 columns]

```
[92]: irisdata_df_enc.corr()
```

```

[92]:      sepal_length  sepal_width  petal_length  petal_width  species
sepal_length    1.000000   -0.109369    0.871754    0.817954   0.782561
sepal_width     -0.109369    1.000000   -0.420516   -0.356544  -0.419446
petal_length     0.871754   -0.420516    1.000000    0.962757   0.949043
petal_width      0.817954   -0.356544    0.962757    1.000000   0.956464
species          0.782561   -0.419446    0.949043    0.956464   1.000000

```

Correlation heatmap Choose the color sets from [color map](#).


```

[174]: # increase the size of the heatmap
plt.figure(figsize=(16, 6))

# store heatmap object in a variable to easily access it
# when you want to include more features (such as title)
# set the range of values to be displayed on the colormap from -1 to 1,
# and set 'annotation=True' to display the correlation values on the heatmap
heatmap = sns.heatmap(irisdata_df_enc.corr(), vmin=-1, vmax=1,
                      annot=True, cmap='PRGn_r')

# give a title to the heatmap
# 'pad=12' defines the distance of the title from the top of the heatmap
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
plt.show()

```



Triangle correlation heatmap When looking at the correlation heatmaps above, you would not lose any information by **cutting** away half of it **along the diagonal** line marked by 1-s.

The **numpy** function `np.triu()` can be used to isolate the upper triangle of a matrix while turning all the values in the lower triangle into 0.

```
[94]: import numpy as np

np.triu(np.ones_like(irisdata_df_enc.corr()))
```

```
[94]: array([[1., 1., 1., 1., 1.],
            [0., 1., 1., 1., 1.],
            [0., 0., 1., 1., 1.],
            [0., 0., 0., 1., 1.],
            [0., 0., 0., 0., 1.]])
```


Use this mask to cut the heatmap along the diagonal:

```
[173]: plt.figure(figsize=(16, 6))

# define the mask to set the values in the upper triangle to 'True'
mask = np.triu(np.ones_like(irisdata_df_enc.corr(), dtype=bool))

heatmap = sns.heatmap(irisdata_df_enc.corr(), mask=mask,
                      vmin=-1, vmax=1, annot=True, cmap='PRGn_r')

heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
plt.show()
```



As a result from the **heatmaps** we can see, that the shape of the **petals** are the **most correlated columns** (0.96) with the **type of flowers** (species classes).

Somewhat lower correlates **sepal length** with **petal length** (0.87).

4.6.2 Visualise data with scatter plot


In the following, **Seaborn** is applied which is a library for making statistical graphics in Python. It is built on top of matplotlib and closely integrated with pandas data structures.

To investigate whether there are dependencies (e.g. correlations) in `irisdata_df` between individual variables in the data set, it is advisable to plot them in a **scatter plot**.

```
[181]: # There are five preset seaborn themes: darkgrid, whitegrid, dark, white, and ticks.
sns.set_style("whitegrid")
# set scale of fonts
sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.5})

# 'sepal_length', 'petal_length' are iris feature data
# 'height' used to define height of graph
# 'hue' stores the class/label of iris dataset
sns.FacetGrid(irisdata_df, hue="species",
               height=7).map(plt.scatter,
                             'petal_width',
                             'petal_length').add_legend()

plt.title('Scatterplot of petal length and width')
plt.show()
```



4.6.3 Visualise data with pairs plot

For systematic investigation of dependencies, all variables (each against each) are plotted in separate scatter plots.

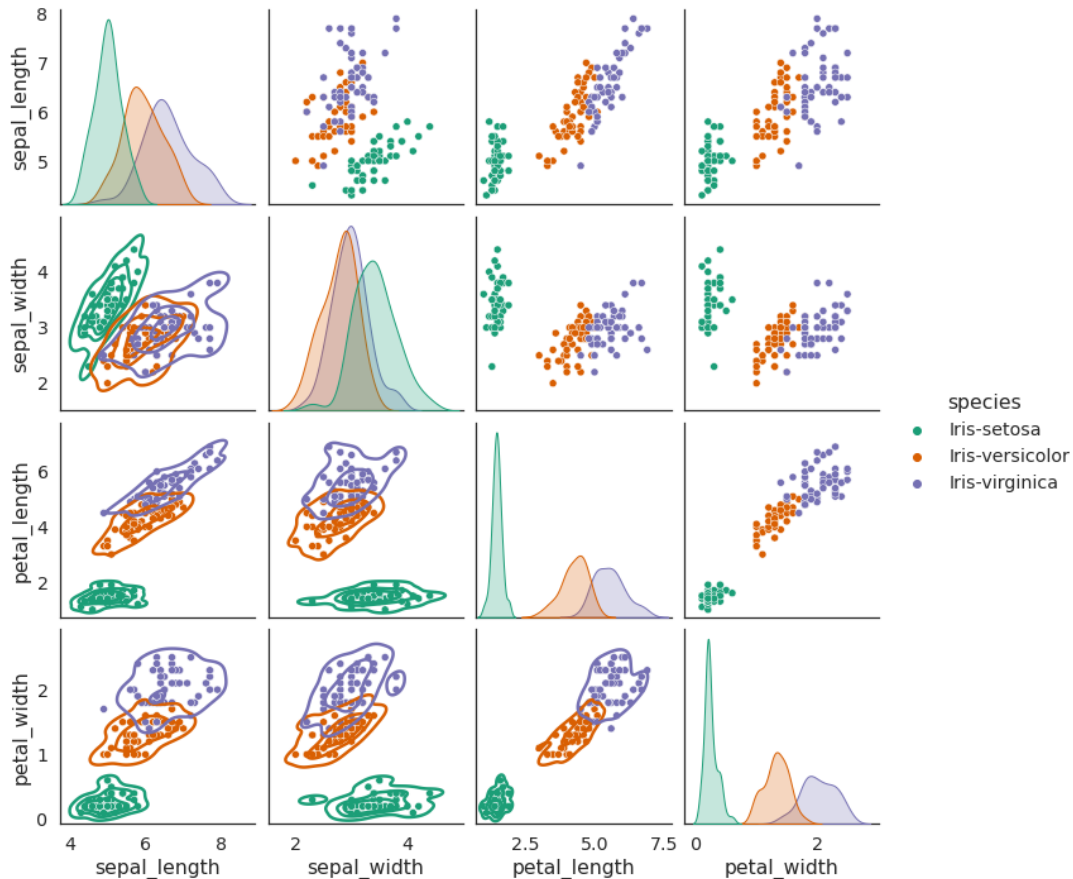
With this so called **pairs plot** it is possible to see both **relationships** between two variables and **distribution** of single variables.

This function will create a grid of Axes such that **each numeric variable** in `irisdata_df` will be shared in the y-axis across a single row and in the x-axis across a single column.

```
[184]: sns.set_style("white")
g = sns.pairplot(irisdata_df, diag_kind="kde", hue='species',
                 palette='Dark2', height=2.5)

g.map_lower(sns.kdeplot, levels=4, color=".2")

plt.show()
```



5 STEP 2: Prepare the data

Through the intensive exploration of the data in Step 1 ([STEP 1: Exploring the data](#)), we know that special **preparation** of the data is **not necessary**. The values are **complete** and **without gaps** and there are **no duplicates**. The values are in similar ranges, which **does not require normalization** of the data.

Furthermore, we know that the **classes** are very **evenly distributed** and thus bias tendencies should be avoided.

6 STEP 3: Classify by support vector classifier - SVC

6.1 Operating principal

Support Vectors Classifier tries to **find the best hyperplane to separate** the different classes by maximizing the distance between sample points and the hyperplane (source: [In Depth: Parameter tuning for SVC](#)).

Following graphic shows the operating principal of SVC: the hyperplane $H1$ does not separate the classes. $H2$ does, but only with a small margin. $H3$ separates them with the maximal margin (source: [Support-vector machine](#)).

```
[7]: display(HTML("<figure><img src='../images/SVM_separating_hyperplanes.svg'
width='400px'> \
        <figcaption>SVC separate the data in classes by finding the best
hyperplane (source: <a href='https://en.wikipedia.org/wiki/File:
Svm_separating_hyperplanes_(SVG).svg'>Svm separating hyperplanes (SVG).svg</a></
figcaption> \
        </figure>"))
```

<IPython.core.display.HTML object>

6.2 Split the dataset

In the next very important step, the dataset is split into **2 subsets**: a **training dataset** and a **test dataset**. As the names suggest, the training dataset is used to train the ML algorithm. The test data set is then used to check the quality of the trained ML algorithm (here the **recognition rate**). For this purpose, the **class labels** are **removed** from the training data set - after all, these are to be predicted.

Typically, the **test dataset** should contain about **20%** of the entire dataset.

```
[43]: from sklearn.model_selection import train_test_split

X = irisdata_df.drop('species', axis=1)
y = irisdata_df['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

For training, do not use only the variables that correlate best with each other, but all of them.

Otherwise, the result of the prediction would be significantly worse. Maybe this is already an indication of **overfitting** of the ML model.

```
[42]: # DO NOT USE THIS!!
X_train, X_test, y_train, y_test = train_test_split(X[['sepal_length',
                                                    'sepal_width']],
                                                    y, test_size = 0.20)
```

6.3 Create the SVM model

In this step we create the SVC model and fit it to our training data.

```
[44]: from sklearn.svm import SVC
classifier = SVC(kernel = 'linear', random_state = 0)

# fit the model for the data
classifier.fit(X_train, y_train)
```

```
[44]: SVC(kernel='linear', random_state=0)
```

6.4 Make predictions

```
[45]: y_pred = classifier.predict(X_test)
      #X_test
```

7 STEP 4: Evaluate the results - metrics

And finally for checking the accuracy of the model, the **confusion matrix** is used for the **cross validation**.

By using the function `sklearn.metrics.confusion_matrix()` a confusion matrix of the true digit values versus the predicted digit values is plotted.

7.1 Textual confusion matrix

```
[46]: cm = metrics.confusion_matrix(y_test, y_pred)
      print(cm)
```

```
[[ 8  0  0]
 [ 0  7  0]
 [ 0  0 15]]
```

7.2 Colored confusion matrix


The function `sklearn.metrics.ConfusionMatrixDisplay()` plots a colored confusion matrix.

```
[58]: sns.set_style("white")

      # print colored confusion matrix
      cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

      #cm_colored.figure_.suptitle("Confusion Matrix")
      cm_colored.figure_.set_figwidth(7)
      cm_colored.figure_.set_figheight(6)

      cm_colored.confusion_matrix
      plt.show()
```



```
[185]: from sklearn.model_selection import cross_val_score

accuracies = cross_val_score(estimator = classifier, X = X_train,
                              y = y_train, cv = 10)

print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

Accuracy: 96.67 %
Standard Deviation: 5.53 %

8 STEP 5: Vary parameters

This section was inspired by [In Depth: Parameter tuning for SVC](#)

In this section, the 4 SVC parameters `kernel`, `gamma`, `C` and `degree` will be introduced one by one. Furthermore, their influence on the classification result by varying these single parameters will be shown.

Disclaimer: In order to show the effects of varying the individual parameters in 2D graphs, only the best correlating variables `petal_length` and `petal_width` are used to train the SVC.

8.1 Prepare dataset

```
[186]: # import iris dataset again
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')

# encode the class column from class strings to integer equivalents
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
```

```
irisdata_df_enc                                "Iris-virginica":2}})
```

```
[186]:      sepal_length  sepal_width  petal_length  petal_width  species
0           5.1         3.5         1.4         0.2         0
1           4.9         3.0         1.4         0.2         0
2           4.7         3.2         1.3         0.2         0
3           4.6         3.1         1.5         0.2         0
4           5.0         3.6         1.4         0.2         0
..          ...          ...          ...          ...          ...
145          6.7         3.0         5.2         2.3         2
146          6.3         2.5         5.0         1.9         2
147          6.5         3.0         5.2         2.0         2
148          6.2         3.4         5.4         2.3         2
149          5.9         3.0         5.1         1.8         2
```

```
[150 rows x 5 columns]
```

```
[116]: # copy only 2 feature columns
# and convert pandas dataframe to numpy array
X = irisdata_df_enc[['petal_length', 'petal_width']].to_numpy(copy=True)
#X = irisdata_df_enc[['sepal_length', 'sepal_width']].to_numpy(copy=True)
#X
```

```
[118]: # convert pandas dataframe to numpy array
# and get a flat 1D copy of 2D numpy array
y = irisdata_df_enc[['species']].to_numpy(copy=True).flatten()
#y
```

8.2 Plotting function

This function helps to visualize the modifications by varying the individual SVC parameters.

```
[101]: def plotSVC(title, xlabel, ylabel):
# create a mesh to plot in
x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

# prevent division by zero
if x_min == 0.0:
    x_min = 0.1

h = (x_max / x_min)/1000
xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

plt.subplot(1, 1, 1)
Z = svc.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
plt.xlabel(xlabel)
plt.ylabel(ylabel)
plt.xlim(xx.min(), xx.max())
plt.title(title)
plt.show()
```


8.3 Vary kernel parameter



The `kernel` parameter selects the type of hyperplane that is used to separate the data. Using `linear` ([linear classifier](#)) kernel will use a linear hyperplane (a line in the case of 2D data). The `rbf` ([radial basis function kernel](#)) and `poly` ([polynomial kernel](#)) kernel use non linear hyperplanes.

```
[102]: kernels = ['linear', 'rbf', 'poly']

xlabel = 'Petal length'
ylabel = 'Petal width'

for kernel in kernels:
    svc = svm.SVC(kernel=kernel).fit(X, y)
    plotSVC('kernel = ' + str(kernel), xlabel, ylabel)
```





8.4 Vary gamma parameter


The **gamma** parameter is used for non linear hyperplanes. The higher the **gamma** value it tries to exactly fit the training data set.


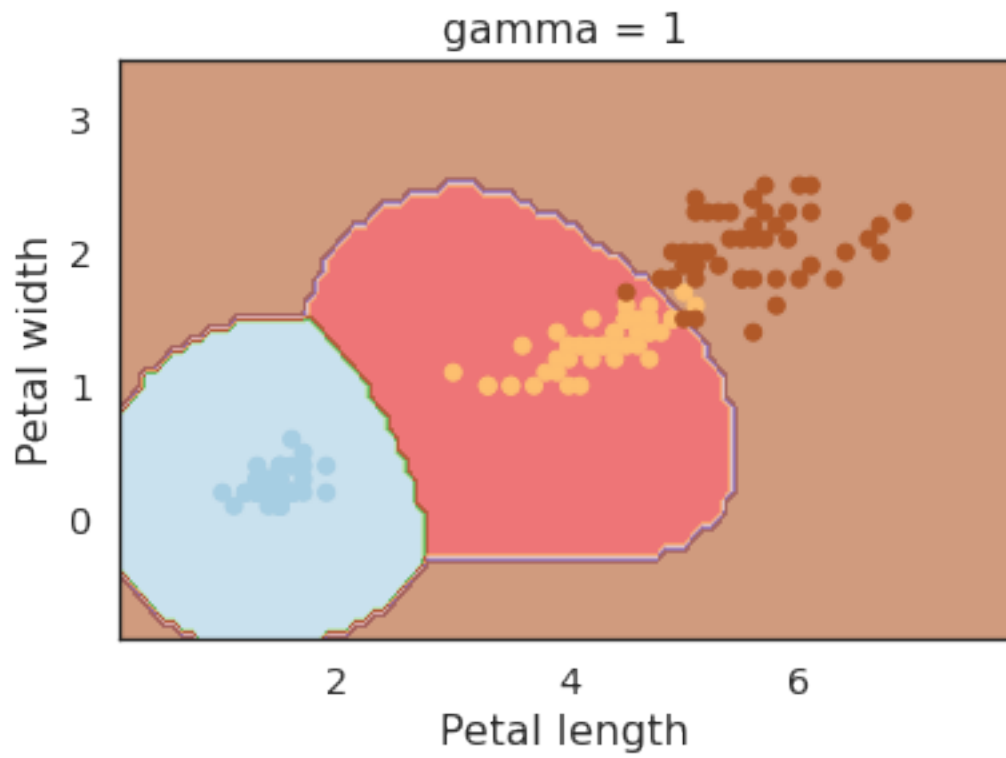
As we can see, increasing **gamma** leads to **overfitting** as the classifier tries to perfectly fit the training data.


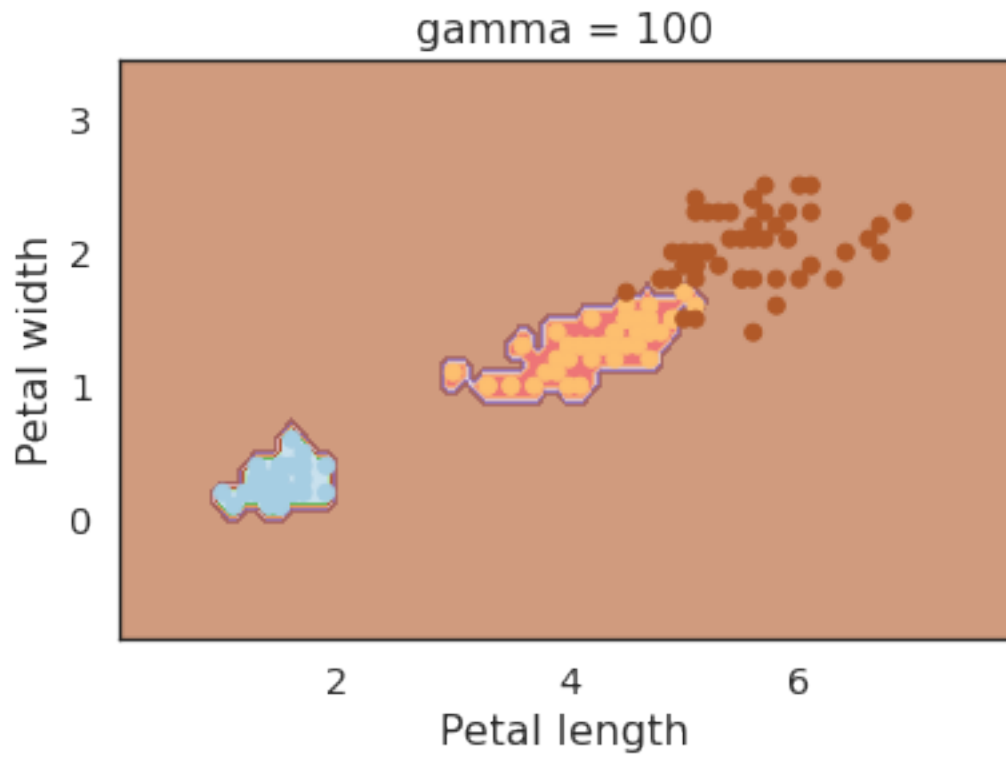
```
[106]: gammas = [0.1, 1, 10, 100, 200]

xlabel = 'Petal length'
ylabel = 'Petal width'

for gamma in gammas:
    svc = svm.SVC(kernel='rbf', gamma=gamma).fit(X, y)
    plotSVC('gamma = ' + str(gamma), xlabel, ylabel)
```







8.5 Vary C parameter


The **C** parameter is the **penalty** of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly.



But be careful: to high **C** values may lead to **overfitting** the training data.



```
[108]: cs = [0.1, 1, 10, 100, 1000, 10000]


xlabel = 'Petal length'
ylabel = 'Petal width'

for c in cs:
    svc = svm.SVC(kernel='rbf', C=c).fit(X, y)
    plotSVC('C = ' + str(c), xlabel, ylabel)
```









8.6 Vary degree parameter


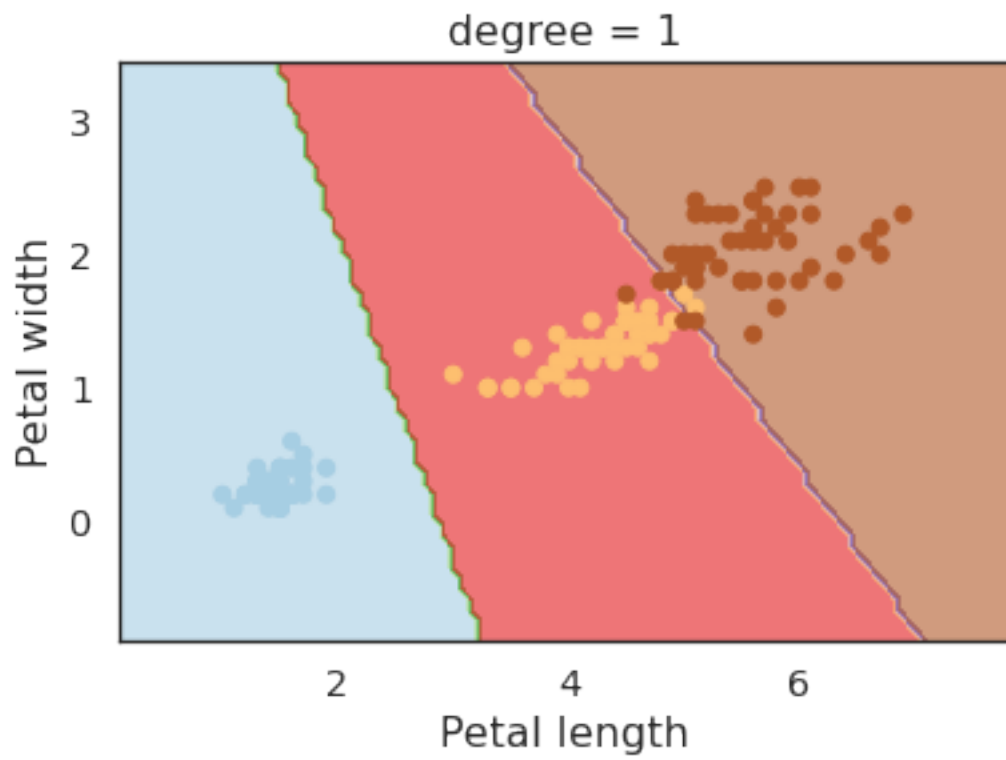
The `degree` parameter is used when the `kernel` is set to `poly`. It's basically the **degree of the polynomial** used to find the hyperplane to split the data.



Using `degree = 1` is the same as using a `linear` kernel. Also, increasing this parameters leads to **higher training times**.

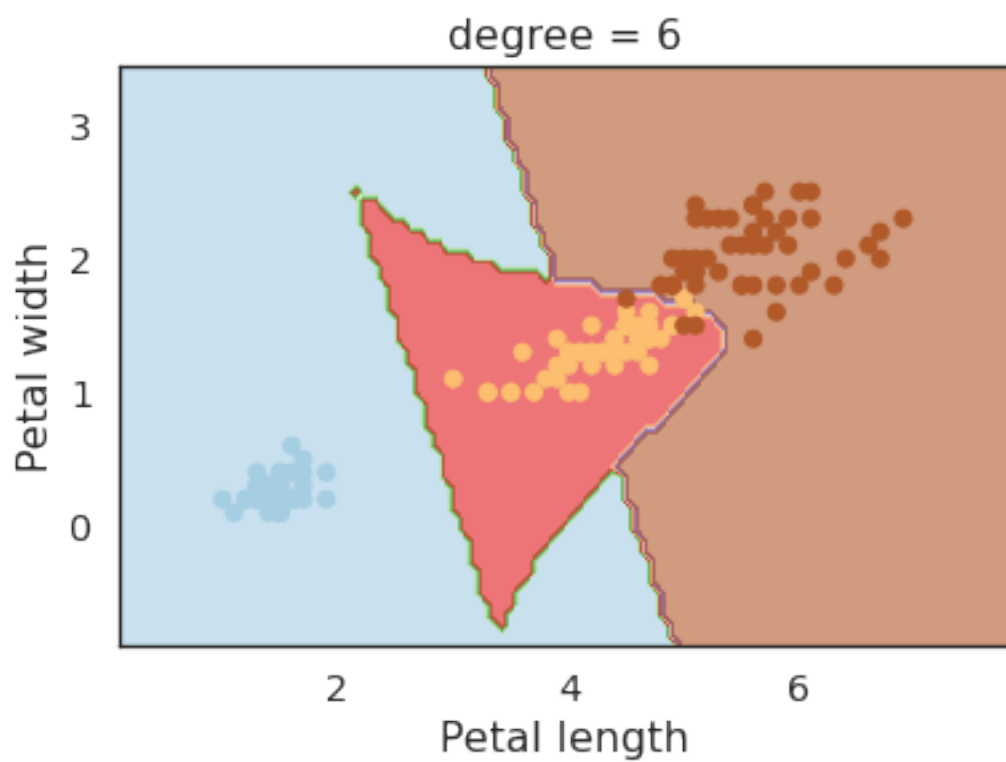

```
[113]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]



xlabel = 'Petal length'
ylabel = 'Petal width'

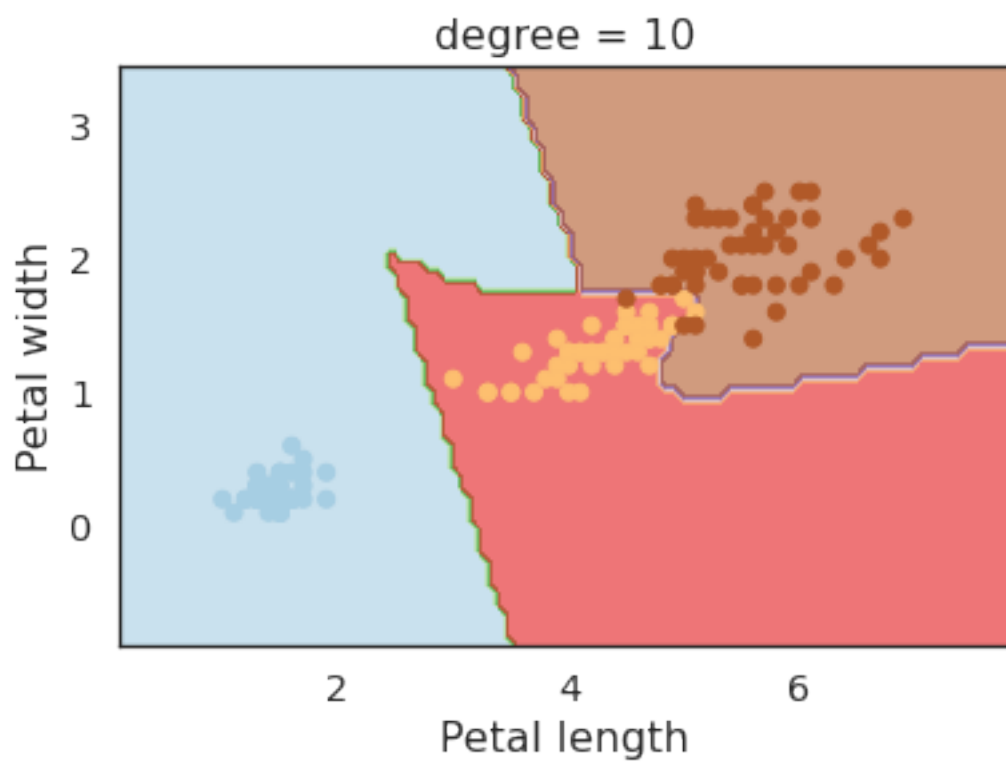

for degree in degrees:
    svc = svm.SVC(kernel='poly', degree=degree).fit(X, y)
    plotSVC('degree = ' + str(degree), xlabel, ylabel)
```











```
[ ]:
```