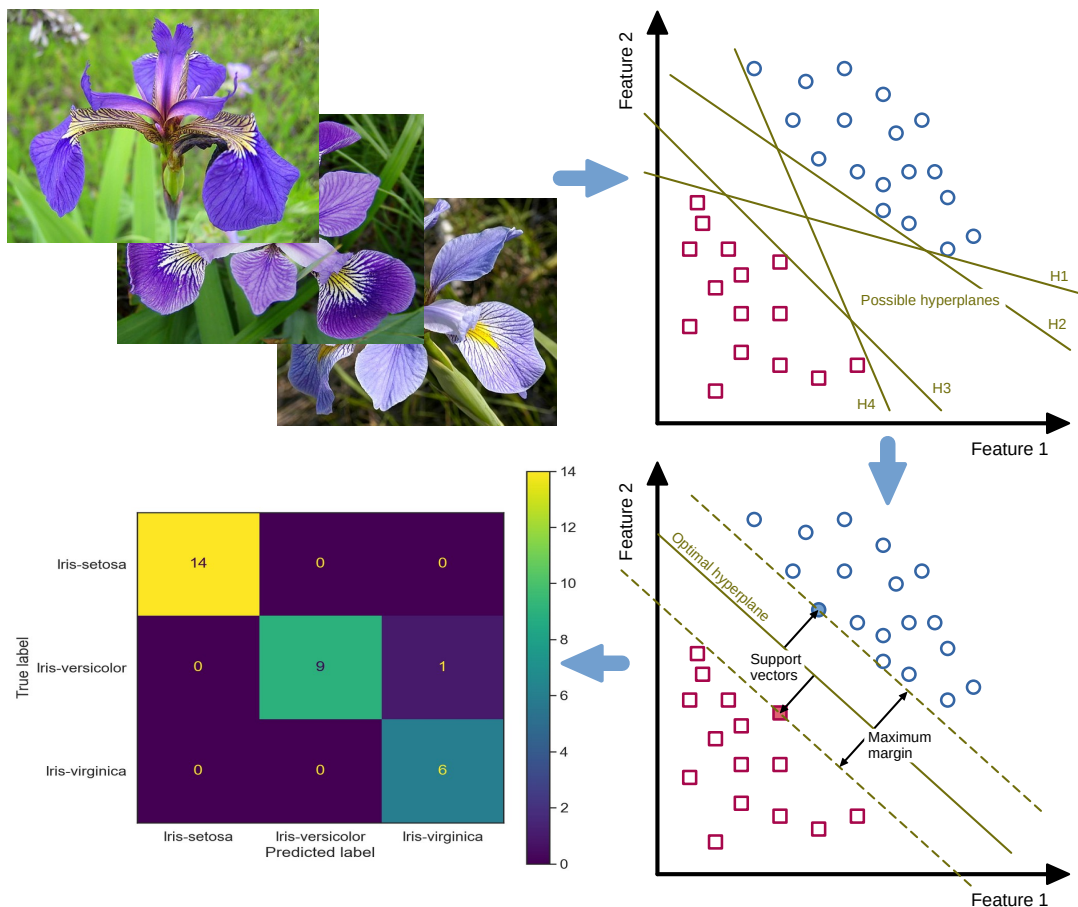


Getting started with Machine Learning (ML) and Support Vector Classifiers (SVC) - A systematic step-by-step approach

Dipl.-Ing. Björn Kasper (kasper.bjoern@bgetem.de)

Test and Certification Body for Electrical Engineering at BG ETEM

August 7, 2022



Anyone who wants to seriously deal with the hypothetical topic of our time “Artificial Intelligence (AI)” or “Machine Learning (ML)” cannot avoid dealing with the basic ML algorithms, corresponding software tools, libraries and programming systems. However, someone who opens the door for the first time to this equally very exciting as well as arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. Here, it is a good idea to consult introductory and systematic tutorials. Therefore, this Getting Started tutorial systematically demonstrates the typical ML work process step-by-step using the very powerful and performant “Support Vector Classifier (SVC)” and the widely known and very beginner-friendly “Iris Dataset”. Furthermore, the selection of the “correct” SVC kernel and its parameters are described and their effect on the classification result is shown.



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/) (CC BY-SA 4.0).

Contents

1	Introduction	2
1.1	English introduction	2
1.2	German introduction	4
1.3	Steps of the systematic ML process	5
2	Load globally used libraries and set plot parameters	5
3	STEP 0: Get the dataset	6
4	STEP 1: Exploring the dataset	6
4.1	Goals of exploration	6
4.2	Clarify the origins history	6
4.3	Overview of the internal structure and organisation of the data	7
4.3.1	Inspect structure of dataframe	7
4.3.2	Get data types	11
4.3.3	Get data ranges with Boxplots	11
4.4	Identify anomalies in the data sets	13
4.4.1	Find gaps in dataset	13
4.4.2	Find and remove duplicates in dataset	19
4.5	Avoidance of tendencies due to bias	24
4.5.1	Count occurrences of unique values	24
4.5.2	Display Histogram	24
4.6	First idea of correlations in data set	25
4.6.1	Visualise data with correlation heatmap	25
4.6.2	Visualise data with scatter plot	27
4.6.3	Visualise data with pairs plot	28
5	STEP 2: Prepare the dataset	29
6	STEP 3: Classify by support vector classifier - SVC	30
6.1	Operating principal	30
6.2	Split the dataset	30
6.3	Create the SVM model	31
6.4	Make predictions	31
7	STEP 4: Evaluate the classification results - metrics	31
7.1	Textual confusion matrix	31
7.2	Colored confusion matrix	31
7.3	Classification accuracy	32
8	STEP 5: Select SVC kernel and vary parameters	32
8.1	Prepare dataset	32
8.1.1	Prepare datasets for parameter variation and plotting	33
8.1.2	Prepare dataset for prediction and evaluation	33
8.2	Plotting functions	33
8.3	Vary kernel of SVC	34
8.4	Vary gamma parameter	36
8.5	Vary C parameter	38
8.6	Vary degree parameter	40

1 Introduction

1.1 English introduction

In the **digitised work environment**, there is an increasing demand for **Work equipment** to be able to adapt independently and in a task-related manner to changing work situations. This **situational adaptivity** can often only be realised through the use of **Artificial Intelligence (AI)** or **Machine**

Learning (ML), depending on the degree of flexibility. Examples of such AI applications in the world of work can range from comparatively simple **voice assistance systems** (similar, for example, to Siri or Alexa from the private sphere) to partially or even **fully autonomous systems**. Such fully autonomous systems are, for example, autonomously driving logistics vehicles in larger industrial plants (so-called **driverless transport systems**).

In addition to the many very interesting advantages in terms of economic efficiency, workload reduction, etc., such fully autonomous systems are characterised by a very high level of technical complexity. This concerns both their **operating functions** (e.g. autonomous navigation through complex industrial environments with shared use of the roadways by other human-controlled vehicles) and their **safety functions** (e.g. evaluation of complex, interconnected, mostly imaging safety sensors for monitoring the driving space).

Very high demands are placed on such autonomous systems and the AI algorithms used for them with regard to **functional safety**. However, when assessing their safety, one quickly comes up against clear limits with regard to the **transparency** and **explainability** of the decisions made by AI as well as limits to the **recognition rates** and thus their **reliability**. In particular, the detection rates achievable by AI even under the most convenient conditions very often do not meet the requirements for realising higher safety levels (e.g. Performance Level d (PLd) according to ISO 13849).

An appropriate assessment or even **testing** with regard to the required functional safety according to uniform and ideally standardised criteria has many implications for the future orientation of technical **occupational safety and health (OSH)** in Germany and in Europe. In addition to the currently still very difficult algorithmic evaluability, a significant aspect is that the previous clear separation between **placing on the market law** (see e.g. Machinery Directive) and **occupational health and safety law** (see European Occupational Health and Safety Framework Directive and German Ordinance on Occupational Safety and Health) can no longer be continued in this way. The reason for this is that the **safety-relevant properties** of the autonomous systems will change due to new or **adapted behaviours** learned during operation.

For these reasons, those involved in technical occupational safety and health who will be involved in the testing of work equipment in the future should deal with AI and ML algorithms in depth as early as possible. This is the only way to ensure that the rapid development of adaptive systems capable of learning can be accompanied by OSH and its testing institutes in a constructive, critical and technically appropriate manner. If this is not done, the OSH system will be ruthlessly circumvented or undermined by the economic interests of globally operating software giants. This would have the consequence that serious or fatal occupational accidents are likely to occur due to inadequately designed AI-based work systems.

Anyone seeking a serious technical entrance into the world of **Artificial Intelligence (AI)** or **Machine Learning (ML)** will not be able to avoid dealing with the basic ML algorithms, corresponding software tools, libraries and programming systems.

However, someone who opens the door for the first time to this equally very exciting as well as arbitrarily complex and, at first glance, confusing world will very quickly be overwhelmed. Here, it is a good idea to consult introductory and systematic tutorials.

The aim of this Getting Started tutorial is to systematically demonstrate the typical ML working process step-by-step based on the example of the very powerful and performant **Support Vector Classifier (SVC)**.

This tutorial will be presented as part of a workshop at the DGUV symposium **Artificial Intelligence**, probably in November 2022 in Dresden. The workshop addresses interested ML novices in the technical occupational safety and health of the social accident insurance institutions.

For the target audience in the workshop, the SVC algorithm was intentionally chosen to show that there are many other very powerful and performant ML algorithms apart from the **deep neural networks** that are very present in the media. On the other hand, a necessary and comprehensible introduction to neural networks and the technical background to perceptrons, activation functions etc. for newcomers would not be possible within the time frame given for the workshop.

Furthermore, this tutorial does *not* address the generation or acquisition of ML-ready datasets. Reason for this is that a newcomer to ML will (or should) first try to familiarize himself with ML algorithms,

tools, libraries and programming systems. Only then it makes sense to explore one's own environment with respect to ML-suitable applications and to acquire suitable datasets from them.

Therefore, this tutorial demonstrates the usage of selected ML tools in the form of Python libraries as well as the systematic approach to the widely known and very beginner-friendly **Iris dataset**. According to the literature, the Support Vector Classifier is particularly well suited for the classification of the iris dataset in terms of recognition rate and performance. Alternatively, decision tree-based ML algorithms such as the **Random Forests Classifier** could be used.

After the classification of the iris dataset by the SVC initially with standard parameters, the selection of the "correct" SVC kernel with its setting parameters is furthermore described and the effect on the classification result is shown.

1.2 German introduction

Von den **Arbeitsmitteln** in der **digitalisierten Arbeitswelt** wird immer stärker gefordert, dass sie sich selbstständig und aufgabenbezogen an sich ändernde Arbeitssituationen anpassen können. Diese **situative Adaptivität** kann je nach Stärke des Flexibilisierungsgrades oft nur durch Anwendung von **Artificial Intelligence (AI)** oder **Machine Learning (ML)** realisiert werden.

Als Beispiele für solche KI-Anwendungen in der Arbeitswelt können vergleichsweise einfache **Sprachassistentensysteme** (ähnlich z. B. Siri oder Alexa aus dem privaten Umfeld) bis hin zu teil- oder gar **vollautonomen Systemen** genannt werden. Solche vollautonomen Systemen sind beispielsweise autonom fahrende Logistikfahrzeuge in größeren Industrieanlagen (sog. **fahrerlosen Transportsystemen**).

Neben den vielen sehr interessanten Vorteilen bzgl. Wirtschaftlichkeit, Arbeitserleichterung usw. kennzeichnet solche vollautonomen Systeme eine sehr hohe technische Komplexität. Diese betrifft sowohl ihre **Betriebsfunktionen** (z. B. autonome Navigation durch komplexe industrielle Umgebungen bei gemeinsamer Nutzung der Fahrwege durch andere menschlich gesteuerte Fahrzeuge) als auch seiner **Sicherheitsfunktionen** (z. B. Auswertung komplexer, miteinander verknüpfter, meist bildgebender Sicherheitssensorik zur Überwachung des Fahrtraums).

An solche autonomen Systeme und die hierfür eingesetzten KI-Algorithmen werden sehr hohe Anforderungen hinsichtlich der **funktionalen Sicherheit** gestellt. Jedoch stößt man bei ihrer sicherheitstechnischen Bewertung heute noch sehr schnell an deutliche Grenzen hinsichtlich der **Transparenz** und **Erklärbarkeit** der durch KI getroffenen Entscheidungen sowie Grenzen der **Erkennungsraten** und damit ihrer **Zuverlässigkeit**. Insbesondere erfüllen die durch KI selbst unter günstigsten Bedingungen erreichbaren Erkennungsraten sehr oft nicht die Anforderungen, um höhere Safety-Level (z. B. Performance Level d (PLd) nach ISO 13849) zu realisieren.

Eine hinsichtlich der geforderten funktionalen Sicherheit angemessene Bewertung oder gar **Prüfung** nach einheitlichen und idealerweise genormten Maßstäben hat viele Implikationen auf die zukünftige Ausrichtung des **technischen Arbeitsschutzes** in Deutschland und in Europa. Neben der derzeit noch sehr schwierigen algorithmischen Bewertbarkeit ist ein wesentlicher Aspekt, dass die bisherige klare Trennung zwischen **Inverkehrbringensrecht** (siehe z. B. Maschinenrichtlinie) und **betrieblichem Arbeitsschutzrecht** (siehe Arbeitsschutzrahmenrichtlinie und Betriebssicherheitsverordnung) so nicht mehr aufrechterhalten werden kann. Grund hierfür ist, dass sich die **sicherheitsrelevanten Eigenschaften** der autonomen Systeme durch während des Betriebs erlernte, neue oder **angepasste Verhaltensweisen** verändern werden.

Aus diesen Gründen sollten sich insbesondere die zukünftig mit der Prüfung befassten Akteure des technischen Arbeitsschutzes möglichst frühzeitig mit den KI- bzw. ML-Algorithmen vertieft auseinandersetzen. Nur dadurch lässt sich erreichen, dass die stürmische Entwicklung lernfähiger, adaptiver Systeme durch den Arbeitsschutz und deren Prüfinstitute konstruktiv, kritisch und fachlich angemessen begleitet werden kann. Wird dies versäumt, wird das Arbeitsschutzsystem durch die wirtschaftlichen Interessen global agierender Softwaregiganten skrupellos umgangen oder ausgehebelt werden. Dies hätte die Folge, dass schwere oder tödliche Arbeitsunfälle auf Grund unzulänglich gestalteter KI-basierter Arbeitssysteme wahrscheinlich werden.

Wer einen ernsthaften fachlichen Einstieg in die Welt von **Künstlicher Intelligenz (KI)** bzw. **Machine Learning (ML)** sucht, wird nicht umhin kommen, sich mit den grundlegenden ML-Algorithmen, entsprechenden Software-Werkzeugen, Bibliotheken und Programmiersystemen auseinander zu setzen.

Wer jedoch zum ersten Mal die Tür zu dieser ebenso spannenden wie beliebig komplexen und auf den ersten Blick verwirrenden Welt öffnet, wird sehr schnell überfordert sein. Hier empfiehlt es sich, einführende und systematische Anleitungen zu Rate zu ziehen.

Ziel dieses Getting-Started-Tutorials ist es, den typischen ML-Arbeitsablauf systematisch und Schritt-für-Schritt am Beispiel des sehr leistungsfähigen **Support Vector Classifier (SVC)** zu demonstrieren.

Dieses Tutorial wird im Rahmen eines Workshops auf der DGUV-Fachtagung **Künstliche Intelligenz** voraussichtlich im November 2022 in Dresden vorgestellt. Der Workshop richtet sich an interessierte ML-Neulinge im technischen Arbeitsschutz der gesetzlichen Unfallversicherungsträger.

Für die Zielgruppe des Workshops wurde der SVC-Algorithmus bewusst gewählt, um zu zeigen, dass es neben den **tiefen neuronalen Netzen**, die in den Medien sehr präsent sind, noch viele andere sehr leistungsfähige ML-Algorithmen gibt. Andererseits wäre eine notwendige und verständliche Einführung in neuronale Netze und die technischen Hintergründe zu Perzeptronen, Aktivierungsfunktionen etc. für Neulinge in dem für den Workshop vorgegebenen Zeitrahmen nicht möglich gewesen.

Außerdem befasst sich dieses Tutorial *nicht* mit der Erzeugung oder Akquisition von ML-tauglichen Datensätzen. Der Grund dafür ist, dass ein ML-Neuling zunächst versuchen wird (oder sollte), sich mit den ML-Algorithmen, Werkzeugen, Bibliotheken und Programmiersystemen vertraut zu machen. Erst dann ist es sinnvoll, die eigene Umgebung auf ML-taugliche Anwendungen hin zu untersuchen und daraus geeignete Datensätze zu gewinnen.

Daher demonstriert dieses Tutorial die Verwendung ausgewählter ML-Tools in Form von Python-Bibliotheken sowie die systematische Herangehensweise an den weithin bekannten und sehr einsteigerfreundlichen **Iris-Datensatz**. Laut Fachliteratur ist für die Klassifikation des Iris-Datensatzes der Support Vector Classifier hinsichtlich Erkennungsrate als auch Performanz besonders gut geeignet. Alternativ könnten auch entscheidungsbaum-basierte ML-Algorithmen wie z. B. der **Random-forests-Klassifikator** eingesetzt werden.

Nach der Klassifikation des Iris-Datensatzes durch den SVC zunächst mit Standard-Parametern wird darüber hinaus die Auswahl des "richtigen" SVC-Kernels mit seinen Einstellparametern beschrieben und die Auswirkung auf das Klassifikationsergebnis wird gezeigt.

1.3 Steps of the systematic ML process

The following steps of the systematic ML process are covered in the next main sections:

- STEP 0: Get the dataset
- STEP 1: Exploring the dataset
- STEP 2: Prepare the dataset
- STEP 3: Classify by support vector classifier - SVC
- STEP 4: Evaluate the classification results - metrics
- STEP 5: Select SVC kernel and vary parameters

2 Load globally used libraries and set plot parameters

```
[1]: import time

from IPython.display import HTML

import pandas as pd
import matplotlib.pyplot as plt
from sklearn import svm, metrics
import seaborn as sns
%matplotlib inline
```

3 STEP 0: Get the dataset

Since this is intended to be an introduction to the world of machine learning (ML), this step does NOT deal with the design of an application suitable for ML and the acquisition of valid measurement data.

In order to get to know the typical work steps and ML tools, the use of **well-known and well-researched data sets** is clearly **recommended**.

In the further course, the famous [Iris flower data sets](#) will be used. It can be downloaded on [Iris Flower Dataset | Kaggle](#). Furthermore, the dataset is included in Python in the machine learning package [Scikit-learn](#), so that users can access it without having to find a special source for it.

```
[3]: # import some data to play with
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

4 STEP 1: Exploring the dataset

4.1 Goals of exploration

The objectives of the exploration of the dataset are as follows:

1. Clarify the **origins history**:
 - Where did the data come from? => Contact persons and licensing permissions?
 - Who obtained the data and with which (measurement) methods? => Did systematic errors occur during the acquisition?
 - What were they originally intended for? => Can they be used for my application?
2. Overview of the internal **structure and organisation** of the data:
 - Which columns are there? => With which methods can they be read in (e.g. import of CSV files)?
 - What do they contain for (physical) measured variables? => Which technical or physical correlations exist?
 - Which data formats or types are there? => Do they have to be converted?
 - In which value ranges do the measurement data vary? => Are normalizations necessary?
3. Identify **anomalies** in the data sets:
 - Do the data have **gaps or duplicates**? => Does the data set needs to be cleaned?
 - Are there obvious erroneous entries or measurement outliers? => Does (statistical) filtering have to be carried out?
4. Avoidance of **tendencies due to bias**:
 - Are all possible classes included in the dataset and equally distributed? => Does the data set need to be enriched with additional data for balance?
5. Find a first rough **idea of which correlations** could be in the data set

4.2 Clarify the origins history

The **Iris flower data sets** is a multivariate data set introduced by the British statistician and biologist *Ronald Fisher* in his paper “The use of multiple measurements in taxonomic problems as an example of linear discriminant analysis” (1936). It is sometimes called *Anderson’s Iris data set* because Edgar Anderson collected the data to quantify the morphologic variation of Iris flowers of three related species (source: [Iris flower data set](#)).

The dataset is published in Public Domain with a [CC0-License](#).

This dataset became a typical test case for many statistical classification techniques in machine learning such as **support vector machines**.

[..] measurements of the flowers of fifty plants each of the two species *Iris setosa* and *I. versicolor*, found **growing together in the same colony** and measured by Dr E. Anderson [..] (source: R. A. Fisher (1936). “The use of multiple measurements in taxonomic problems”. [Annals of Eugenics](#))

[..] *Iris virginica*, differs from the two other samples in **not being taken from the same natural colony** [..] (source: ibidem)

4.3 Overview of the internal structure and organisation of the data

The data set consists of 50 samples from each of three species of Iris (*Iris setosa*, *Iris virginica* and *Iris versicolor*), so there are 150 total samples. Four features were measured from each sample: the length and the width of the **sepals** and **petals**, in centimetres.

Here is a principle illustration of a flower with sepal and petal:




Figure 1: Principle illustration of a flower with sepal and petal (source: [Mature_flower_diagram.svg](#), license: public domain)

Here are pictures of the three different Iris species (*Iris setosa*, *Iris virginica* and *Iris versicolor*). Given the dimensions of the flower, it will be possible to predict the class of the flower.



Figure 2: Left: *Iris setosa* (source: [Irissetosa1.jpg](#), license: public domain); middle: *Iris versicolor* (source: [Iris_versicolor_3.jpg](#), license: CC-SA 3.0); right: *Iris virginica* (source: [Iris_virginica.jpg](#), license: CC-SA 2.0)

4.3.1 Inspect structure of dataframe

Print first or last 5 rows of dataframe:

```
[4]: irisdata_df.head()
```



```
[4]:      sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2  Iris-setosa
1           4.9           3.0           1.4           0.2  Iris-setosa
2           4.7           3.2           1.3           0.2  Iris-setosa
3           4.6           3.1           1.5           0.2  Iris-setosa
4           5.0           3.6           1.4           0.2  Iris-setosa
```

```
[5]: irisdata_df.tail()
```

```
[5]:      sepal_length  sepal_width  petal_length  petal_width  species
145           6.7           3.0           5.2           2.3  Iris-virginica
146           6.3           2.5           5.0           1.9  Iris-virginica
147           6.5           3.0           5.2           2.0  Iris-virginica
148           6.2           3.4           5.4           2.3  Iris-virginica
149           5.9           3.0           5.1           1.8  Iris-virginica
```

While printing a dataframe - only an abbreviated view of the dataframe is shown :(
Default setting in the pandas library makes it to display only 5 lines from head and from tail.

```
[6]: irisdata_df
```

```
[6]:      sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2  Iris-setosa
1           4.9           3.0           1.4           0.2  Iris-setosa
2           4.7           3.2           1.3           0.2  Iris-setosa
3           4.6           3.1           1.5           0.2  Iris-setosa
4           5.0           3.6           1.4           0.2  Iris-setosa
..          ...           ...           ...           ...      ...
145          6.7           3.0           5.2           2.3  Iris-virginica
146          6.3           2.5           5.0           1.9  Iris-virginica
147          6.5           3.0           5.2           2.0  Iris-virginica
148          6.2           3.4           5.4           2.3  Iris-virginica
149          5.9           3.0           5.1           1.8  Iris-virginica
```

[150 rows x 5 columns]

To print all rows of a dataframe, the option `display.max_rows` has to set to `None` in pandas:

```
[7]: pd.set_option('display.max_rows', None)
irisdata_df
```

```
[7]:      sepal_length  sepal_width  petal_length  petal_width  species
0           5.1           3.5           1.4           0.2  Iris-setosa
1           4.9           3.0           1.4           0.2  Iris-setosa
2           4.7           3.2           1.3           0.2  Iris-setosa
3           4.6           3.1           1.5           0.2  Iris-setosa
4           5.0           3.6           1.4           0.2  Iris-setosa
5           5.4           3.9           1.7           0.4  Iris-setosa
6           4.6           3.4           1.4           0.3  Iris-setosa
7           5.0           3.4           1.5           0.2  Iris-setosa
8           4.4           2.9           1.4           0.2  Iris-setosa
9           4.9           3.1           1.5           0.1  Iris-setosa
10          5.4           3.7           1.5           0.2  Iris-setosa
11          4.8           3.4           1.6           0.2  Iris-setosa
12          4.8           3.0           1.4           0.1  Iris-setosa
13          4.3           3.0           1.1           0.1  Iris-setosa
14          5.8           4.0           1.2           0.2  Iris-setosa
15          5.7           4.4           1.5           0.4  Iris-setosa
```

16	5.4	3.9	1.3	0.4	Iris-setosa
17	5.1	3.5	1.4	0.3	Iris-setosa
18	5.7	3.8	1.7	0.3	Iris-setosa
19	5.1	3.8	1.5	0.3	Iris-setosa
20	5.4	3.4	1.7	0.2	Iris-setosa
21	5.1	3.7	1.5	0.4	Iris-setosa
22	4.6	3.6	1.0	0.2	Iris-setosa
23	5.1	3.3	1.7	0.5	Iris-setosa
24	4.8	3.4	1.9	0.2	Iris-setosa
25	5.0	3.0	1.6	0.2	Iris-setosa
26	5.0	3.4	1.6	0.4	Iris-setosa
27	5.2	3.5	1.5	0.2	Iris-setosa
28	5.2	3.4	1.4	0.2	Iris-setosa
29	4.7	3.2	1.6	0.2	Iris-setosa
30	4.8	3.1	1.6	0.2	Iris-setosa
31	5.4	3.4	1.5	0.4	Iris-setosa
32	5.2	4.1	1.5	0.1	Iris-setosa
33	5.5	4.2	1.4	0.2	Iris-setosa
34	4.9	3.1	1.5	0.1	Iris-setosa
35	5.0	3.2	1.2	0.2	Iris-setosa
36	5.5	3.5	1.3	0.2	Iris-setosa
37	4.9	3.1	1.5	0.1	Iris-setosa
38	4.4	3.0	1.3	0.2	Iris-setosa
39	5.1	3.4	1.5	0.2	Iris-setosa
40	5.0	3.5	1.3	0.3	Iris-setosa
41	4.5	2.3	1.3	0.3	Iris-setosa
42	4.4	3.2	1.3	0.2	Iris-setosa
43	5.0	3.5	1.6	0.6	Iris-setosa
44	5.1	3.8	1.9	0.4	Iris-setosa
45	4.8	3.0	1.4	0.3	Iris-setosa
46	5.1	3.8	1.6	0.2	Iris-setosa
47	4.6	3.2	1.4	0.2	Iris-setosa
48	5.3	3.7	1.5	0.2	Iris-setosa
49	5.0	3.3	1.4	0.2	Iris-setosa
50	7.0	3.2	4.7	1.4	Iris-versicolor
51	6.4	3.2	4.5	1.5	Iris-versicolor
52	6.9	3.1	4.9	1.5	Iris-versicolor
53	5.5	2.3	4.0	1.3	Iris-versicolor
54	6.5	2.8	4.6	1.5	Iris-versicolor
55	5.7	2.8	4.5	1.3	Iris-versicolor
56	6.3	3.3	4.7	1.6	Iris-versicolor
57	4.9	2.4	3.3	1.0	Iris-versicolor
58	6.6	2.9	4.6	1.3	Iris-versicolor
59	5.2	2.7	3.9	1.4	Iris-versicolor
60	5.0	2.0	3.5	1.0	Iris-versicolor
61	5.9	3.0	4.2	1.5	Iris-versicolor
62	6.0	2.2	4.0	1.0	Iris-versicolor
63	6.1	2.9	4.7	1.4	Iris-versicolor
64	5.6	2.9	3.6	1.3	Iris-versicolor
65	6.7	3.1	4.4	1.4	Iris-versicolor
66	5.6	3.0	4.5	1.5	Iris-versicolor
67	5.8	2.7	4.1	1.0	Iris-versicolor
68	6.2	2.2	4.5	1.5	Iris-versicolor
69	5.6	2.5	3.9	1.1	Iris-versicolor
70	5.9	3.2	4.8	1.8	Iris-versicolor
71	6.1	2.8	4.0	1.3	Iris-versicolor
72	6.3	2.5	4.9	1.5	Iris-versicolor

73	6.1	2.8	4.7	1.2	Iris-versicolor
74	6.4	2.9	4.3	1.3	Iris-versicolor
75	6.6	3.0	4.4	1.4	Iris-versicolor
76	6.8	2.8	4.8	1.4	Iris-versicolor
77	6.7	3.0	5.0	1.7	Iris-versicolor
78	6.0	2.9	4.5	1.5	Iris-versicolor
79	5.7	2.6	3.5	1.0	Iris-versicolor
80	5.5	2.4	3.8	1.1	Iris-versicolor
81	5.5	2.4	3.7	1.0	Iris-versicolor
82	5.8	2.7	3.9	1.2	Iris-versicolor
83	6.0	2.7	5.1	1.6	Iris-versicolor
84	5.4	3.0	4.5	1.5	Iris-versicolor
85	6.0	3.4	4.5	1.6	Iris-versicolor
86	6.7	3.1	4.7	1.5	Iris-versicolor
87	6.3	2.3	4.4	1.3	Iris-versicolor
88	5.6	3.0	4.1	1.3	Iris-versicolor
89	5.5	2.5	4.0	1.3	Iris-versicolor
90	5.5	2.6	4.4	1.2	Iris-versicolor
91	6.1	3.0	4.6	1.4	Iris-versicolor
92	5.8	2.6	4.0	1.2	Iris-versicolor
93	5.0	2.3	3.3	1.0	Iris-versicolor
94	5.6	2.7	4.2	1.3	Iris-versicolor
95	5.7	3.0	4.2	1.2	Iris-versicolor
96	5.7	2.9	4.2	1.3	Iris-versicolor
97	6.2	2.9	4.3	1.3	Iris-versicolor
98	5.1	2.5	3.0	1.1	Iris-versicolor
99	5.7	2.8	4.1	1.3	Iris-versicolor
100	6.3	3.3	6.0	2.5	Iris-virginica
101	5.8	2.7	5.1	1.9	Iris-virginica
102	7.1	3.0	5.9	2.1	Iris-virginica
103	6.3	2.9	5.6	1.8	Iris-virginica
104	6.5	3.0	5.8	2.2	Iris-virginica
105	7.6	3.0	6.6	2.1	Iris-virginica
106	4.9	2.5	4.5	1.7	Iris-virginica
107	7.3	2.9	6.3	1.8	Iris-virginica
108	6.7	2.5	5.8	1.8	Iris-virginica
109	7.2	3.6	6.1	2.5	Iris-virginica
110	6.5	3.2	5.1	2.0	Iris-virginica
111	6.4	2.7	5.3	1.9	Iris-virginica
112	6.8	3.0	5.5	2.1	Iris-virginica
113	5.7	2.5	5.0	2.0	Iris-virginica
114	5.8	2.8	5.1	2.4	Iris-virginica
115	6.4	3.2	5.3	2.3	Iris-virginica
116	6.5	3.0	5.5	1.8	Iris-virginica
117	7.7	3.8	6.7	2.2	Iris-virginica
118	7.7	2.6	6.9	2.3	Iris-virginica
119	6.0	2.2	5.0	1.5	Iris-virginica
120	6.9	3.2	5.7	2.3	Iris-virginica
121	5.6	2.8	4.9	2.0	Iris-virginica
122	7.7	2.8	6.7	2.0	Iris-virginica
123	6.3	2.7	4.9	1.8	Iris-virginica
124	6.7	3.3	5.7	2.1	Iris-virginica
125	7.2	3.2	6.0	1.8	Iris-virginica
126	6.2	2.8	4.8	1.8	Iris-virginica
127	6.1	3.0	4.9	1.8	Iris-virginica
128	6.4	2.8	5.6	2.1	Iris-virginica
129	7.2	3.0	5.8	1.6	Iris-virginica

130	7.4	2.8	6.1	1.9	Iris-virginica
131	7.9	3.8	6.4	2.0	Iris-virginica
132	6.4	2.8	5.6	2.2	Iris-virginica
133	6.3	2.8	5.1	1.5	Iris-virginica
134	6.1	2.6	5.6	1.4	Iris-virginica
135	7.7	3.0	6.1	2.3	Iris-virginica
136	6.3	3.4	5.6	2.4	Iris-virginica
137	6.4	3.1	5.5	1.8	Iris-virginica
138	6.0	3.0	4.8	1.8	Iris-virginica
139	6.9	3.1	5.4	2.1	Iris-virginica
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

4.3.2 Get data types

```
[8]: irisdata_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null    float64
1   sepal_width     150 non-null    float64
2   petal_length    150 non-null    float64
3   petal_width     150 non-null    float64
4   species         150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
[9]: irisdata_df.describe()
```

```
[9]:
```

	sepal_length	sepal_width	petal_length	petal_width
count	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.054000	3.758667	1.198667
std	0.828066	0.433594	1.764420	0.763161
min	4.300000	2.000000	1.000000	0.100000
25%	5.100000	2.800000	1.600000	0.300000
50%	5.800000	3.000000	4.350000	1.300000
75%	6.400000	3.300000	5.100000	1.800000
max	7.900000	4.400000	6.900000	2.500000

4.3.3 Get data ranges with Boxplots

Boxplots can be used to explore the data ranges in the dataset. These also provide information about outliers.

```
[10]: sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.0})
sns.set_style("whitegrid")
#sns.set_style("white")

fig, axs = plt.subplots(2, 2, figsize=(12, 10))

fn = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']
cn = ['Iris-setosa', 'Iris-versicolor', 'Iris-virginica']
box1 = sns.boxplot(x = 'species', y = 'sepal_length',
                   data = irisdata_df, order = cn, ax = axs[0,0])
box2 = sns.boxplot(x = 'species', y = 'sepal_width',
                   data = irisdata_df, order = cn, ax = axs[0,1])
box3 = sns.boxplot(x = 'species', y = 'petal_length',
                   data = irisdata_df, order = cn, ax = axs[1,0])
box4 = sns.boxplot(x = 'species', y = 'petal_width',
                   data = irisdata_df, order = cn, ax = axs[1,1])

# add some spacing between subplots
fig.tight_layout(pad=2.0)

plt.show()
```

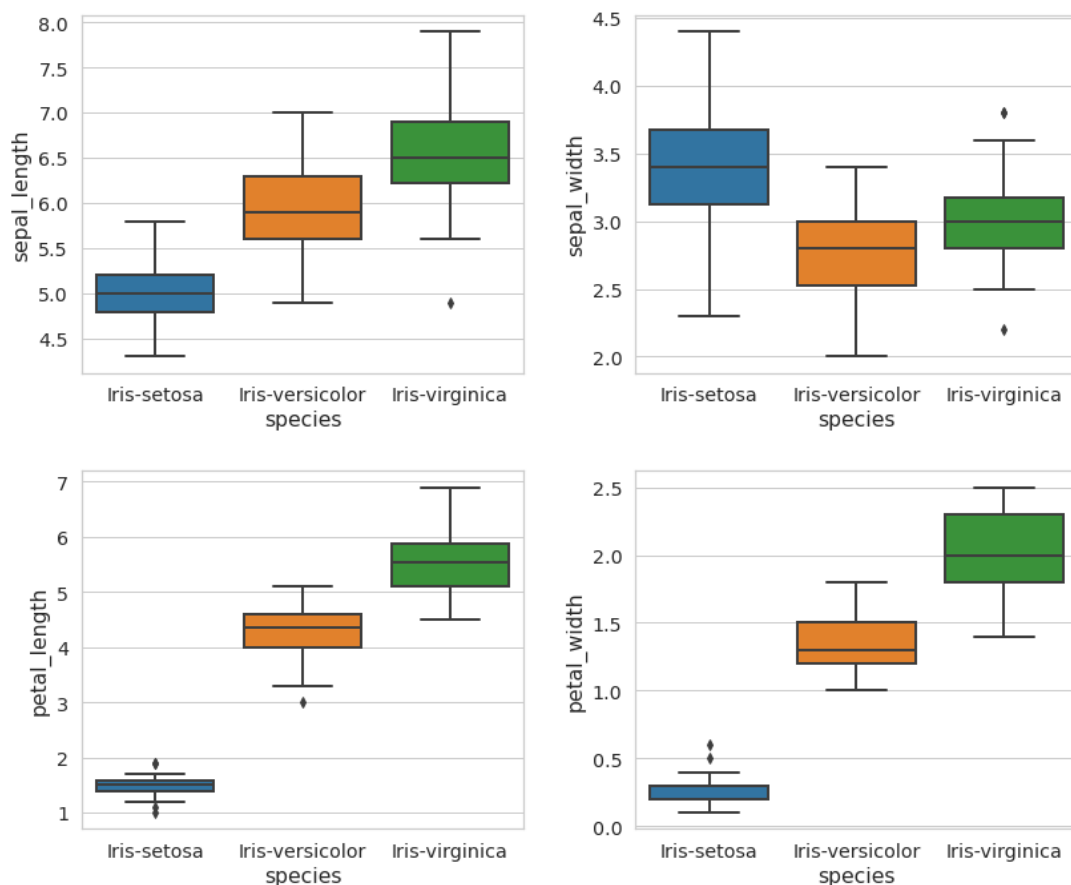


Figure 3: Boxplots used to explore the data ranges in the Iris dataset

4.4 Identify anomalies in the data sets

4.4.1 Find gaps in dataset

This section was inspired by [Working with Missing Data in Pandas](#).

Checking for missing values using `isnull()` In order to check for missing values in Pandas DataFrame, we use the function `isnull()`. This function returns a dataframe of Boolean values which are True for NaN values.

```
[11]: pd.set_option('display.max_rows', 40)
      pd.set_option('display.min_rows', 30)
```

```
[12]: irisdata_df.isnull()
```

```
[12]:      sepal_length  sepal_width  petal_length  petal_width  species
0             False           False           False           False      False
1             False           False           False           False      False
2             False           False           False           False      False
3             False           False           False           False      False
4             False           False           False           False      False
5             False           False           False           False      False
6             False           False           False           False      False
7             False           False           False           False      False
8             False           False           False           False      False
9             False           False           False           False      False
10            False           False           False           False      False
11            False           False           False           False      False
12            False           False           False           False      False
13            False           False           False           False      False
14            False           False           False           False      False
..            ...             ...             ...             ...             ...
135           False           False           False           False      False
136           False           False           False           False      False
137           False           False           False           False      False
138           False           False           False           False      False
139           False           False           False           False      False
140           False           False           False           False      False
141           False           False           False           False      False
142           False           False           False           False      False
143           False           False           False           False      False
144           False           False           False           False      False
145           False           False           False           False      False
146           False           False           False           False      False
147           False           False           False           False      False
148           False           False           False           False      False
149           False           False           False           False      False
```

[150 rows x 5 columns]

Show only the gaps:

```
[13]: irisdata_df_gaps = irisdata_df[irisdata_df.isnull().any(axis=1)]
      irisdata_df_gaps
```

```
[13]: Empty DataFrame
      Columns: [sepal_length, sepal_width, petal_length, petal_width, species]
      Index: []
```

Fine - this dataset seems to be complete :)

So let's look for something else for exercise: [employees.csv](#)

```
[14]: # import data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")

employees_df
```

```
[14]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	
998	George	Male	6/21/2013	5:47 PM	98874	4479.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance

```

...
989          False          Legal
990          False      Marketing
991           True          Finance
992          False      Human Resources
993          False          Legal
994           True      Client Services
995           True          Marketing
996           True          Finance
997           True      Engineering
998           True          Marketing
999          False      Distribution
1000         False          Finance
1001         False          Product
1002         False      Business Development
1003         True           Sales

```

[1004 rows x 8 columns]

Show only the gaps from this gappy dataset again:

```
[15]: employees_df_gaps = employees_df[employees_df.isnull().any(axis=1)]
      employees_df_gaps
```

```
[15]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
20	Lois	NaN	4/22/1995	7:18 PM	64714	4934.00	
22	Joshua	NaN	3/8/2012	1:58 AM	90816	18816.00	
23	NaN	Male	6/14/2012	4:19 PM	125792	5042.00	
25	NaN	Male	10/8/2012	1:12 AM	37076	18576.00	
27	Scott	NaN	7/11/1991	6:58 PM	122367	5218.00	
31	Joyce	NaN	2/20/2005	2:40 PM	88657	12752.00	
32	NaN	Male	8/21/1998	2:27 PM	122340	6417.00	
39	NaN	Male	1/29/2016	2:33 AM	122173	7797.00	
41	Christine	NaN	6/28/2015	1:08 AM	66582	11308.00	
49	Chris	NaN	1/24/1980	12:13 PM	113590	3055.00	
51	NaN	NaN	12/17/2011	8:29 AM	41126	14009.00	
53	Alan	NaN	3/3/2014	1:28 PM	40341	17578.00	
..		
916	Joe	Male	12/8/1998	10:28 AM	126120	1.02	
927	Irene	NaN	2/28/1991	10:23 PM	135369	4.38	
929	NaN	Female	8/23/2000	4:19 PM	95866	19388.00	
941	Aaron	NaN	1/22/1986	7:39 PM	63126	18424.00	
942	Mark	NaN	9/9/2006	12:27 PM	44836	2657.00	
943	Ralph	NaN	7/28/1995	6:53 PM	70635	2147.00	
949	Gerald	NaN	4/15/1989	12:44 PM	93712	17426.00	
950	NaN	Female	9/15/1985	1:50 AM	133472	16941.00	
951	NaN	Male	7/30/2012	3:07 PM	107351	5329.00	
955	NaN	Female	9/14/2010	5:19 AM	143638	9662.00	
965	Antonio	NaN	6/18/1989	9:37 PM	103050	3.05	
976	Victor	NaN	7/28/2006	2:49 PM	76381	11159.00	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00	
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00	

Senior Management

Team


```

1          True          NaN
7          NaN          Finance
10         True          NaN
20         True          Legal
22         True    Client Services
23         NaN          NaN
25         NaN    Client Services
27         False         Legal
31         False         Product
32         NaN          NaN
39         NaN    Client Services
41         True  Business Development
49         False         Sales
51         NaN          Sales
53         True          Finance
..         ...          ...
916        False          NaN
927        False  Business Development
929         NaN          Sales
941        False    Client Services
942        False    Client Services
943        False    Client Services
949         True    Distribution
950         NaN    Distribution
951         NaN    Marketing
955         NaN          NaN
965        False         Legal
976         True         Sales
989        False         Legal
993        False         Legal
999        False    Distribution

```

[237 rows x 8 columns]

Fill the missing values with fillna() Now we are going to fill all the null (NaN) values in Gender column with “No Gender”.

Attention: We are doing that directly in this dataframe with `inplace = True` - we don't make a deep copy!

```

[16]: # filling a null values using fillna()
employees_df["Gender"].fillna("No Gender", inplace = True)
employees_df

```

```

[16]:   First Name  Gender  Start Date  Last Login Time  Salary  Bonus % \
0    Douglas    Male    8/6/1993      12:42 PM    97308    6945.00
1    Thomas    Male    3/31/1996      6:53 AM     61933      4.17
2    Maria    Female  4/23/1993      11:17 AM    130590   11858.00
3    Jerry    Male    3/4/2005       1:00 PM    138705     9.34
4    Larry    Male    1/24/1998      4:47 PM    101004   1389.00
5    Dennis    Male    4/18/1987      1:35 AM    115163   10125.00
6    Ruby    Female    8/17/1987      4:20 PM     65476   10012.00
7    NaN    Female    7/20/2015     10:43 AM     45906   11598.00
8    Angela    Female  11/22/2005      6:29 AM     95570   18523.00
9    Frances    Female    8/8/2002      6:51 AM    139852    7524.00
10   Louise    Female    8/12/1980      9:01 AM     63241   15132.00
11   Julie    Female   10/26/1997      3:19 PM    102508   12637.00

```

12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
...
989	Stephen	No Gender	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	No Gender	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[1004 rows x 8 columns]

Dropping missing values using dropna() In order to drop null values from a dataframe, we use dropna() function. This function drops rows or columns of datasets with NaN values in different ways.

Default is to drop rows with at least 1 null value (NaN). Giving the parameter `how = 'all'` the function drops rows with all data missing or contain null values (NaN).

```
[17]: # making a new dataframe with dropped NaN values
employees_df_dropped = employees_df.dropna(axis = 0, how ='any')
employees_df_dropped
```

```
[17]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
15	Lillian	Female	6/5/2016	6:09 AM	59414	1256.00	
16	Jeremy	Male	9/21/2010	5:56 AM	90370	7369.00	
17	Shawn	Male	12/7/1986	7:45 PM	111737	6414.00	
...	
989	Stephen	No Gender	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	
993	Justin	No Gender	2/10/1991	4:58 PM	38344	3794.00	
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00	
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00	
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00	
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04	
998	George	Male	6/21/2013	5:47 PM	98874	4479.00	
999	Henry	No Gender	11/23/2014	6:09 AM	132483	16655.00	
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00	
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00	
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00	
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00	

	Senior Management	Team
0	True	Marketing
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
8	True	Engineering
9	True	Business Development
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
15	False	Product
16	False	Human Resources
17	False	Product
...
989	False	Legal

```

990          False          Marketing
991           True           Finance
992          False    Human Resources
993          False           Legal
994           True    Client Services
995           True           Marketing
996           True           Finance
997           True           Engineering
998           True           Marketing
999          False    Distribution
1000          False           Finance
1001          False           Product
1002          False    Business Development
1003           True           Sales

```

```
[903 rows x 8 columns]
```

Finally we compare the sizes of dataframes so that we learn how many rows had at least 1 Null value.

```
[18]: print("Old data frame length:", len(employees_df))
      print("New data frame length:", len(employees_df_dropped))
      print("Number of rows with at least 1 NaN value: ",
            (len(employees_df)-len(employees_df_dropped)))
```

```
Old data frame length: 1004
```

```
New data frame length: 903
```

```
Number of rows with at least 1 NaN value: 101
```

4.4.2 Find and remove duplicates in dataset

This section was inspired by: - [How to Find Duplicates in Pandas DataFrame \(With Examples\)](#) - [How to Drop Duplicate Rows in a Pandas DataFrame](#)

Checking for duplicate values using duplicated() In order to check for duplicate values in Pandas DataFrame, we use a function `duplicated()`. This function can be used in two ways: - find duplicate rows across **all columns** with `duplicateRows = df[df.duplicated()]` - find duplicate rows across **specific columns** `duplicateRows = df[df.duplicated(subset=['col1', 'col2'])]`

Find duplicate rows across **all columns**:

```
[19]: # import (again) data to dataframe from csv file
      employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[20]: # find duplicate rows across all columns
      duplicateRows = employees_df[employees_df.duplicated()]
      duplicateRows
```

```
[20]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
112	Karen	Female	11/30/1999	7:46 AM	102488	17653.0	
127	Linda	Female	5/25/2000	5:45 PM	119009	12506.0	
296	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.0	
580	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.0	

	Senior Management	Team
112	True	Product
127	True	Business Development
296	False	Business Development

580 True Human Resources

```
[21]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(keep='last')]
duplicateRows
```

```
[21]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
55      Karen  Female  11/30/1999      7:46 AM  102488  17653.0
92      Linda  Female   5/25/2000      5:45 PM  119009  12506.0
153     Brandon   NaN   11/3/1997      8:17 PM  121333  15295.0
442    Nicholas   Male   3/1/2013      9:26 PM  101036   2826.0

      Senior Management      Team
55      True      Product
92      True  Business Development
153     False  Business Development
442      True   Human Resources
```

Find duplicate rows across **specific columns**:

```
[22]: # identify duplicate rows across 'First Name' and 'Last Login Time' columns
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'])]
duplicateRows
```

```
[22]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
112      Karen  Female  11/30/1999      7:46 AM  102488  17653.0
127      Linda  Female   5/25/2000      5:45 PM  119009  12506.0
296     Brandon   NaN   11/3/1997      8:17 PM  121333  15295.0
577      NaN  Female   1/13/2009      1:01 PM  118736   7421.0
580    Nicholas   Male   3/1/2013      9:26 PM  101036   2826.0
632      NaN   NaN   9/2/1988      12:49 PM  147309   1702.0
881      NaN   Male   9/5/1980      7:36 AM  114896  13823.0
929      NaN  Female   8/23/2000      4:19 PM   95866  19388.0
934      Nancy  Female   9/10/2001     11:57 PM   85213   2386.0
973      Linda  Female   2/4/2010      8:49 PM   44486  17308.0

      Senior Management      Team
112      True      Product
127      True  Business Development
296     False  Business Development
577      NaN   Client Services
580      True   Human Resources
632      NaN   Distribution
881      NaN   Client Services
929      NaN      Sales
934      True   Marketing
973      True   Engineering
```

```
[23]: # argument keep='last' displays the first duplicate rows instead of the last
duplicateRows = employees_df[employees_df.duplicated(
    subset=['First Name', 'Last Login Time'], keep='last')]
duplicateRows
```

```
[23]:      First Name  Gender  Start Date  Last Login Time  Salary  Bonus %  \
23      NaN   Male   6/14/2012      4:19 PM  125792   5042.00
37      Linda  Female  10/19/1981      8:49 PM   57427   9557.00
```

55	Karen	Female	11/30/1999	7:46 AM	102488	17653.00
66	Nancy	Female	12/15/2012	11:57 PM	125250	2672.00
92	Linda	Female	5/25/2000	5:45 PM	119009	12506.00
153	Brandon	NaN	11/3/1997	8:17 PM	121333	15295.00
222	NaN	Female	6/17/1991	12:49 PM	71945	5.56
269	NaN	Male	2/4/2005	1:01 PM	40451	16044.00
442	Nicholas	Male	3/1/2013	9:26 PM	101036	2826.00
778	NaN	Female	6/18/2000	7:36 AM	106428	10867.00

	Senior Management	Team
23	NaN	NaN
37	True	Client Services
55	True	Product
66	True	Business Development
92	True	Business Development
153	False	Business Development
222	NaN	Marketing
269	NaN	Distribution
442	True	Human Resources
778	NaN	NaN

Dropping duplicate values using drop_duplicates() In order to drop duplicate values from a dataframe, we use drop_duplicates() function.

This function can be used in two ways: - remove duplicate rows across **all columns** with df.drop_duplicates() - find duplicate rows across **specific columns** df.drop_duplicates(subset=['col1', 'col2'])

Attention: We are doing that directly in this dataframe with inplace = True - we don't make a deep copy!

Remove duplicate rows across **all columns**:

```
[24]: # remove duplicate rows across all columns
employees_df.drop_duplicates(inplace=True)
employees_df
```

```
[24]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	
3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34	
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00	
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00	
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00	
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00	
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00	
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00	
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00	
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00	
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00	
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00	
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00	
...	
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00	
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00	
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00	
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00	

993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution
1000	False	Finance
1001	False	Product
1002	False	Business Development
1003	True	Sales

[1000 rows x 8 columns]

Remove duplicate rows across **specific columns**:

```
[25]: # remove duplicate rows across 'First Name' and 'Last Login Time' columns
employees_df.drop_duplicates(
    subset=['First Name', 'Last Login Time'], keep='last', inplace=True)
employees_df
```

```
[25]:
```

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus %	\
0	Douglas	Male	8/6/1993	12:42 PM	97308	6945.00	
1	Thomas	Male	3/31/1996	6:53 AM	61933	4.17	
2	Maria	Female	4/23/1993	11:17 AM	130590	11858.00	

3	Jerry	Male	3/4/2005	1:00 PM	138705	9.34
4	Larry	Male	1/24/1998	4:47 PM	101004	1389.00
5	Dennis	Male	4/18/1987	1:35 AM	115163	10125.00
6	Ruby	Female	8/17/1987	4:20 PM	65476	10012.00
7	NaN	Female	7/20/2015	10:43 AM	45906	11598.00
8	Angela	Female	11/22/2005	6:29 AM	95570	18523.00
9	Frances	Female	8/8/2002	6:51 AM	139852	7524.00
10	Louise	Female	8/12/1980	9:01 AM	63241	15132.00
11	Julie	Female	10/26/1997	3:19 PM	102508	12637.00
12	Brandon	Male	12/1/1980	1:08 AM	112807	17492.00
13	Gary	Male	1/27/2008	11:40 PM	109831	5831.00
14	Kimberly	Female	1/14/1999	7:13 AM	41426	14543.00
...
989	Stephen	NaN	7/10/1983	8:10 PM	85668	1909.00
990	Donna	Female	11/26/1982	7:04 AM	82871	17999.00
991	Gloria	Female	12/8/2014	5:08 AM	136709	10331.00
992	Alice	Female	10/5/2004	9:34 AM	47638	11209.00
993	Justin	NaN	2/10/1991	4:58 PM	38344	3794.00
994	Robin	Female	7/24/1987	1:35 PM	100765	10982.00
995	Rose	Female	8/25/2002	5:12 AM	134505	11051.00
996	Anthony	Male	10/16/2011	8:35 AM	112769	11625.00
997	Tina	Female	5/15/1997	3:53 PM	56450	19.04
998	George	Male	6/21/2013	5:47 PM	98874	4479.00
999	Henry	NaN	11/23/2014	6:09 AM	132483	16655.00
1000	Phillip	Male	1/31/1984	6:30 AM	42392	19675.00
1001	Russell	Male	5/20/2013	12:39 PM	96914	1421.00
1002	Larry	Male	4/20/2013	4:45 PM	60500	11985.00
1003	Albert	Male	5/15/2012	6:24 PM	129949	10169.00

	Senior Management	Team
0	True	Marketing
1	True	NaN
2	False	Finance
3	True	Finance
4	True	Client Services
5	False	Legal
6	True	Product
7	NaN	Finance
8	True	Engineering
9	True	Business Development
10	True	NaN
11	True	Legal
12	True	Human Resources
13	False	Sales
14	True	Finance
...
989	False	Legal
990	False	Marketing
991	True	Finance
992	False	Human Resources
993	False	Legal
994	True	Client Services
995	True	Marketing
996	True	Finance
997	True	Engineering
998	True	Marketing
999	False	Distribution

```

1000          False          Finance
1001          False          Product
1002          False  Business Development
1003           True           Sales

```

```
[994 rows x 8 columns]
```

4.5 Avoidance of tendencies due to bias

The description of the Iris dataset says, that it consists of **50 samples** from **each of three species** of Iris (Iris setosa, Iris virginica and Iris versicolor), so there are **150 total samples**.

But how to prove it?

4.5.1 Count occurrences of unique values

To prove whether all possible classes included in the dataset and equally distributed, you can use the function `df.value_counts`.

Following parameters can be used for fine tuning: - `dropna=False` causes that NaN values are included - `normalize=True`: relative frequencies of the unique values are returned - `ascending=False`: sort resulting classes descending

```
[26]: # import (again) data to dataframe from csv file
employees_df = pd.read_csv("../datasets/employees_edit.csv")
```

```
[27]: # count unique values without missing values in a column,
# ordered descending and normalized
irisdata_df['species'].value_counts(ascending=False, dropna=False, normalize=True)
```

```
[27]: Iris-setosa          0.333333
Iris-versicolor         0.333333
Iris-virginica          0.333333
Name: species, dtype: float64
```

```
[28]: # count unique values and missing values in a column,
# ordered descending and not absolute values
employees_df['Team'].value_counts(ascending=False, dropna=False, normalize=False)
```

```
[28]: Client Services      106
Business Development     103
Finance                  102
Marketing                 98
Product                  96
Sales                    94
Engineering               92
Human Resources           92
Distribution              90
Legal                    88
NaN                       43
Name: Team, dtype: int64
```

4.5.2 Display Histogram

This section was inspired by: [Pandas Histogram – DataFrame.hist\(\)](#).

Histograms represent **frequency distributions** graphically. This requires the separation of the data into classes (so-called **bins**).

These classes are represented in the histogram as rectangles of equal or variable width. The height of each rectangle then represents the (relative or absolute) **frequency density**.

```
[29]: employees_df.hist(column=['Salary'])
plt.show()
```




Figure 4: Histogram for frequency distribution of the salary

```
[30]: employees_df.hist(column='Salary', by='Gender')
plt.show()
```

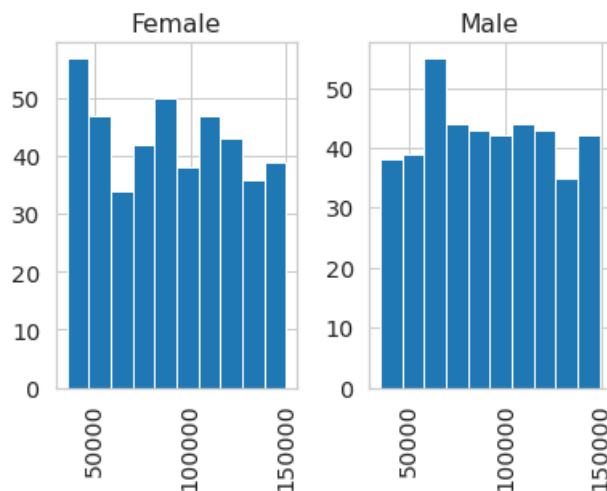


Figure 5: Histogram for the frequency distribution of the salary in comparison between men and women

4.6 First idea of correlations in data set

To get a rough idea of the **dependencies** and **correlations** in the data set, it can be helpful to visualize the whole dataset in a **correlation heatmap**. They show in a glance which variables are correlated, to what degree and in which direction.

Later, 2 particularly well correlated variables are selected from the data set and plotted in a **scatterplot**.

4.6.1 Visualise data with correlation heatmap

This section was inspired by [How to Create a Seaborn Correlation Heatmap in Python?](#).

Correlation matrices are an **essential tool of exploratory data analysis**. Correlation heatmaps contain the same information in a visually appealing way. What more: they show in a glance which variables are correlated, to what degree, in which direction, and alerts us to potential multicollinearity problems (source: ibidem).

Simple correlation matrix Because **string values can never be correlated**, the class names (species) have to be converted first:

```
[187]: # encoding the class column
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
                                                    "Iris-virginica":2}})

#irisdata_df_enc
```

```
[32]: irisdata_df_enc.corr()
```

```
[32]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
sepal_length	1.000000	-0.109369	0.871754	0.817954	0.782561
sepal_width	-0.109369	1.000000	-0.420516	-0.356544	-0.419446
petal_length	0.871754	-0.420516	1.000000	0.962757	0.949043
petal_width	0.817954	-0.356544	0.962757	1.000000	0.956464
species	0.782561	-0.419446	0.949043	0.956464	1.000000

Correlation heatmap Choose the color sets from [color map](#).

```
[33]: # increase the size of the heatmap
plt.figure(figsize=(16, 6))

# store heatmap object in a variable to easily access it
# when you want to include more features (such as title)
# set the range of values to be displayed on the colormap from -1 to 1,
# and set 'annotation=True' to display the correlation values on the heatmap
heatmap = sns.heatmap(irisdata_df_enc.corr(), vmin=-1, vmax=1,
                      annot=True, cmap='PRGn_r')

# give a title to the heatmap
# 'pad=12' defines the distance of the title from the top of the heatmap
heatmap.set_title('Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
plt.show()
```




Figure 6: Correlation heatmap to explore coherences between single variables in the iris dataset

Triangle correlation heatmap When looking at the correlation heatmaps above, you would not lose any information by **cutting** away half of it **along the diagonal** line marked by 1-s.

The **numpy** function `np.triu()` can be used to isolate the upper triangle of a matrix while turning all the values in the lower triangle into 0.

```
[34]: import numpy as np

      np.triu(np.ones_like(irisdata_df_enc.corr()))
```

```
[34]: array([[1., 1., 1., 1., 1.],
            [0., 1., 1., 1., 1.],
            [0., 0., 1., 1., 1.],
            [0., 0., 0., 1., 1.],
            [0., 0., 0., 0., 1.]])
```

Use this mask to cut the heatmap along the diagonal:

```
[35]: plt.figure(figsize=(16, 6))

      # define the mask to set the values in the upper triangle to 'True'
      mask = np.triu(np.ones_like(irisdata_df_enc.corr(), dtype=bool))

      heatmap = sns.heatmap(irisdata_df_enc.corr(), mask=mask,
                           vmin=-1, vmax=1, annot=True, cmap='PRGn_r')

      heatmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18}, pad=16)
      plt.show()
```




Figure 7: Correlation heatmap, which was cut at its main diagonal without losing any information

As a result from the **heatmaps** we can see, that the shape of the **petals** are the **most correlated columns** (0.96) with the **type of flowers** (species classes).

Somewhat lower correlates **sepal length** with **petal length** (0.87).

4.6.2 Visualise data with scatter plot

In the following, **Seaborn** is applied which is a library for making statistical graphics in Python. It is built on top of matplotlib and closely integrated with pandas data structures.

To investigate whether there are dependencies (e.g. correlations) in `irisdata_df` between individual variables in the data set, it is advisable to plot them in a **scatter plot**.

```
[36]: # There are five preset seaborn themes: darkgrid, whitegrid, dark, white, and ticks.
sns.set_style("whitegrid")
# set scale of fonts
sns.set_context("notebook", font_scale=1.3, rc={"lines.linewidth": 2.5})

# 'sepal_length', 'petal_length' are iris feature data
# 'height' used to define height of graph
# 'hue' stores the class/label of iris dataset
sns.FacetGrid(irisdata_df, hue="species",
               height=7).map(plt.scatter,
                             'petal_width',
                             'petal_length').add_legend()

plt.title('Scatterplot of petal length and width')
plt.show()
```

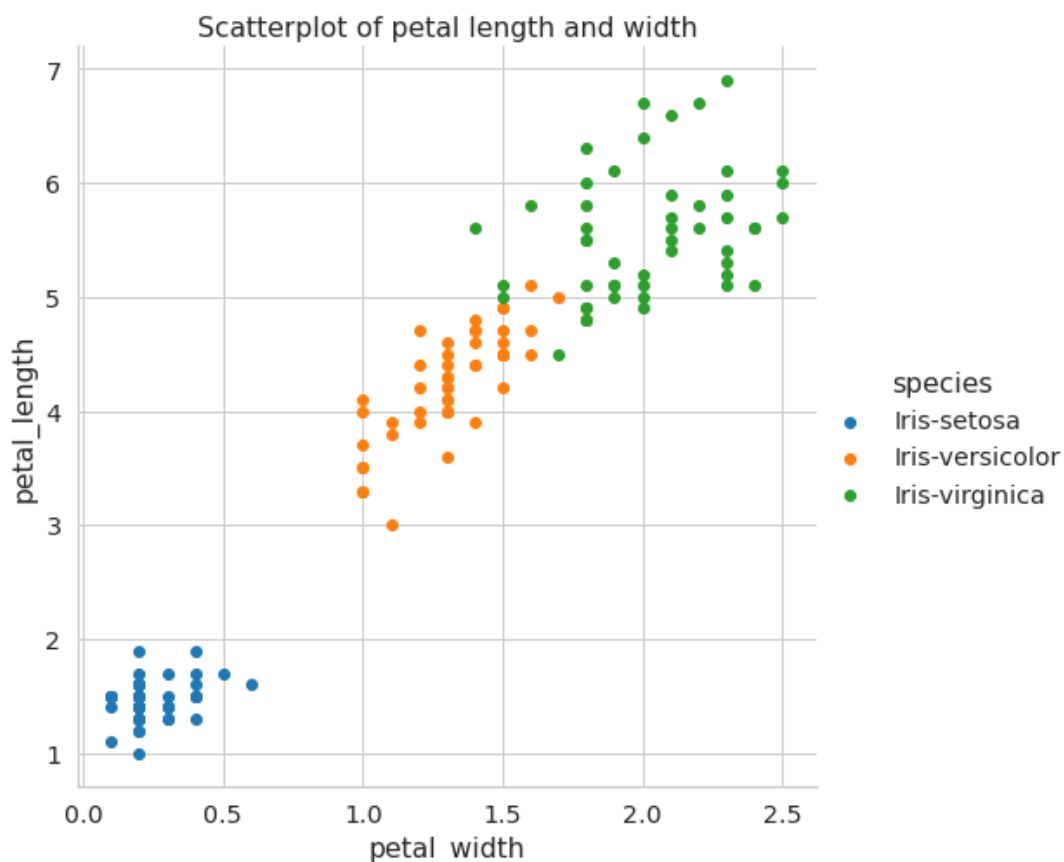


Figure 8: Plotting two individual variables of the iris dataset in the scatterplot to explore the relationships between these two

4.6.3 Visualise data with pairs plot

For systematic investigation of dependencies, all variables (each against each) are plotted in separate scatter plots.

With this so called **pairs plot** it is possible to see both **relationships** between two variables and **distribution** of single variables.

This function will create a grid of Axes such that **each numeric variable** in `irisdata_df` will be shared in the y-axis across a single row and in the x-axis across a single column.

```
[195]: sns.set(font_scale=1.0)
sns.set_style("white")

g = sns.pairplot(irisdata_df, diag_kind="kde", hue='species',
                palette='Dark2', height=2.5)

g.map_lower(sns.kdeplot, levels=4, color=".2")
# y .. padding between title and plot
g.fig.suptitle('Pairs plot of the Iris dataset', y=1.05)
plt.show()
```

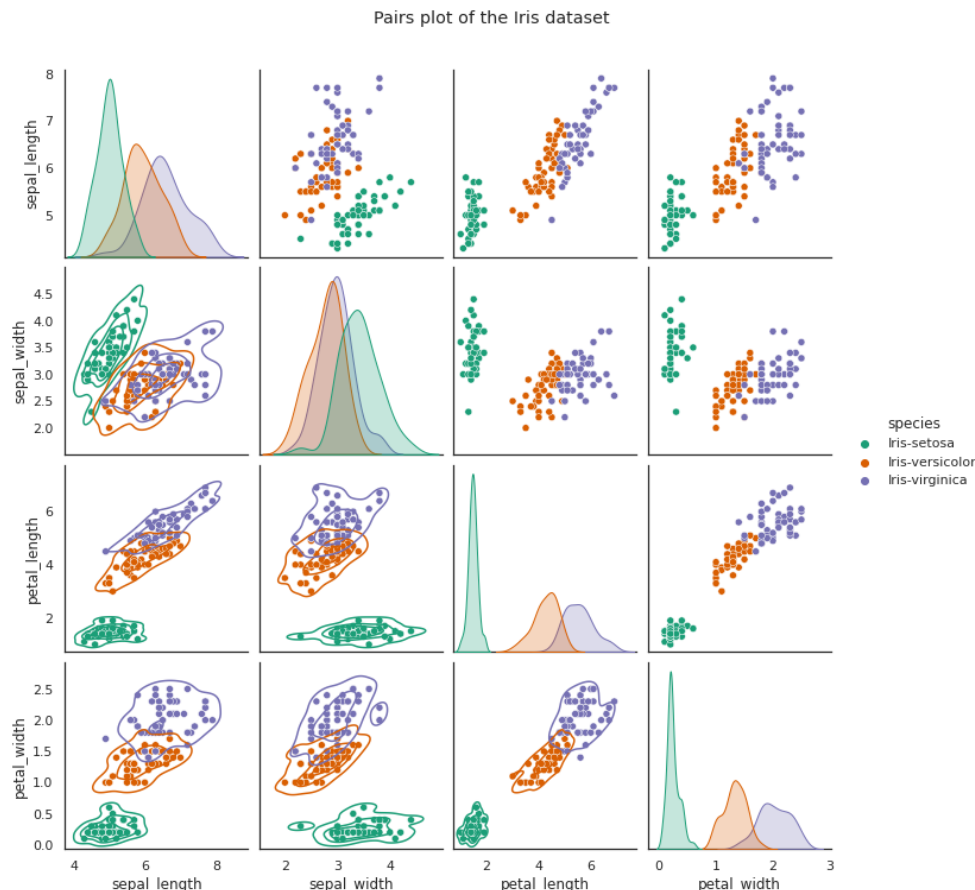


Figure 9: Plot all individual variables of the Iris dataset in pairs plot to see both the relationships between two variables and the distribution of the individual variables

5 STEP 2: Prepare the dataset

Through the intensive exploration of the data in Step 1 ([STEP 1: Exploring the dataset](#)), we know that special **preparation** of the data is **not necessary**. The values are **complete** and **without gaps** and there are **no duplicates**. The values are in similar ranges, which **does not require normalization** of the data.

Furthermore, we know that the **classes** are very **evenly distributed** and thus bias tendencies should be avoided.

6 STEP 3: Classify by support vector classifier - SVC

6.1 Operating principal

Support Vectors Classifier tries to **find the best hyperplane to separate** the different classes by maximizing the distance between sample points and the hyperplane (source: [In Depth: Parameter tuning for SVC](#)).

The figure ?? shows the operating principal of the SVC algorithm: the hyperplanes $H1$ till $H4$ (left graphic) do separate the classes. A good separation is achieved by the hyperplane that has the largest distance to the nearest training-data point of any class (so-called functional margin), since in general the larger the margin, the lower the generalization error of the classifier (source: [Support-vector machine](#)).

The right graphic shows the optimal hyperplane characterized by maximising the margin between the classes. The perpendicular distance of the closest data points to the hyperplane determines their position and orientation. These perpendicular distances are the **support vectors** of the hyperplane - this is how the algorithm got its name.

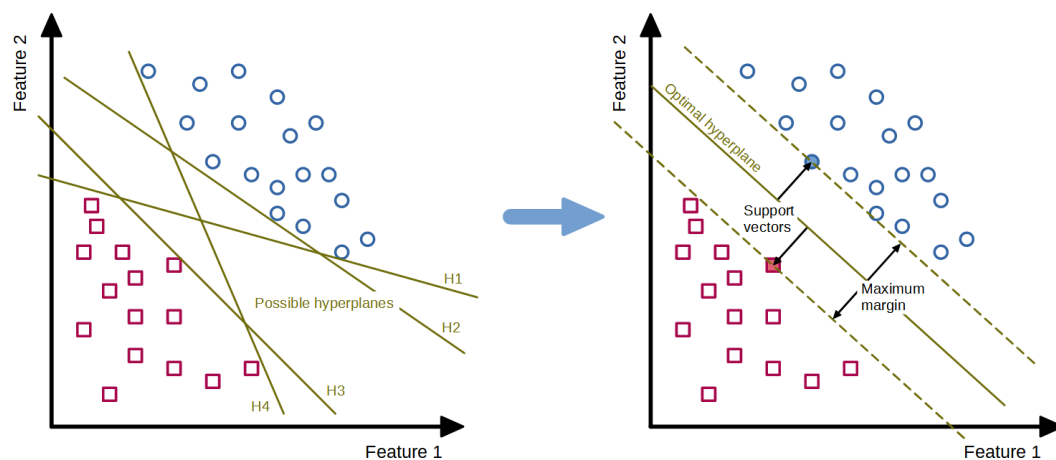


Figure 10: Support Vectors Classifiers (SVC) separate the data points in classes by finding the best hyperplane by maximizing the margin to its support vectors

6.2 Split the dataset

In the next very important step, the dataset is split into **2 subsets**: a **training dataset** and a **test dataset**. As the names suggest, the training dataset is used to train the ML algorithm. The test data set is then used to check the quality of the trained ML algorithm (here the **recognition rate**). For this purpose, the **class labels** are **removed** from the training data set - after all, these are to be predicted.

Typically, the **test dataset** should contain about **20%** of the entire dataset.

```
[63]: from sklearn.model_selection import train_test_split

X = irisdata_df.drop('species', axis=1)
y = irisdata_df['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

For training, do not use only the variables that correlate best with each other, but all of them.

Otherwise, the result of the prediction would be significantly worse. Maybe this is already an indication of **overfitting** of the ML model.

```
[39]: # DO NOT USE THIS!!
X_train, X_test, y_train, y_test = train_test_split(X[['sepal_length',
                                                    'sepal_width']],
```

```
y, test_size = 0.20)
```

6.3 Create the SVM model

In this step we create the SVC model and fit it to our training data.

```
[54]: from sklearn.svm import SVC
classifier = SVC(kernel = 'linear', random_state = 0)

# fit the model for the data
classifier.fit(X_train, y_train)
```

```
[54]: SVC(kernel='linear', random_state=0)
```

6.4 Make predictions

```
[55]: y_pred = classifier.predict(X_test)
#X_test
```

7 STEP 4: Evaluate the classification results - metrics

And finally for checking the accuracy of the model, the **confusion matrix** is used for the **cross validation**.

By using the function `sklearn.metrics.confusion_matrix()` a confusion matrix of the true digit values versus the predicted digit values is plotted.

7.1 Textual confusion matrix

```
[56]: cm = metrics.confusion_matrix(y_test, y_pred)
print(cm)
```

```
[[ 9  0  0]
 [ 0 10  0]
 [ 0  0 11]]
```

7.2 Colored confusion matrix

The function `sklearn.metrics.ConfusionMatrixDisplay()` plots a colored confusion matrix.

```
[57]: sns.set_style("white")

# print colored confusion matrix
cm_colored = metrics.ConfusionMatrixDisplay.from_predictions(y_test, y_pred)

cm_colored.figure_.suptitle("Colored Confusion Matrix")
cm_colored.figure_.set_figwidth(8)
cm_colored.figure_.set_figheight(7)

cm_colored.confusion_matrix

# save figure as PNG
plt.tight_layout()
plt.savefig('images/confusion_matrix.png', dpi=150, pad_inches=5)
```

```
plt.show()
```




Figure 11: Checking the accuracy of the model by using the confusion matrix for cross-validation

7.3 Classification accuracy

```
[58]: from sklearn.model_selection import cross_val_score

accuracies = cross_val_score(estimator = classifier, X = X_train,
                              y = y_train, cv = 10)

print("Accuracy: {:.2f} %".format(accuracies.mean()*100))
print("Standard Deviation: {:.2f} %".format(accuracies.std()*100))
```

Accuracy: 97.50 %
Standard Deviation: 7.50 %

8 STEP 5: Select SVC kernel and vary parameters

This section was inspired by [In Depth: Parameter tuning for SVC](#)

In this section, the 4 SVC parameters `kernel`, `gamma`, `C` and `degree` will be introduced one by one. Furthermore, their influence on the classification result by varying these single parameters will be shown.

Disclaimer: In order to show the effects of varying the individual parameters in 2D graphs, only the best correlating variables `petal_length` and `petal_width` are used to train the SVC.

8.1 Prepare dataset

```
[142]: from sklearn.svm import SVC
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_val_score
import numpy as np

# import iris dataset again
irisdata_df = pd.read_csv('./datasets/IRIS_flower_dataset_kaggle.csv')
```

```
# encode the class column from class strings to integer equivalents
irisdata_df_enc = irisdata_df.replace({"species": {"Iris-setosa":0,
                                                    "Iris-versicolor":1,
                                                    "Iris-virginica":2}})

#irisdata_df_enc
```

8.1.1 Prepare datasets for parameter variation and plotting

These datasets will be used for parameter variation and plotting only. In particular, for later **2D plotting** of the effects of parameter variation, only **2 variables** of the iris dataset can be used.

However, as seen in the previous section, this selection is very much at the expense of detection accuracy. Therefore, it is not useful to make predictions with this subset of data - it is not necessary to divide it into a training and a test data set.

```
[143]: # copy only 2 feature columns
# and convert pandas dataframe to numpy array
X_plot = irisdata_df_enc[['petal_length', 'petal_width']].to_numpy(copy=True)
#X_plot = irisdata_df_enc[['sepal_length', 'sepal_width']].to_numpy(copy=True)
#X_plot
```

```
[144]: # convert pandas dataframe to numpy array
# and get a flat 1D copy of 2D numpy array
y_plot = irisdata_df_enc[['species']].to_numpy(copy=True).flatten()
#y_plot
```

8.1.2 Prepare dataset for prediction and evaluation

To **evaluate the recognition accuracy** by parameter variation, the complete iris data set with all variables must be used. To make predictions with test data, the data set is again divided into a training and a test data set.

```
[145]: X = irisdata_df.drop('species', axis=1)
y = irisdata_df['species']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.20)
```

8.2 Plotting functions

This function helps to visualize the modifications by varying the individual SVC parameters:

```
[146]: def plotSVC(title, svc, X, y, xlabel, ylabel):
    # create a mesh to plot in
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1

    # prevent division by zero
    if x_min == 0.0:
        x_min = 0.1

    h = (x_max / x_min)/1000
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))

    plt.subplot(1, 1, 1)
    Z = svc.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
```

```
plt.contourf(xx, yy, Z, cmap=plt.cm.Paired, alpha=0.6)
plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
plt.xlabel(xlabel)
plt.ylabel(ylabel)
plt.xlim(xx.min(), xx.max())
plt.title(title)
plt.show()
```

This function cares for cross validation:

```
[168]: def crossValSVC(X_train, y_train, kernel='rbf', gamma='scale', C=1.0, degree=3):
        # train the SVC
        svc = svm.SVC(kernel=kernel,
                        gamma=gamma,
                        C=C,
                        degree=degree).fit(X_train, y_train)
        # calculate accuracies
        accuracies = cross_val_score(estimator = svc, X = X_train,
                                      y = y_train, cv = 10)

        accuracy = accuracies.mean()*100
        return accuracy
```

This function plots the variation of the SVC parameters against the prediction accuracy to show the effect of variation and its limits regarding the phenomenon **overfitting**:

```
[182]: def plotParamsAcc(param_list, acc_list, param_name, log_scale=False):
        fig, ax = plt.subplots(figsize=(10,6))
        title_str = 'Variation of {} parameter '.format(param_name) \
                    + 'and its effect to prediction accuracy'
        plt.title(title_str)
        ax.plot(param_list, accuracy_list)
        if log_scale:
            # set the X axis scale to logarithmic
            ax.set_xscale('log')
        plt.xlabel(param_name)
        plt.ylabel('accuracy [%]')
        plt.grid()
        plt.show()
```

8.3 Vary kernel of SVC

The **kernel** parameter selects the type of hyperplane that is used to separate the data. Using **linear** ([linear classifier](#)) kernel will use a linear hyperplane (a line in the case of 2D data). The **rbf** ([radial basis function kernel](#)) and **poly** ([polynomial kernel](#)) kernel use non linear hyperplanes. The **default** is **kernel=rbf**.

```
[149]: kernels = ['linear', 'rbf', 'poly', 'sigmoid']

xlabel = 'Petal length'
ylabel = 'Petal width'

for kernel in kernels:
    svc_plot = svm.SVC(kernel=kernel).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel=kernel)
    title_str = 'kernel: \''+str(kernel)+'\'', '+Acc. prediction: {:.2f}%'.
    format(accuracy)
```

```
plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```




Figure 12: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')




Figure 13: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')




Figure 14: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')




Figure 15: This group of images shows the effect on the classification by the choice of the different SVC kernels ('linear', 'rbf', 'poly' and 'sigmoid')

8.4 Vary gamma parameter

The **gamma** parameter is used for **non linear hyperplanes**. The higher the **gamma** float value it tries to exactly fit the training data set. The **default** is **gamma='scale'**.

```
[175]: gammas = [0.1, 1, 10, 100, 200]

xlabel = 'Petal length'
ylabel = 'Petal width'

for gamma in gammas:
    svc_plot = svm.SVC(kernel='rbf', gamma=gamma).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', gamma=gamma)
    title_str = 'gamma: \''+str(gamma)+'\'', ' \
                + 'Acc. prediction: {:.2f}%'.format(accuracy)
    plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```




Figure 16: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

Show the variation of the SVC parameter **gamma** against the **prediction accuracy**.

As we can see, increasing **gamma** leads to **overfitting** as the classifier tries to perfectly fit the training data.




Figure 17: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel




Figure 18: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel




Figure 19: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

```
[184]: gammas = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1, 10, 100, 200]
accuracy_list = list()
```




Figure 20: This group of images shows the effect on the classification by the variation of the parameter 'gamma' of the 'rbf' kernel

```
for gamma in gammas:
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', gamma=gamma)
    accuracy_list.append(accuracy)

plotParamsAcc(gammas, accuracy_list, 'gamma', log_scale=True)
```




Figure 21: The plot shows the variation of the SVC parameter 'gamma' against the prediction accuracy

8.5 Vary C parameter

The C parameter is the **penalty** of the error term. It controls the trade off between smooth decision boundary and classifying the training points correctly. The **default** is $C=1.0$.

```
[172]: cs = [0.1, 1, 5, 10, 100, 1000, 10000]
```

```

xlabel = 'Petal length'
ylabel = 'Petal width'

for c in cs:
    svc_plot = svm.SVC(kernel='rbf', C=c).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', C=c)
    title_str = 'C: \''+str(c)+'\\', ' \
                + 'Acc. prediction: {:.2f}%'.format(accuracy)
    plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)

```




Figure 22: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel




Figure 23: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

Show the variation of the SVC parameter **C** against the **prediction accuracy**.

But be careful: to high **C** values may lead to **overfitting** the training data.

```

[185]: cs = [0.1, 1, 5, 6, 7, 8, 10, 100, 1000, 10000]

accuracy_list = list()
for c in cs:
    accuracy = crossValSVC(X_train, y_train, kernel='rbf', C=c)
    accuracy_list.append(accuracy)

plotParamsAcc(cs, accuracy_list, 'C', log_scale=True)

```




Figure 24: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel




Figure 25: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel




Figure 26: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

8.6 Vary degree parameter

The **degree** parameter is used when the **kernel** is set to **poly** and is ignored by all other kernels. It's basically the **degree of the polynomial** used to find the hyperplane to split the data. The **default** is




Figure 27: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel




Figure 28: This group of images shows the effect on the classification by the variation of the parameter 'C' of the 'rbf' kernel

degree=3.

Using `degree = 1` is the same as using a linear kernel. Also, increasing this parameters leads to **higher training times**.

```
[173]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

xlabel = 'Petal length'
ylabel = 'Petal width'

for degree in degrees:
    svc_plot = svm.SVC(kernel='poly', degree=degree).fit(X_plot, y_plot)
    accuracy = crossValSVC(X_train, y_train, kernel='poly', degree=degree)
    title_str = 'degree: \''+str(degree)+'\'', ' \
                + 'Acc. prediction: {:.2f}%'.format(accuracy)
    plotSVC(title_str, svc_plot, X_plot, y_plot, xlabel, ylabel)
```

Show the variation of the SVC parameter `degree` against the **prediction accuracy**.

As we can see, increasing the `degree` of the polynomial hyperplane leads to **overfitting** the training data.




Figure 29: The plot shows the variation of the SVC parameter 'C' against the prediction accuracy




Figure 30: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel

```
[186]: degrees = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]

accuracy_list = list()
for degree in degrees:
    accuracy = crossValSVC(X_train, y_train, kernel='poly', degree=degree)
    accuracy_list.append(accuracy)

plotParamsAcc(degrees, accuracy_list, 'degree', log_scale=False)
```

```
[ ]:
```




Figure 31: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 32: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 33: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 34: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 35: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 36: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 37: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 38: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 39: This group of images shows the effect on the classification by the variation of the parameter 'degree' of the 'poly' kernel




Figure 40: The plot shows the variation of the SVC parameter 'degree' against the prediction accuracy